

Advanced Data Structures and Algorithms

Min Cut and Karger's Algorithm


This Module

- Minimum Cuts!
 - Karger's algorithm
 - Karger-Stein algorithm
- Back to **randomized algorithms!**

*For today, all graphs are **undirected** and **unweighted**.

Recall: cuts in graphs


- A cut is a partition of the vertices into two **nonempty** parts.




*For today, all graphs are **undirected** and **unweighted**.

Recall: cuts in graphs


- A cut is a partition of the vertices into two **nonempty** parts.



This is not a cut




This is a cut



This is a cut


These edges **cross the cut**.

- They go from one part to the other.




A (global) minimum cut

is a cut that has the fewest edges possible crossing it.



A (global) minimum cut


is a cut that has the fewest edges possible crossing it.



Why might we care about global minimum cuts?


Why might we care about global minimum cuts?

- Clustering:




Why might we care about global minimum cuts?

- Clustering:




- Image Segmentation




Why might we care about global minimum cuts?

- Clustering:



- Image Segmentation



Karger's algorithm

- Finds **global minimum cuts** in undirected graphs
- Randomized algorithm

Karger's algorithm

- Finds **global minimum cuts** in undirected graphs
- Randomized algorithm
- Karger's algorithm **might be wrong**.
 - Compare to QuickSort, which just might be slow.


Karger's algorithm

- Finds **global minimum cuts** in undirected graphs
- Randomized algorithm
- Karger's algorithm **might be wrong**.
 - Compare to QuickSort, which just might be slow.
- Why would we want an algorithm that might be wrong?

Karger's algorithm


- Finds **global minimum cuts** in undirected graphs
- Randomized algorithm
- Karger's algorithm **might be wrong**.
 - Compare to QuickSort, which just might be slow.
- Why would we want an algorithm that might be wrong?
 - **With high probability it won't be wrong.**
 - Maybe the stakes are low and the cost of a deterministic algorithm is high.

Karger's algorithm




Karger's algorithm

- Pick a random edge.




Karger's algorithm

- Pick a random edge.
- **Contract** it.




Karger's algorithm

- Pick a random edge.
- **Contract** it.
- Repeat until you only have two vertices left.




Karger's algorithm

- Pick a random edge.
- **Contract** it.
- Repeat until you only have two vertices left.

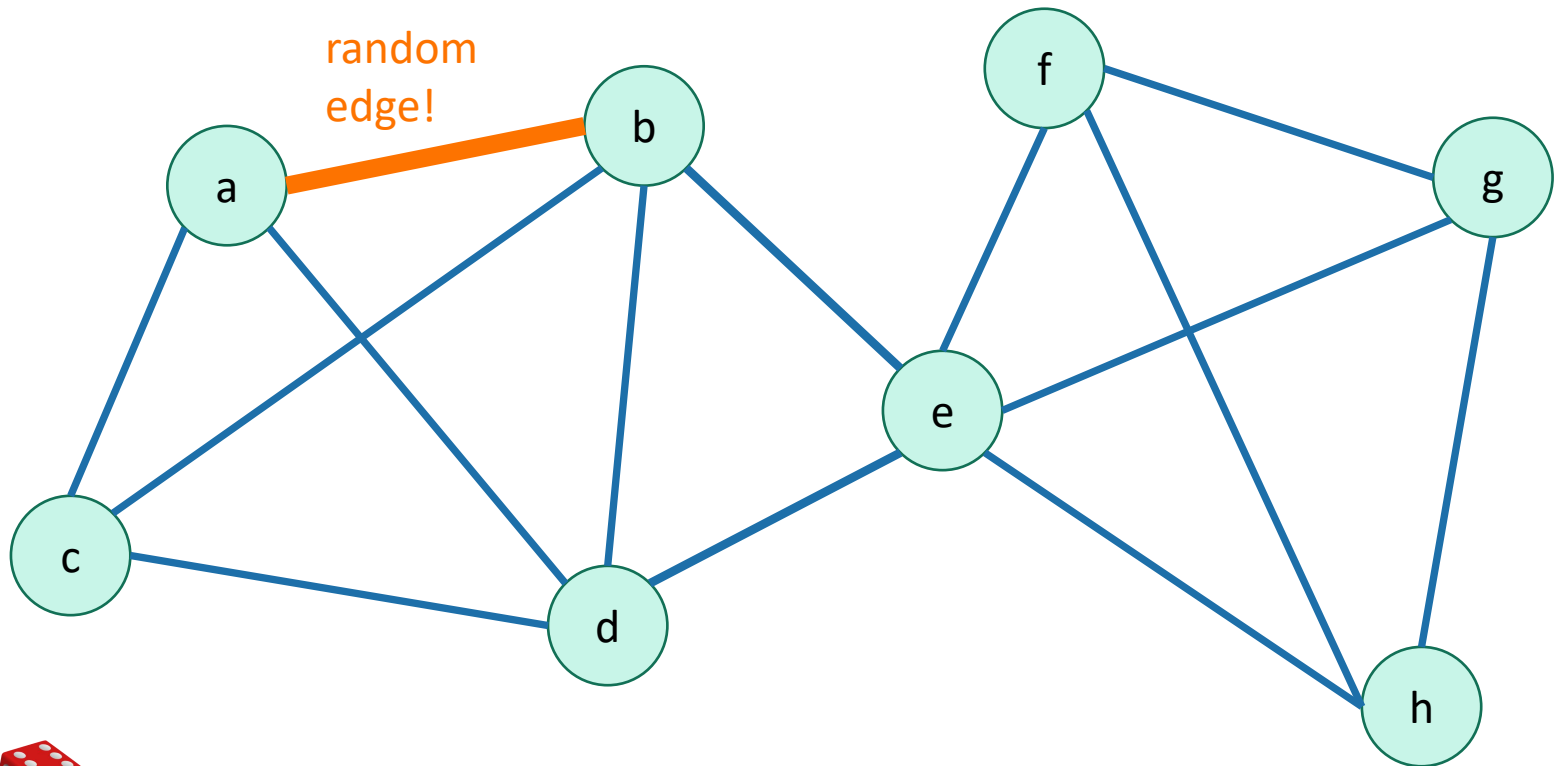


Why is this a good idea? We'll see shortly.

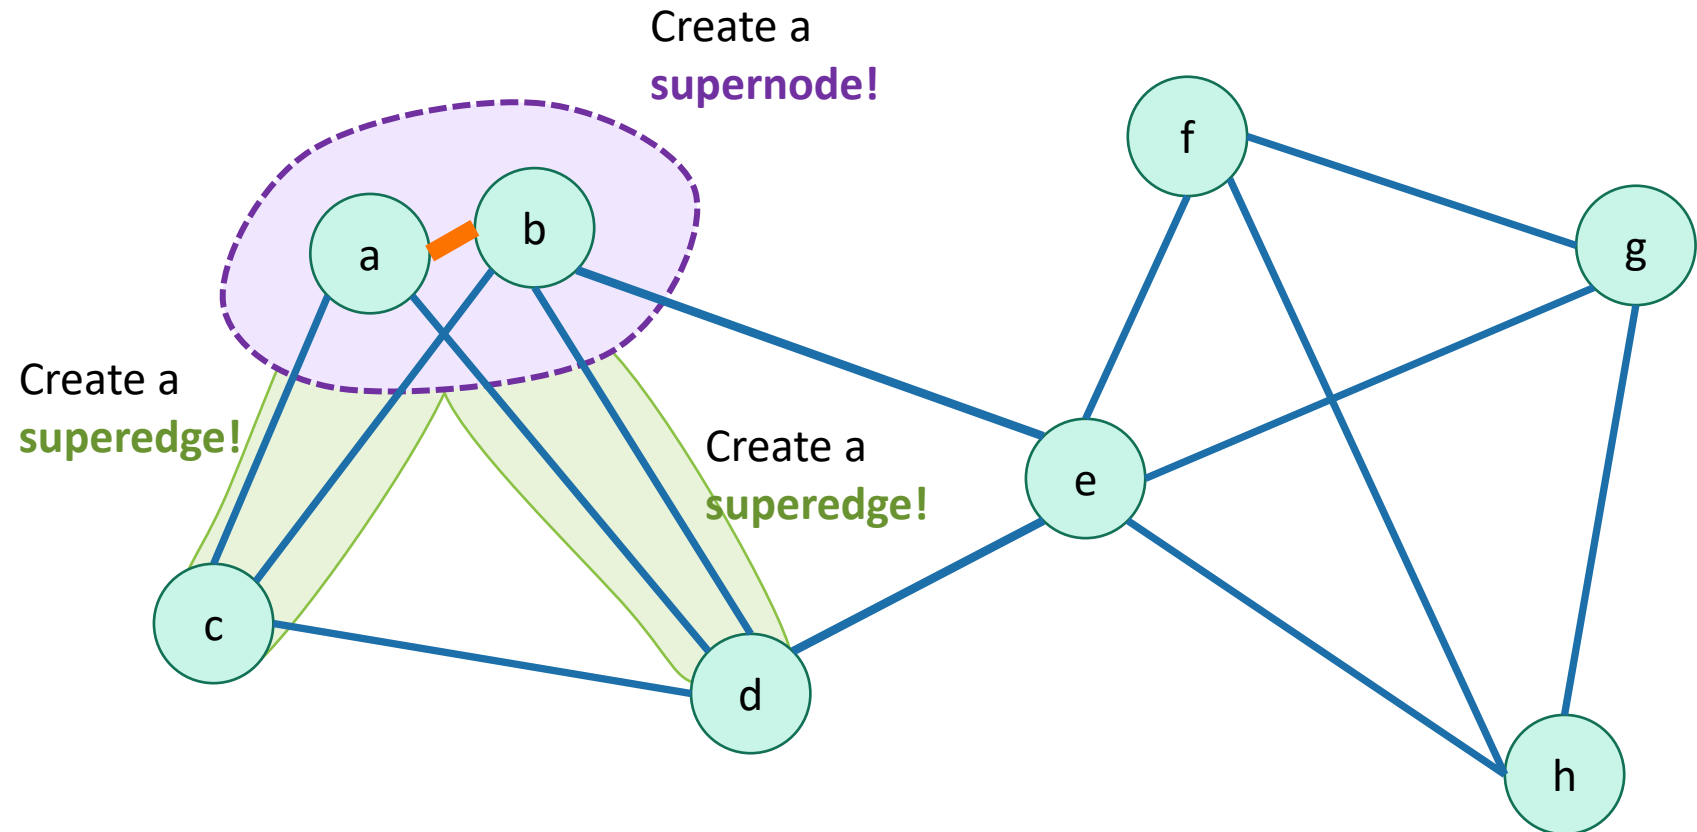
Karger's algorithm



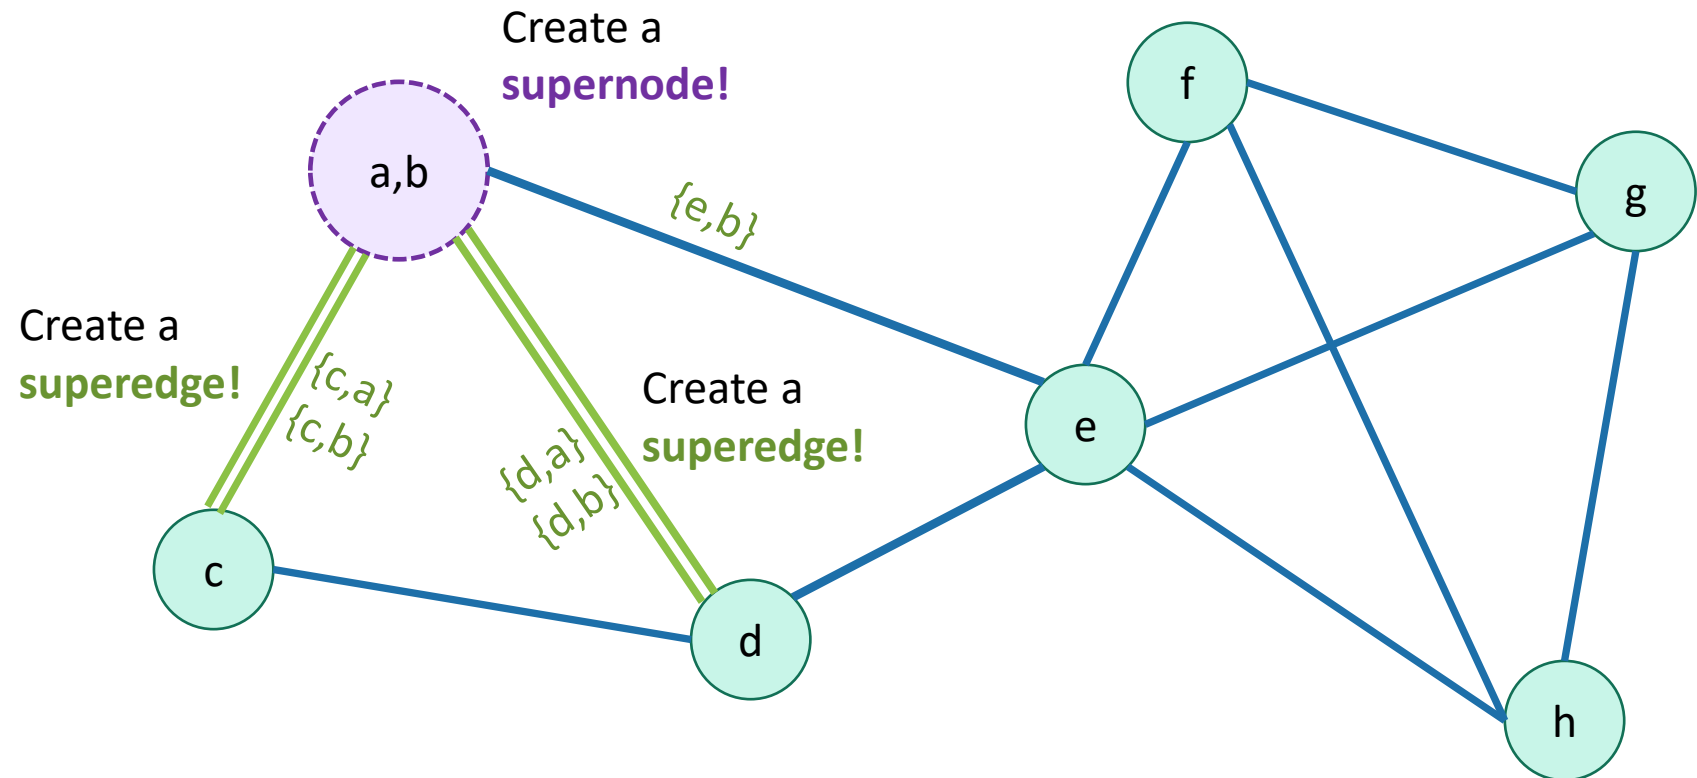
Karger's algorithm




Karger's algorithm



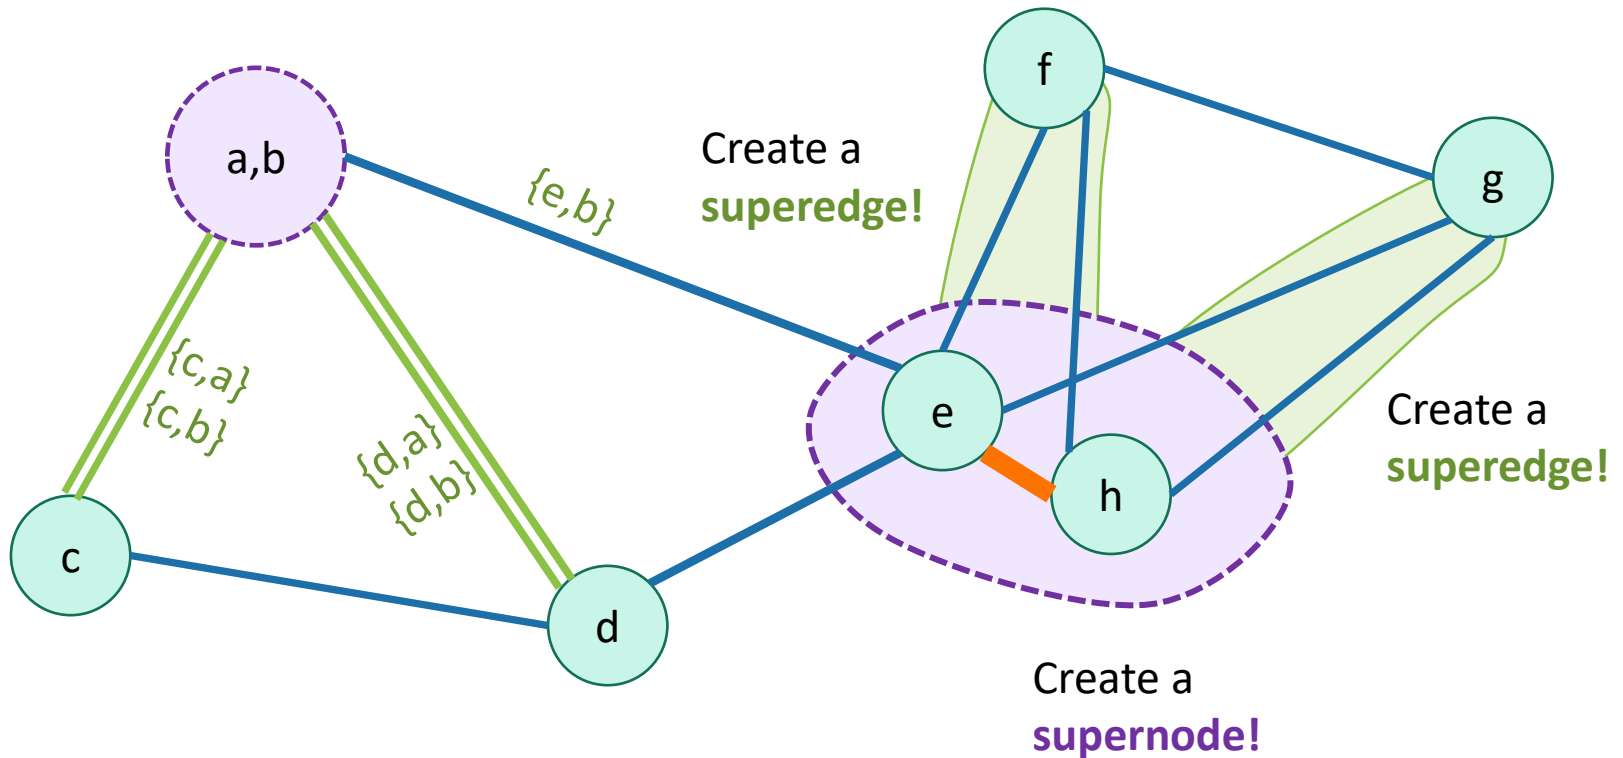
Karger's algorithm



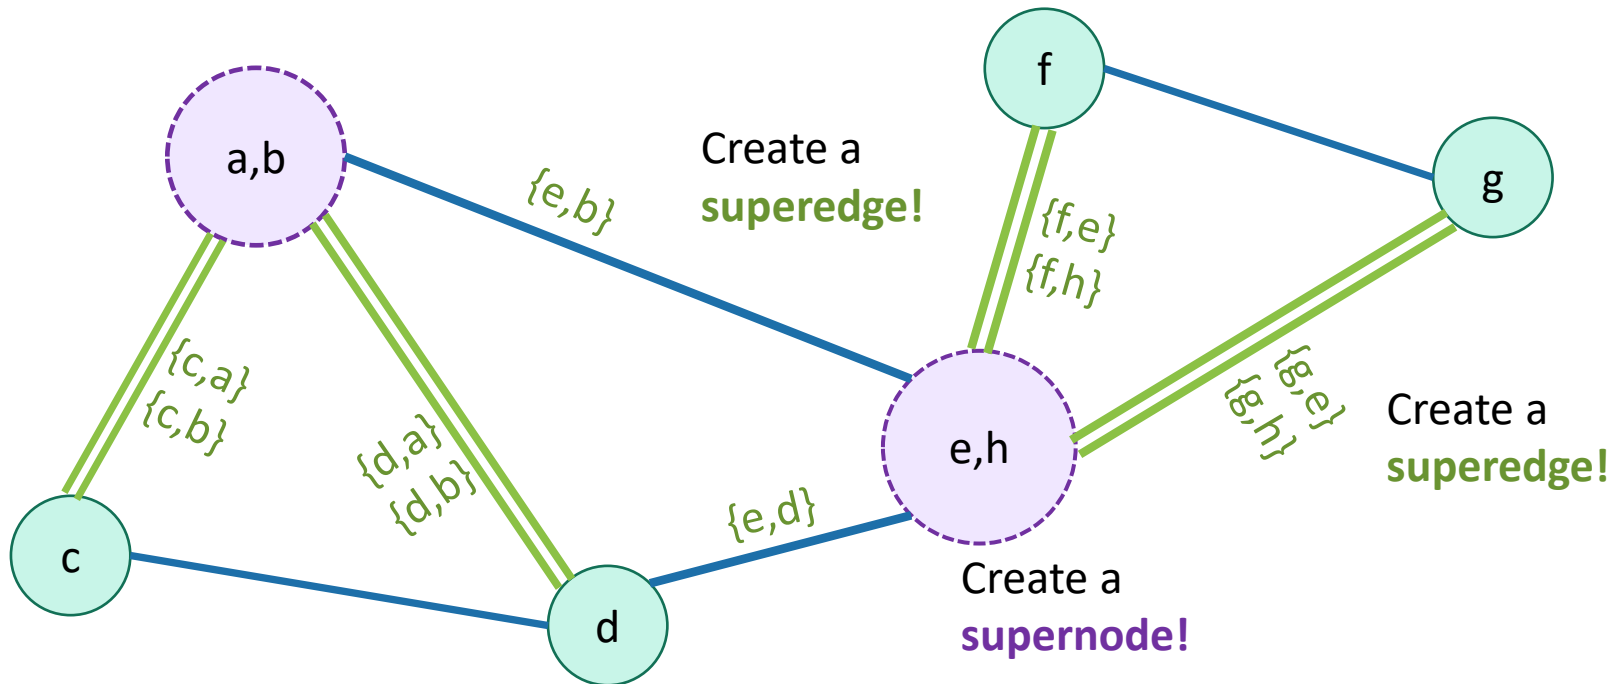
Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm




Karger's algorithm



Karger's algorithm




Karger's algorithm



Karger's algorithm

Now stop!

- There are only two nodes left.




Karger's algorithm

Now stop!

- There are only two nodes left.

The **minimum cut** is given by the remaining super-nodes:


- $\{a,b,c,d\}$ and $\{e,h,f,g\}$



Karger's algorithm

The **minimum cut** is given by the remaining super-nodes:

- $\{a,b,c,d\}$ and $\{e,h,f,g\}$



Karger's algorithm

- Does it work?

- Is it fast? 

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger($G=(V,E)$):**

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$

// one supernode for each vertex

// one superedge for each edge

// we'll choose randomly from F

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger**($G=(V,E)$):

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F
- **while** $|\Gamma| > 2$:
 - $(u,v) \leftarrow$ uniformly random edge in F

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger**($G=(V,E)$):

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F
- **while** $|\Gamma| > 2$:
 - $(u,v) \leftarrow$ uniformly random edge in F
 - **merge**(u, v)
// merge the SuperNode containing u with the SuperNode containing v .

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger**($G=(V,E)$):

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F
- **while** $|\Gamma| > 2$:
 - $(u,v) \leftarrow$ uniformly random edge in F
 - **merge**(u, v)
// merge the SuperNode containing u with the SuperNode containing v .

- **merge**(u, v):

- $\bar{x} = \text{SuperNode}(\bar{u} \cup \bar{v})$ // merge also knows about Γ and the $E_{\bar{u},\bar{v}}$'s
- **for each** \bar{w} in $\Gamma - \{\bar{u}, \bar{v}\}$:
// create a new supernode
 - $E_{\bar{x},\bar{w}} = E_{\bar{u},\bar{w}} \cup E_{\bar{v},\bar{w}}$
- Remove \bar{u} and \bar{v} from Γ and add \bar{x} .

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger($G=(V,E)$):**

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F
- **while** $|\Gamma| > 2$:
 - $(u,v) \leftarrow$ uniformly random edge in F
 - **merge**(u, v)
 - // merge the SuperNode containing u with the SuperNode containing v .
 - $F \leftarrow F - E_{\bar{u},\bar{v}}$
 - // remove all the edges in the SuperEdge between those SuperNodes.

- **merge**(u, v):

- // merge also knows about Γ and the $E_{\bar{u},\bar{v}}$'s
- $\bar{x} = \text{SuperNode}(\bar{u} \cup \bar{v})$ // create a new supernode
- for each w in $\Gamma - \{\bar{u}, \bar{v}\}$:
 - $E_{\bar{x},\bar{w}} = E_{\bar{u},\bar{w}} \cup E_{\bar{v},\bar{w}}$
- Remove \bar{u} and \bar{v} from Γ and add \bar{x} .

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger($G=(V,E)$):**

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F
- **while** $|\Gamma| > 2$:
 - $(u,v) \leftarrow$ uniformly random edge in F
 - **merge**(u, v)
 - // merge the SuperNode containing u with the SuperNode containing v .
 - $F \leftarrow F - E_{\bar{u},\bar{v}}$
 - // remove all the edges in the SuperEdge between those SuperNodes.
- **return** the cut given by the remaining two superNodes.

- **merge**(u, v):

- // merge also knows about Γ and the $E_{\bar{u},\bar{v}}$'s
- $\bar{x} = \text{SuperNode}(\bar{u} \cup \bar{v})$ // create a new supernode
- for each w in $\Gamma - \{\bar{u}, \bar{v}\}$:
 - $E_{\bar{x},\bar{w}} = E_{\bar{u},\bar{w}} \cup E_{\bar{v},\bar{w}}$
- Remove \bar{u} and \bar{v} from Γ and add \bar{x} .

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger($G=(V,E)$):**

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F

- **while** $|\Gamma| > 2$:

- $(u,v) \leftarrow$ uniformly random edge in F
- **merge**(u, v)

The **while** loop runs $n-2$ times

merge takes time $O(n)$ naively

// merge the SuperNode containing u with the SuperNode containing v .

- $F \leftarrow F - E_{\bar{u},\bar{v}}$

// remove all the edges in the SuperEdge between those SuperNodes.

- **return** the cut given by the remaining two superNodes.

- **merge**(u, v):

// merge also knows about Γ and the $E_{\bar{u},\bar{v}}$'s

- $\bar{x} = \text{SuperNode}(\bar{u} \cup \bar{v})$ // create a new supernode
- for each w in $\Gamma - \{\bar{u}, \bar{v}\}$:
 - $E_{\bar{x},\bar{w}} = E_{\bar{u},\bar{w}} \cup E_{\bar{v},\bar{w}}$
- Remove \bar{u} and \bar{v} from Γ and add \bar{x} .

Pseudocode

Let \bar{u} denote the SuperNode in Γ containing u
Say $E_{\bar{u},\bar{v}}$ is the SuperEdge between \bar{u}, \bar{v} .

- **Karger**($G=(V,E)$):

- $\Gamma = \{ \text{SuperNode}(v) : v \text{ in } V \}$ // one supernode for each vertex
- $E_{\bar{u},\bar{v}} = \{(u,v)\}$ for (u,v) in E // one superedge for each edge
- $E_{\bar{u},\bar{v}} = \{\}$ for (u,v) not in E .
- $F = \text{copy of } E$ // we'll choose randomly from F

- **while** $|\Gamma| > 2$:

- $(u,v) \leftarrow$ uniformly random edge in F
- **merge**(u, v)

The **while** loop runs $n-2$ times

merge takes time $O(n)$ naively

// merge the SuperNode containing u with the SuperNode containing v .

- $F \leftarrow F - E_{\bar{u},\bar{v}}$

// remove all the edges in the SuperEdge between those SuperNodes.

- **return** the cut given by the remaining two superNodes.

- **merge**(u, v):

// merge also knows about Γ and the $E_{\bar{u},\bar{v}}$'s

- $\bar{x} = \text{SuperNode}(\bar{u} \cup \bar{v})$ // create a new supernode
- for each \bar{w} in $\Gamma - \{\bar{u}, \bar{v}\}$:
 - $E_{\bar{x},\bar{w}} = E_{\bar{u},\bar{w}} \cup E_{\bar{v},\bar{w}}$
- Remove \bar{u} and \bar{v} from Γ and add \bar{x} .

total runtime $O(n^2)$

We can do a bit better with fancy data structures, but let's go with this for now.

How do we implement this?

- Maintain a secondary “superGraph” which keeps track of superNodes and superEdges
- Running time?
 - We contract $n-2$ edges
 - Each time we contract an edge we get rid of a vertex, and we get rid of $n - 2$ vertices total.
 - Naively each contraction takes time $O(n)$
 - Maybe there are $\Omega(n)$ nodes in the superNodes that we are merging. (We can do better with fancy data structures).
 - So total running time $O(n^2)$.
 - We can do $O(m \cdot \alpha(n))$ with a union-find data structure, but $O(n^2)$ is good enough for today.

Karger's algorithm

- Does it work?
- Is it fast?
 - $O(n^2)$

Karger's algorithm


- Does it work?



- Is it fast?
 - $O(n^2)$


Karger's algorithm

- Does it work?



- Is it fast?

- $O(n^2)$




Algorithm:

- Randomly contract edges until there are only two supernodes left.

Karger's algorithm


- Does it work?



Think-pair-share!

- Is it fast?

- $O(n^2)$




Algorithm:

- Randomly contract edges until there are only two supernodes left.

Karger's algorithm

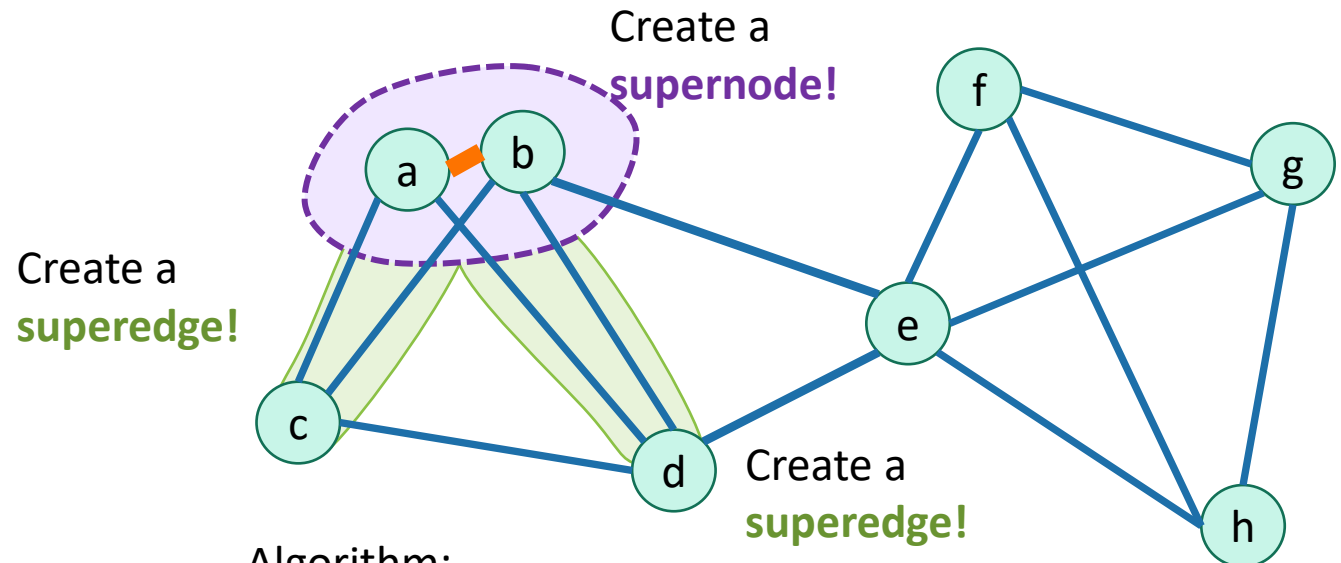
- Does it work?



No?

- Is it fast?

- $O(n^2)$




Algorithm:

- Randomly contract edges until there are only two supernodes left.

Why did that work?


Why did that work?

- We got really lucky!
- This could have gone wrong in so many ways.




Karger's algorithm

Say we had chosen this edge



Karger's algorithm


Say we had chosen this edge



Karger's algorithm

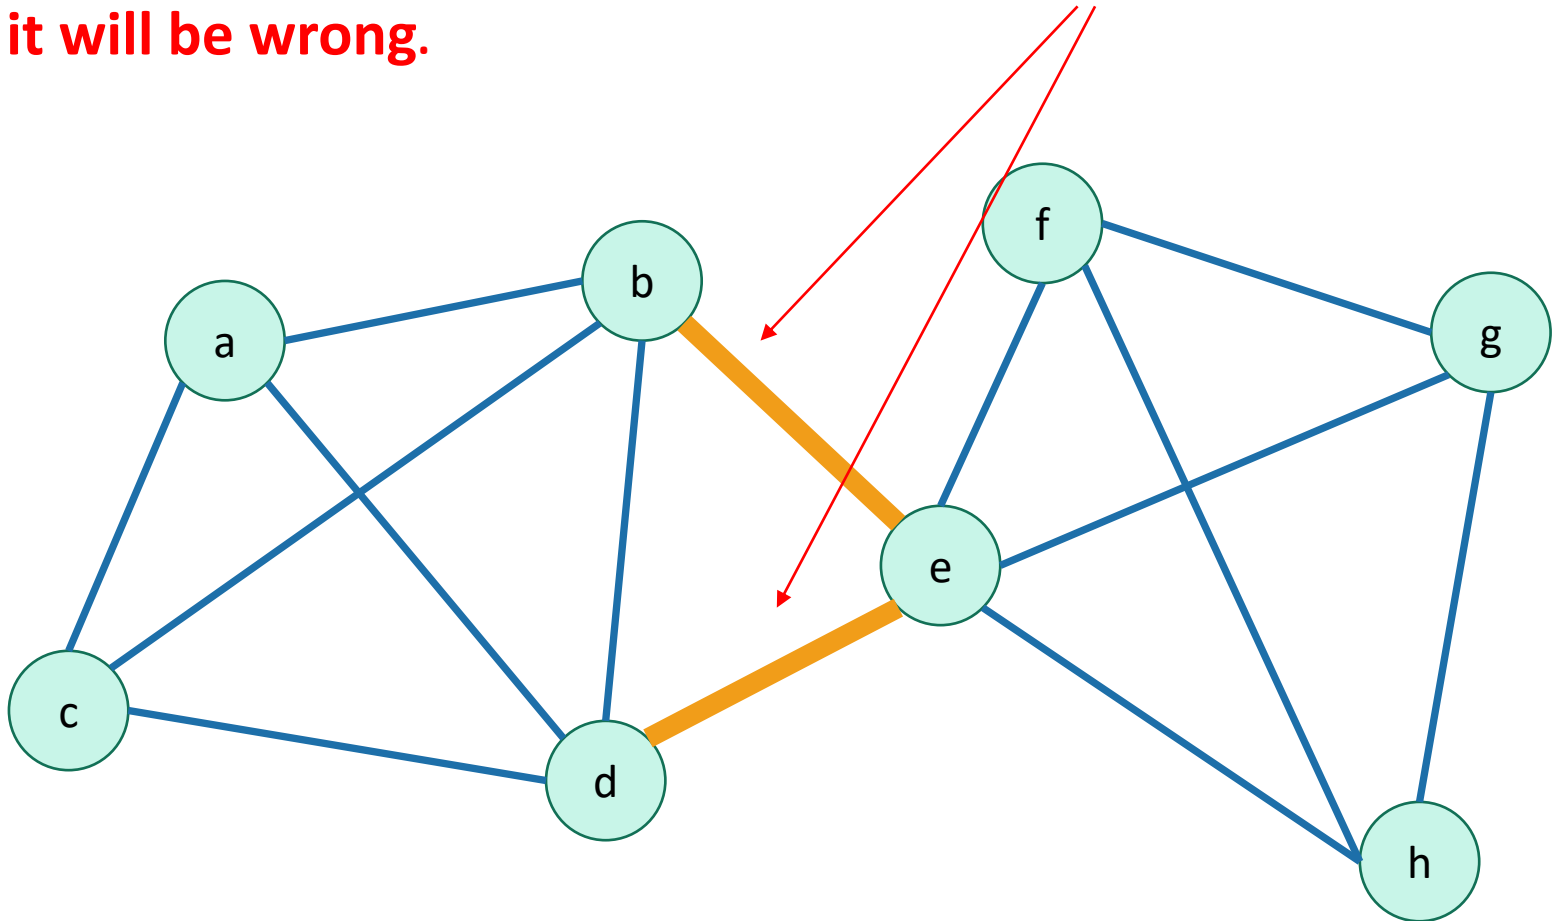
Say we had chosen this edge

Now there is **no way** we could return a cut that separates b and e.




Even worse


If the algorithm **EVER** chooses either of **these edges**,
it will be wrong.




How likely is that?




- For this particular graph, if we do 10,000 times:



How likely is that?



- For this particular graph, if we do 10,000 times:




**The algorithm
is only correct
about 20% of
the time!**

That doesn't sound good

That doesn't sound good

- Too see why it's good after all, we'll do a case study of this graph.




That doesn't sound good

- Too see why it's good after all, we'll do a case study of this graph.

The plan:

- See that 20% chance of correctness is actually nontrivial.
- Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct 99% of the time.




That doesn't sound good

- To see why it's good after all, we'll do a case study of this graph.

The plan:

- See that 20% chance of correctness is actually nontrivial.
- Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct 99% of the time.



- To see the first point, let's compare Karger's algorithm to the algorithm:


*Choose a completely random cut
and hope that it's a minimum cut.*

Uniformly random cut

- Pick a random way to split the vertices into two parts:


Uniformly random cut

- Pick a random way to split the vertices into two parts:



Uniformly random cut

- Pick a random way to split the vertices into two parts:




Uniformly random cut

- Pick a random way to split the vertices into two parts:
- The probability of choosing the minimum cut is*...

$$\frac{\text{number of min cuts in that graph}}{\text{number of ways to split 8 vertices in 2 parts}} = \frac{2}{2^8 - 2} \approx 0.008$$


- Aka, we get a minimum cut **0.8% of the time.**

Karger is better than completely random!



Karger's alg. is correct about 20% of the time


Completely random is correct about 0.8% of the time



What's going on?


What's going on?

- Which is more likely?




A: The algorithm never chooses either of the edges in **the minimum cut**.

B: The algorithm never chooses any of the edges in **this big cut**.




What's going on?

- Which is more likely?



A: The algorithm never chooses either of the edges in **the minimum cut**.


B: The algorithm never chooses any of the edges in **this big cut**.



- Neither A nor B are very likely, **but A is more likely than B.**

What's going on?

- Which is more likely?




A: The algorithm never chooses either of the edges in **the minimum cut**.

Thing 1: It's unlikely that Karger will hit the min cut since it's so small!




B: The algorithm never chooses any of the edges in **this big cut**.




- Neither A nor B are very likely, **but A is more likely than B.**

What's going on?




A: This cut can be returned
by Karger's algorithm.

What's going on?




A: This cut can be returned by Karger's algorithm.

B: This cut can't be returned by Karger's algorithm!
(Because how would a and g end up in the same super-node?)




What's going on?



A: This cut can be returned by Karger's algorithm.


B: This cut can't be returned by Karger's algorithm!
(Because how would a and g end up in the same super-node?)



This cut actually separates the graph into three pieces, so it's not minimal – either half of it is a smaller cut.


What's going on?

Thing 2: By only contracting edges we are ignoring certain really-not-minimal cuts.



A: This cut can be returned by Karger's algorithm.

B: This cut can't be returned by Karger's algorithm!
(Because how would a and g end up in the same super-node?)



This cut actually separates the graph into three pieces, so it's not minimal – either half of it is a smaller cut.

Why does that help?

- Okay, so it's better than completely random...
- We're still wrong about 80% of the time.

The plan:



- See that 20% chance of correctness is actually nontrivial.
- Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct 99% of the time.

Why does that help?

- Okay, so it's better than completely random...
- We're still wrong about 80% of the time.
- The main idea: **repeat!**
 - If I'm wrong 20% of the time, then if I repeat it a few times I'll eventually get it right.

The plan:



- See that 20% chance of correctness is actually nontrivial.
- Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct 99% of the time.

Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?

4



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



Thought experiment

- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



4 3 3 2 5 2 5

Thought experiment


- Suppose you have a magic button that produces one of 5 numbers, $\{a,b,c,d,e\}$, uniformly at random when you push it.
- You don't know what $\{a,b,c,d,e\}$ are.
- Q: What is the minimum of a,b,c,d,e ?



4 3 3 2 5 2 5

How many times do you have to push the button, in expectation, before you see the minimum value?

What is the probability that you have to push it more than 5 times? 10 times?



$$\Pr[\text{We push the button once and don't ever get the min}] = (1 - 0.2) = 0.8$$

$$\Pr[\text{We push the button once and don't ever get the min}] = (1 - 0.2) = 0.8$$

$$E[\text{Number of times we push the button until we get the minimum value}] = 1/(0.20) = 5$$

$$\Pr[\text{We push the button once and don't ever get the min}] = (1 - 0.2) = 0.8$$

$$E[\text{Number of times we push the button until we get the minimum value}] = 1/(0.20) = 5$$

$$\Pr[\text{We push the button } t \text{ times and don't ever get the min}] = (1 - 0.2)^t$$

$$\Pr[\text{We push the button once and don't ever get the min}] = (1 - 0.2) = 0.8$$

$$E[\text{Number of times we push the button until we get the minimum value}] = 1/(0.20) = 5$$

$$\Pr[\text{We push the button } t \text{ times and don't ever get the min}] = (1 - 0.2)^t$$

$$\Pr[\text{We push the button 5 times and don't ever get the min}] = (1 - 0.2)^5 \approx 0.33$$

$$\Pr[\text{We push the button once and don't ever get the min}] = (1 - 0.2) = 0.8$$

$$E[\text{Number of times we push the button until we get the minimum value}] = 1/(0.20) = 5$$

$$\Pr[\text{We push the button } t \text{ times and don't ever get the min}] = (1 - 0.2)^t$$

$$\Pr[\text{We push the button 5 times and don't ever get the min}] = (1 - 0.2)^5 \approx 0.33$$

$$\Pr[\text{We push the button 10 times and don't ever get the min}] = (1 - 0.2)^{10} \approx 0.1$$

In this context

(Lets Run Karger's!, 5 times)

In this context

(Lets Run Karger's!, 5 times)



- Run Karger's! The cut size is 4!

In this context

(Lets Run Karger's!, 5 times)



- Run Karger's! The cut size is 4!



- Run Karger's! The cut size is 3!

In this context

(Lets Run Karger's!, 5 times)



- Run Karger's! The cut size is 4!



- Run Karger's! The cut size is 3!



- Run Karger's! The cut size is 3!

In this context

(Lets Run Karger's!, 5 times)



- Run Karger's! The cut size is 4!



- Run Karger's! The cut size is 3!



- Run Karger's! The cut size is 3!



- Run Karger's! The cut size is 2!

In this context

(Lets Run Karger's!, 5 times)



• Run Karger's! The cut size is 4!



• Run Karger's! The cut size is 3!



• Run Karger's! The cut size is 3!



• Run Karger's! The cut size is 2!



• Run Karger's! The cut size is 5!

In this context

(Lets Run Karger's!, 5 times)



- Run Karger's! The cut size is 4!




- Run Karger's! The cut size is 3!



- Run Karger's! The cut size is 3!



- Run Karger's! The cut size is 2!



Correct!



- Run Karger's! The cut size is 5!

In this context

(Lets Run Karger's!, 5 times)



• Run Karger's! The cut size is 4!



• Run Karger's! The cut size is 3!



• Run Karger's! The cut size is 3!



• Run Karger's! The cut size is 2!




Correct!



• Run Karger's! The cut size is 5!

If the success probability is about 20%, then if you run Karger's algorithm 5 times and take the best answer you get, that will likely be correct! (with probability about 0.66)

For this particular graph




- Repeat Karger's algorithm about 5 times, and we will get a min cut with decent probability.
 - In contrast, we'd have to choose a random cut about $1/0.008 = 125$ times!

The plan:

- ✓ • See that 20% chance of correctness is actually nontrivial.
- ✓ • Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct most of the time.

For this particular graph



- Repeat Karger's algorithm about 5 times, and we will get a min cut with decent probability.
 - In contrast, we'd have to choose a random cut about $1/0.008 = 125$ times!


Hang on! This “20%” figure just came from running experiments on this particular graph. What about general graphs? Can we prove something?



The plan:

- ✓ See that 20% chance of correctness is actually nontrivial.
- ✓ Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct most of the time.

For this particular graph



- Repeat Karger's algorithm about 5 times, and we will get a min cut with decent probability.
 - In contrast, we'd have to choose a random cut about $1/0.008 = 125$ times!

Hang on! This “20%” figure just came from running experiments on this particular graph. What about general graphs? Can we prove something?



Also, we should be a bit more precise about this “about 5 times” statement.

The plan:

- ✓ See that 20% chance of correctness is actually nontrivial.
- ✓ Use repetition to boost an algorithm that's correct 20% of the time to an algorithm that's correct most of the time.

Questions



To generalize this approach to all graphs


1. What is the probability that Karger's algorithm returns a minimum cut?
2. How many times should we run Karger's algorithm to "probably" succeed?
 - Say, with probability 0.99?
 - Or more generally, probability $1 - \delta$?

Answer to Question 1

Claim:


The probability that Karger's algorithm returns a minimum cut is

at least $1/\binom{n}{2}$



In this case, $1/\binom{8}{2} = 0.036$, so we are guaranteed to win at least 3.6% of the time.

Questions





1. What is the probability that Karger's algorithm returns a minimum cut?

According to the claim, at least $\frac{1}{\binom{n}{2}}$

2. How many times should we run Karger's algorithm to “probably” succeed?
 - Say, with probability 0.99?
 - Or more generally, probability $1 - \delta$?

Before we prove the Claim

2. How many times should we run Karger's algorithm to succeed with probability $1 - \delta$?



A computation

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.

A computation

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}] = (1 - p)^T$

A computation

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.

A computation


- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}]$
 $= (1 - p)^T$

A computation

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}]$
 - $= (1 - p)^T$
 - $\leq (e^{-p})^T$

A computation


- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials}]$
 $= (1 - p)^T$
 $\leq (e^{-p})^T$



$$1 - p \leq e^{-p}$$

A computation


- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }]$
 - $= (1 - p)^T$
 - $\leq (e^{-p})^T$
 - $= e^{-pT}$



$$1 - p \leq e^{-p}$$

A computation


- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials}]$
 - $= (1 - p)^T$
 - $\leq (e^{-p})^T$
 - $= e^{-pT}$
 - $= e^{-\ln(1/\delta)}$



$$1 - p \leq e^{-p}$$

A computation

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0,1)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{ don't return a min cut in } T \text{ trials }]$
 - $= (1 - p)^T$
 - $\leq (e^{-p})^T$
 - $= e^{-pT}$
 - $= e^{-\ln(1/\delta)}$
 - $= \delta$




A computation

Punchline:


If we repeat $T = \binom{n}{2} \ln(1/\delta)$ times, we win with probability at least $1 - \delta$.

- Suppose :
 - the probability of successfully returning a minimum cut is $p \in [0, 1]$,
 - we want failure probability at most $\delta \in (0, 1)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}] = (1 - p)^T$
- So $p = 1/\binom{n}{2}$ by the Claim. Let's choose $T = \binom{n}{2} \ln(1/\delta)$.
- $\Pr[\text{don't return a min cut in } T \text{ trials}]$
 - $= (1 - p)^T$
 - $\leq (e^{-p})^T$
 - $= e^{-pT}$
 - $= e^{-\ln(1/\delta)}$
 - $= \delta$



$$1 - p \leq e^{-p}$$

Answers



1. What is the probability that Karger's algorithm returns a minimum cut?

According to the claim, at least $\frac{1}{\binom{n}{2}}$

2. How many times should we run Karger's algorithm to “probably” succeed?
 - Say, with probability 0.99?
 - Or more generally, probability $1 - \delta$?

$\binom{n}{2} \ln \left(\frac{1}{\delta} \right)$ times.

Theorem

Assuming the claim about $1/\binom{n}{2} \dots$

- Suppose G has n vertices.
- Consider the following algorithm:
 - $\text{bestCut} = \text{None}$
 - **for** $t = 1, \dots, \binom{n}{2} \ln \left(\frac{1}{\delta} \right) :$
 - $\text{candidateCut} \leftarrow \text{Karger}(G)$
 - **if** candidateCut is smaller than bestCut :
 - $\text{bestCut} \leftarrow \text{candidateCut}$
 - **return** bestCut
- Then $\Pr[\text{this doesn't return a min cut}] \leq \delta$.

What's the running time?

What's the running time?

- $\binom{n}{2} \ln \left(\frac{1}{\delta} \right)$ repetitions, and $O(n^2)$ per repetition.
- So, $O \left(n^2 \cdot \binom{n}{2} \ln \left(\frac{1}{\delta} \right) \right) = O(n^4)$ Treating δ as constant.

What's the running time?

- $\binom{n}{2} \ln \left(\frac{1}{\delta} \right)$ repetitions, and $O(n^2)$ per repetition.
- So, $O \left(n^2 \cdot \binom{n}{2} \ln \left(\frac{1}{\delta} \right) \right) = O(n^4)$ Treating δ as constant.

Again we can do better with a
union-find data structure.



Theorem

Assuming the claim about $1/\binom{n}{2}$...

Suppose G has n vertices. Then [repeating Karger's algorithm] finds a min cut in G with probability at least 0.99 in time $O(n^4)$.


Now let's prove the claim...

Claim

The probability that Karger's algorithm returns a minimum cut is

at least $1/\binom{n}{2}$


Now let's prove that claim
Say that S^* is a minimum cut.



Now let's prove that claim

Say that S^* is a minimum cut.


- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]



Now let's prove that claim

Say that S^* is a minimum cut.


- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]
= **PR**[e_1 doesn't cross S^*]



Now let's prove that claim

Say that S^* is a minimum cut.


- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- $\mathbf{PR[return } S^* \text{]} = \mathbf{PR[none of the } e_i \text{ cross } S^* \text{]}$
= $\mathbf{PR[} e_1 \text{ doesn't cross } S^* \text{]}$
× $\mathbf{PR[} e_2 \text{ doesn't cross } S^* \mid e_1 \text{ doesn't cross } S^* \text{]}$



Now let's prove that claim


Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- $\mathbf{PR}[\text{return } S^*] = \mathbf{PR}[\text{none of the } e_i \text{ cross } S^*]$
= $\mathbf{PR}[e_1 \text{ doesn't cross } S^*]$
× $\mathbf{PR}[e_2 \text{ doesn't cross } S^* \mid e_1 \text{ doesn't cross } S^*]$
...
× $\mathbf{PR}[e_{n-2} \text{ doesn't cross } S^* \mid e_1, \dots, e_{n-3} \text{ don't cross } S^*]$



Focus in on:


$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$


- Suppose: After $j-1$ iterations, we haven't messed up yet!



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$


- Suppose: After $j-1$ iterations, we haven't messed up yet!



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$


- Suppose: After $j-1$ iterations, we haven't messed up yet!
- What's the probability of messing up now?



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$


- Suppose: After $j-1$ iterations, we haven't messed up yet!
- What's the probability of messing up now?



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$


- Suppose: After $j-1$ iterations, we haven't messed up yet!
- What's the probability of messing up now?
- Say there are k edges that cross S^*
- Every supernode has at least k (original) edges coming out.
 - Otherwise we'd have a smaller cut.



Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$

- Suppose: After $j-1$ iterations, we haven't messed up yet!
- What's the probability of messing up now?
- Say there are k edges that cross S^*
- Every supernode has at least k (original) edges coming out.
 - Otherwise we'd have a smaller cut.
- Thus, there are at least $(n-j+1)k/2$ edges total.
 - because there are $n - j + 1$ supernodes left, each with k edges.




Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$

- Suppose: After $j-1$ iterations, we haven't messed up yet!
- What's the probability of messing up now?
- Say there are k edges that cross S^*
- Every supernode has at least k (original) edges coming out.
 - Otherwise we'd have a smaller cut.
- Thus, there are at least $(n-j+1)k/2$ edges total.
 - because there are $n-j+1$ supernodes left, each with k edges.

So the probability that we choose one of the k edges crossing S^* at step j is at most:

$$\frac{k}{\left(\frac{(n-j+1)k}{2}\right)} = \frac{2}{n-j+1}$$




Focus in on:

$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$

- So the probability that we choose one of the k edges crossing S^* at step j is at most:

$$\frac{k}{\binom{(n-j+1)k}{2}} = \frac{2}{n-j+1}$$



Focus in on:


$\text{PR}[e_j \text{ doesn't cross } S^* \mid e_1, \dots, e_{j-1} \text{ don't cross } S^*]$

- So the probability that we choose one of the k edges crossing S^* at step j is at most:

$$\frac{k}{\left(\frac{(n-j+1)k}{2}\right)} = \frac{2}{n-j+1}$$

- The probability we **don't** choose one of the k edges is at least:


$$1 - \frac{2}{n-j+1} = \frac{n-j-1}{n-j+1}$$



Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- $\mathbf{PR}[\text{return } S^*] = \mathbf{PR}[\text{none of the } e_i \text{ cross } S^*]$
= $\mathbf{PR}[e_1 \text{ doesn't cross } S^*]$
× $\mathbf{PR}[e_2 \text{ doesn't cross } S^* \mid e_1 \text{ doesn't cross } S^*]$
...
× $\mathbf{PR}[e_{n-2} \text{ doesn't cross } S^* \mid e_1, \dots, e_{n-3} \text{ don't cross } S^*]$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}

- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$= \binom{n-2}{n} \binom{n-3}{n-1} \binom{n-4}{n-2} \binom{n-5}{n-3} \binom{n-6}{n-4} \cdots \binom{4}{6} \binom{3}{5} \binom{2}{4} \binom{1}{3}$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$= \cancel{\left(\frac{n-2}{n}\right)} \left(\frac{n-3}{n-1}\right) \cancel{\left(\frac{n-4}{n-2}\right)} \left(\frac{n-5}{n-3}\right) \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{4}{6}\right) \left(\frac{3}{5}\right) \left(\frac{2}{4}\right) \left(\frac{1}{3}\right)$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{4}{6}\right) \left(\frac{3}{5}\right) \left(\frac{2}{4}\right) \left(\frac{1}{3}\right)$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \cancel{\left(\frac{n-6}{n-4}\right)} \cdots \left(\frac{4}{6}\right) \left(\frac{3}{5}\right) \left(\frac{2}{4}\right) \left(\frac{1}{3}\right)$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \cancel{\left(\frac{n-6}{n-4}\right)} \dots \cancel{\left(\frac{4}{6}\right)} \cancel{\left(\frac{3}{5}\right)} \cancel{\left(\frac{2}{4}\right)} \cancel{\left(\frac{1}{3}\right)}$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$\begin{aligned} &= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \cancel{\left(\frac{n-6}{n-4}\right)} \cdots \cancel{\left(\frac{4}{6}\right)} \cancel{\left(\frac{3}{5}\right)} \cancel{\left(\frac{2}{4}\right)} \cancel{\left(\frac{1}{3}\right)} \\ &= \left(\frac{2}{n(n-1)}\right) \end{aligned}$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$\begin{aligned} &= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \cancel{\left(\frac{n-6}{n-4}\right)} \cdots \cancel{\left(\frac{4}{6}\right)} \cancel{\left(\frac{3}{5}\right)} \cancel{\left(\frac{2}{4}\right)} \cancel{\left(\frac{1}{3}\right)} \\ &= \left(\frac{2}{n(n-1)}\right) \\ &= \frac{1}{\binom{n}{2}} \end{aligned}$$




Now let's prove that claim

Say that S^* is a minimum cut.

- Suppose the edges that we choose are e_1, e_2, \dots, e_{n-2}
- **PR**[return S^*] = **PR**[none of the e_i cross S^*]

$$\begin{aligned} &= \cancel{\left(\frac{n-2}{n}\right)} \cancel{\left(\frac{n-3}{n-1}\right)} \cancel{\left(\frac{n-4}{n-2}\right)} \cancel{\left(\frac{n-5}{n-3}\right)} \cancel{\left(\frac{n-6}{n-4}\right)} \cdots \cancel{\left(\frac{4}{6}\right)} \cancel{\left(\frac{3}{5}\right)} \cancel{\left(\frac{2}{4}\right)} \cancel{\left(\frac{1}{3}\right)} \\ &= \left(\frac{2}{n(n-1)}\right) \\ &= \frac{1}{\binom{n}{2}} \end{aligned}$$

**CLAIM
PROVED**



Theorem

Assuming the claim about $1/\binom{n}{2}$...

Suppose G has n vertices. Then [repeating Karger's algorithm] finds a min cut in G with probability at least 0.99 in time $O(n^4)$.

**That proves this
Theorem!**

What have we learned?

- If we randomly contract edges:
 - It's unlikely that we'll end up with a min cut.
 - But it's not **TOO** unlikely
 - By repeating, we likely will find a min cut.
- Repeating this process:
 - Finds a **global min cut in time $O(n^4)$, with probability 0.99.**
 - We can run a bit faster if we use a **union-find** data structure.

Here I chose $\delta = 0.01$
just for concreteness.



More generally

- If we have a Monte-Carlo algorithm¹ with a small success probability,
- and we can check how good a solution is,
- Then we can **boost** the success probability by repeating it a bunch and taking the best solution.

¹In computing, a Monte Carlo algorithm is a randomized algorithm whose output may be incorrect with a certain (typically small) probability.



Can we do better?

- Repeating $O(n^2)$ times is pretty expensive.
 - $O(n^4)$ total runtime to get success probability 0.99.

Can we do better?


- Repeating $O(n^2)$ times is pretty expensive.
 - $O(n^4)$ total runtime to get success probability 0.99.
- The **Karger-Stein Algorithm** will do better!
 - The trick is that we'll do the repetitions in a clever way.
 - $O(n^2 \log^2(n))$ runtime for the same success probability.
 - **Warning!** This is a tricky algorithm! We'll sketch the approach here: the important part is the high-level idea, not the details of the computations.

Can we do better?

- Repeating $O(n^2)$ times is pretty expensive.
 - $O(n^4)$ total runtime to get success probability 0.99.
- The **Karger-Stein Algorithm** will do better!
 - The trick is that we'll do the repetitions in a clever way.
 - $O(n^2 \log^2(n))$ runtime for the same success probability.
 - **Warning!** This is a tricky algorithm! We'll sketch the approach here: the important part is the high-level idea, not the details of the computations.

To see how we might save on repetitions,
let's run through Karger's algorithm again.

Karger's algorithm

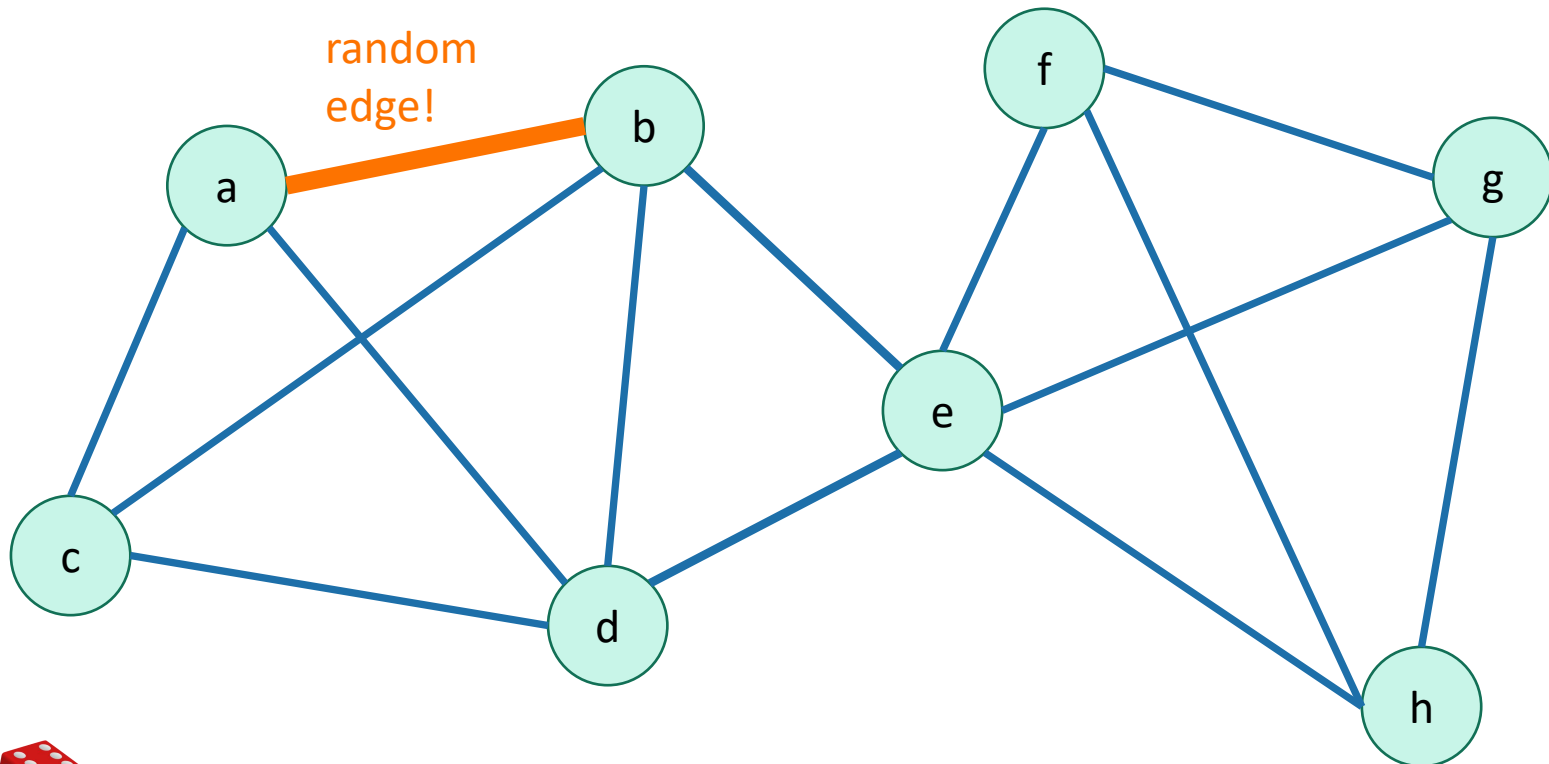


Karger's algorithm

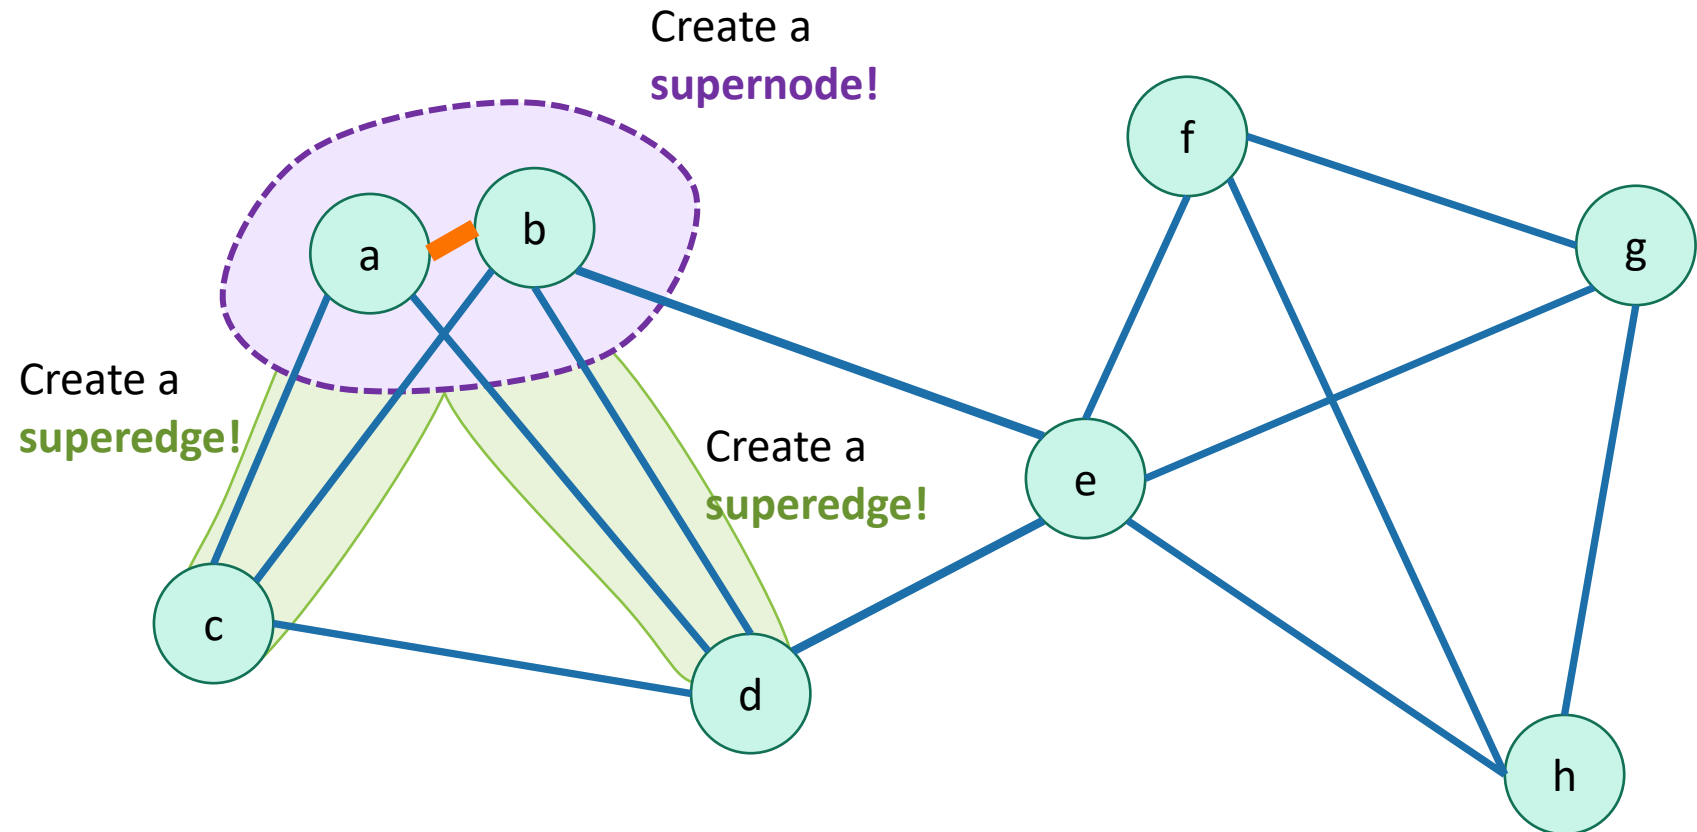
Probability that we didn't mess up:

12/14

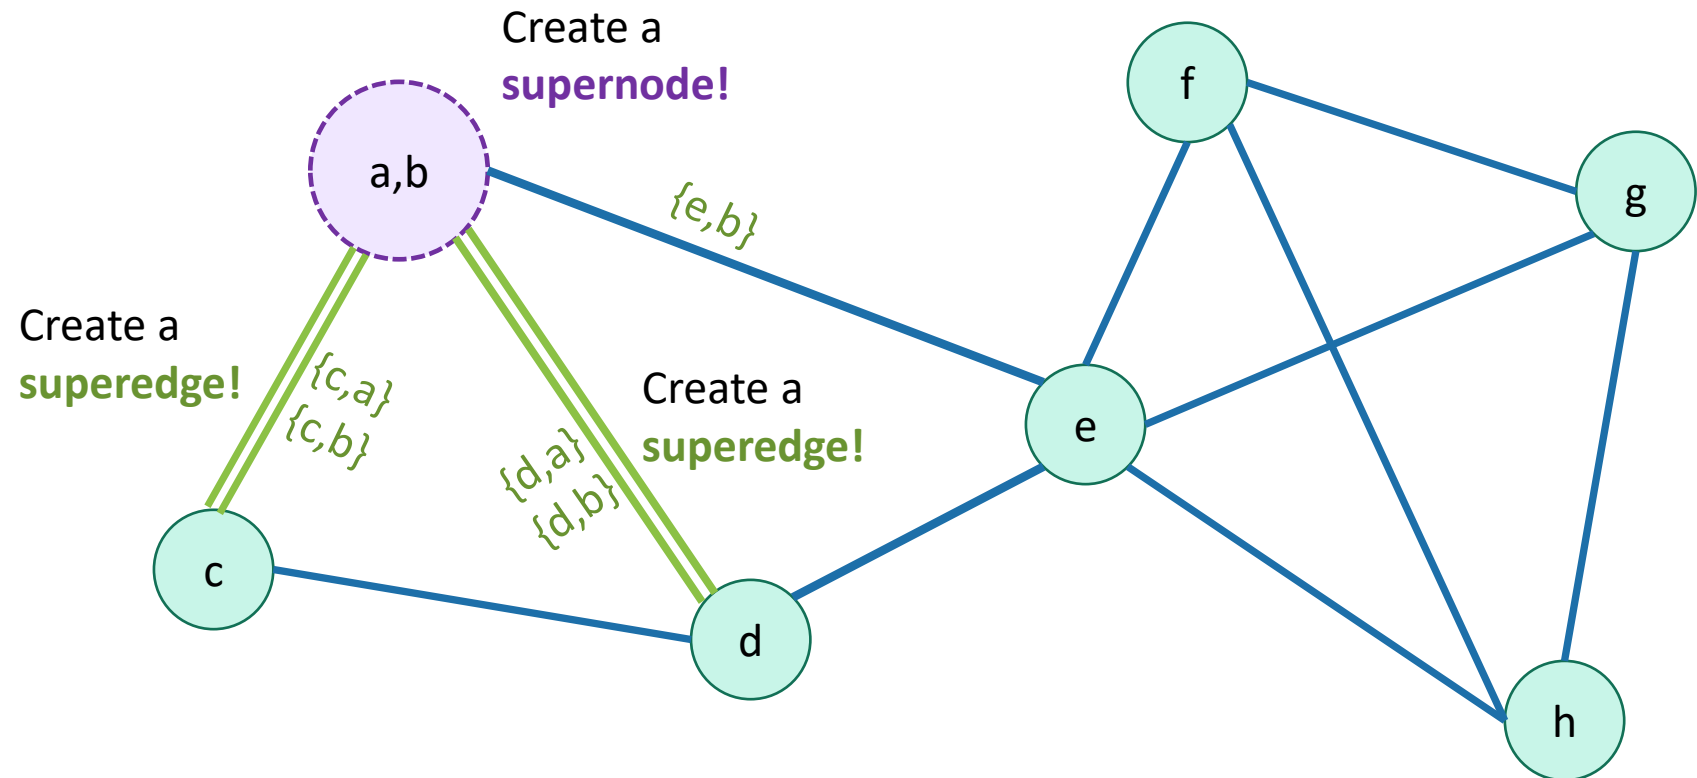
There are 14 edges, 12 of which are good to contract.



Karger's algorithm



Karger's algorithm




Karger's algorithm

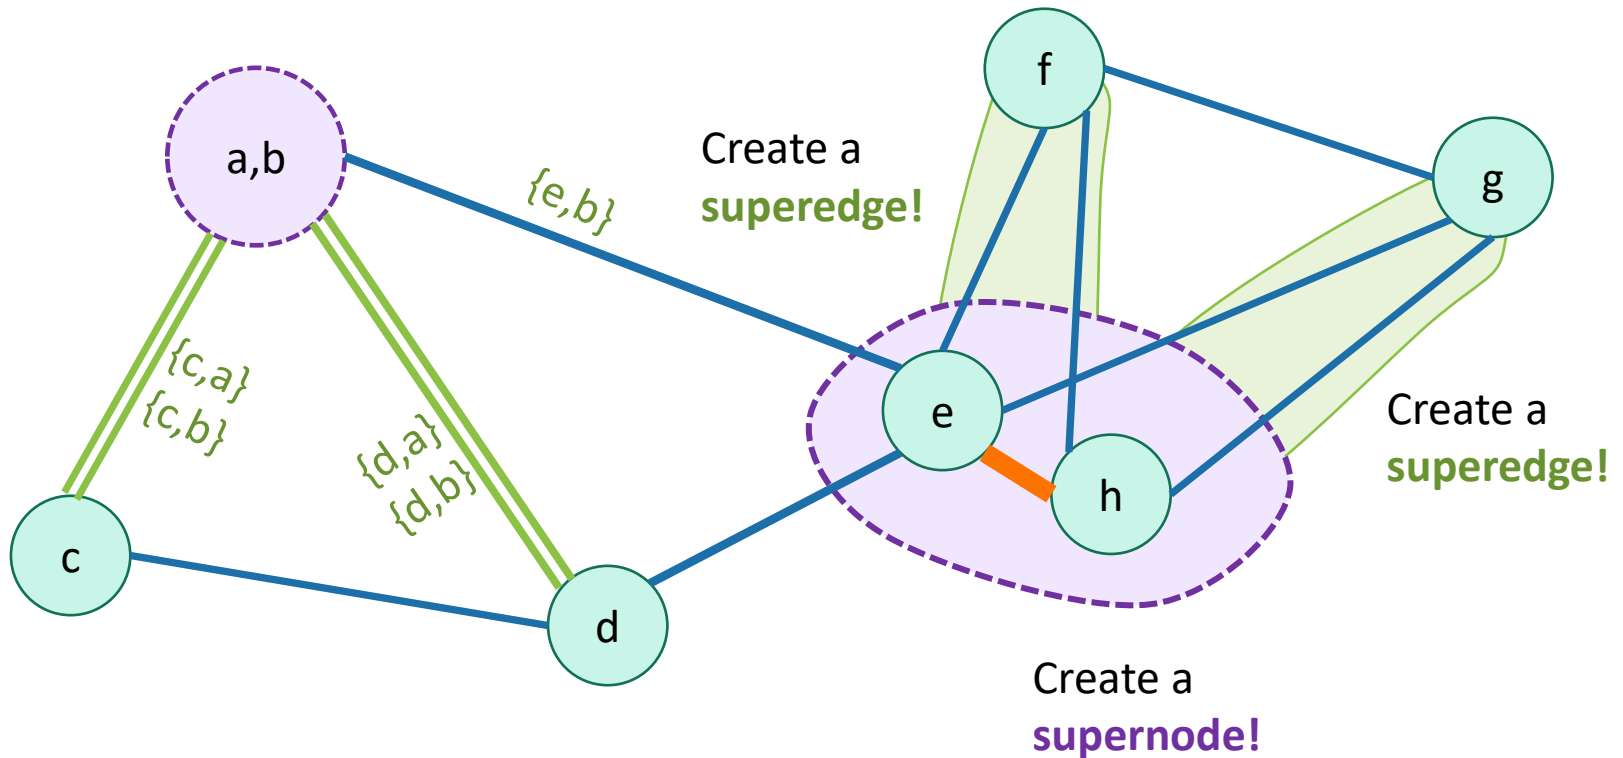
Probability that we didn't mess up:

11/13

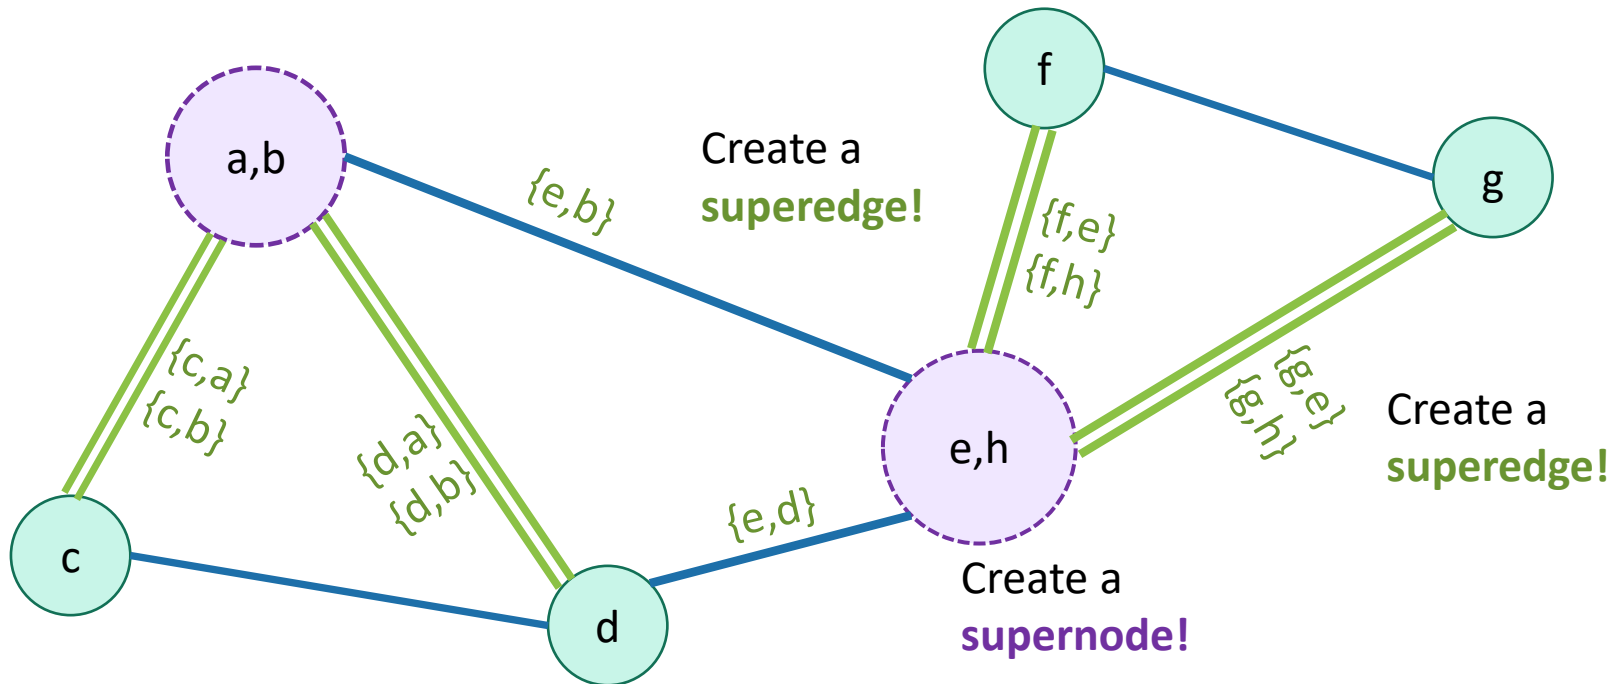
Now there are only 13 edges,
since the edge between a and b
disappeared.



Karger's algorithm




Karger's algorithm




Karger's algorithm

Probability that we didn't mess up:
10/12

Now there are only 12 edges,
since the edge between e and h
disappeared.




Karger's algorithm




Karger's algorithm

Probability that we didn't mess up:
9/11




Karger's algorithm




Karger's algorithm

Probability that we didn't mess up:

$5/7$




Karger's algorithm




Karger's algorithm

Probability that we didn't mess up:

$\frac{3}{5}$




Karger's algorithm



Karger's algorithm


Now stop!

- There are only two nodes left.




Probability of not messing up

- At the beginning, it's pretty likely we'll be fine.
- The probability that we mess up gets worse and worse over time.



Probability of not messing up


- At the beginning, it's pretty likely we'll be fine.
- The probability that we mess up gets worse and worse over time.




Moral:

Repeating the stuff from the beginning of the algorithm is **wasteful!**


Instead...




Instead...




Instead...




Instead...




Contract!




Contract!




Instead...



Contract!




Contract!




FORK!




Instead...





Contract!




Contract!




FORK!




Instead...




Contract!




Contract!





FORK!




Contract!




Contract!




Instead...





Contract!




Contract!




FORK!



Contract!



Contract!




etc


etc




Instead...




Contract!




Contract!




This branch
made a bad
choice.



FORK!




Contract!



Contract!


etc




etc




Instead...



Contract!




Contract!




But it's okay
since this
branch made a
good choice.

This branch
made a bad
choice.

FORK!




Contract!



etc

etc

Contract!



In words

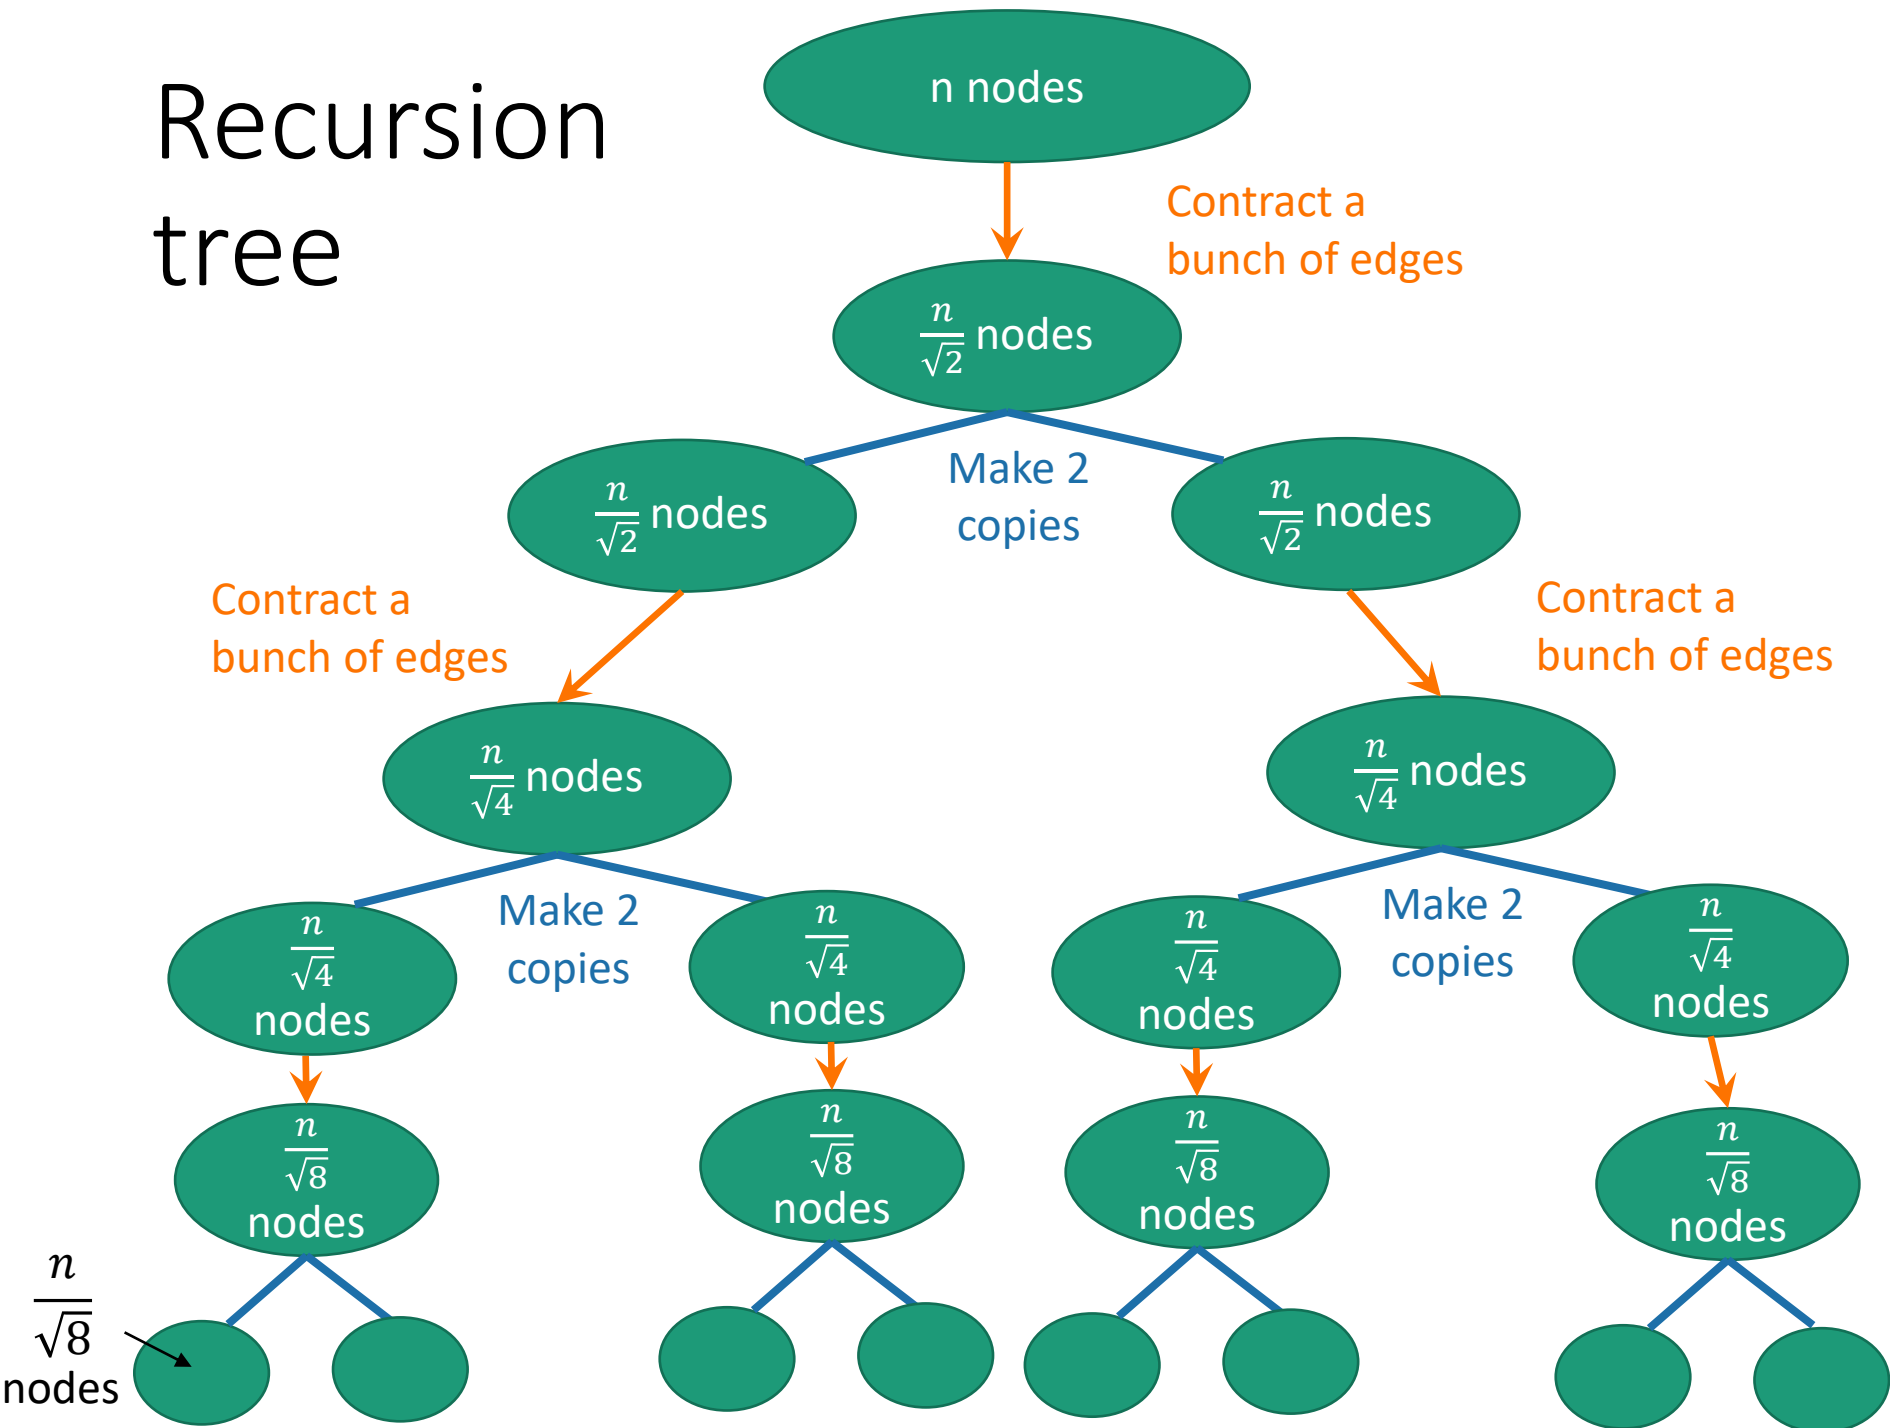
- Run Karger's algorithm on G for a bit.
 - Until there are $\frac{n}{\sqrt{2}}$ supernodes left.
- Then split into two independent copies, G_1 and G_2
- Run Karger's algorithm on each of those for a bit.
 - Until there are $\frac{\left(\frac{n}{\sqrt{2}}\right)}{\sqrt{2}} = \frac{n}{2}$ supernodes left in each.
- Then split each of those into two independent copies...

Why $\frac{n}{\sqrt{2}}$? We'll see later.

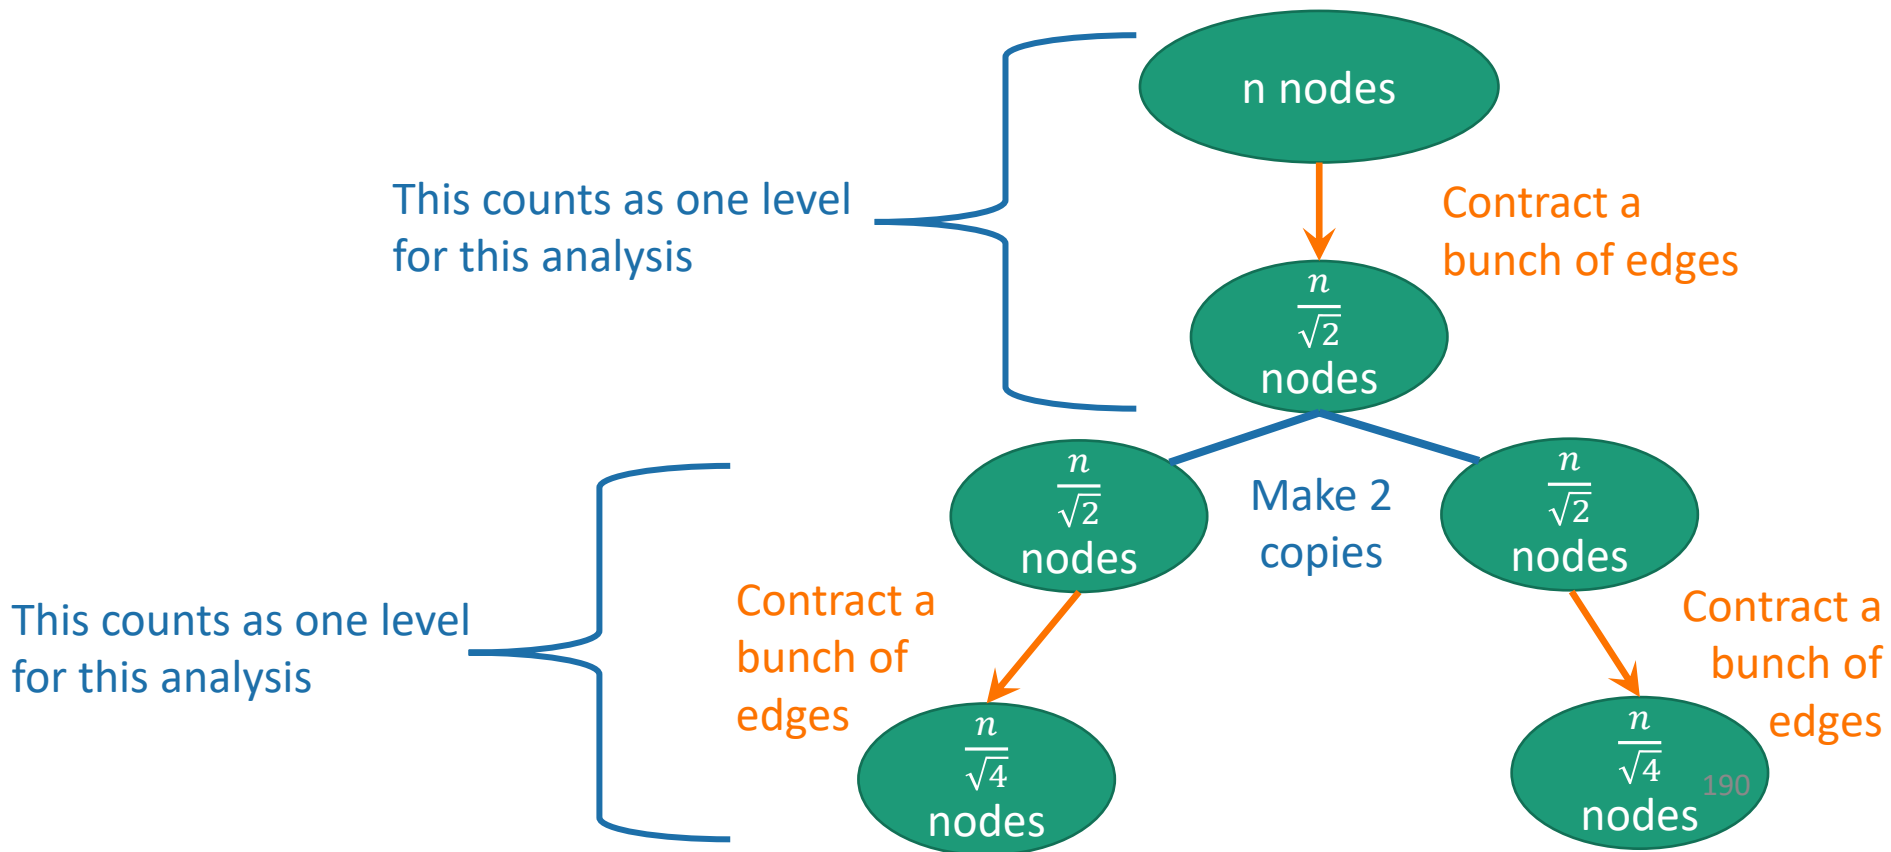
In pseudocode

- **KargerStein**($G = (V, E)$):
 - $n \leftarrow |V|$
 - if $n < 4$:
 - find a min-cut by brute force \\ time $O(1)$
 - Run Karger's algorithm on G with independent repetitions until $\left\lfloor \frac{n}{\sqrt{2}} \right\rfloor$ nodes remain.
 - $G_1, G_2 \leftarrow$ copies of what's left of G
 - $S_1 = \text{KargerStein}(G_1)$
 - $S_2 = \text{KargerStein}(G_2)$
 - **return** whichever of S_1, S_2 is the smaller cut.

Recursion tree

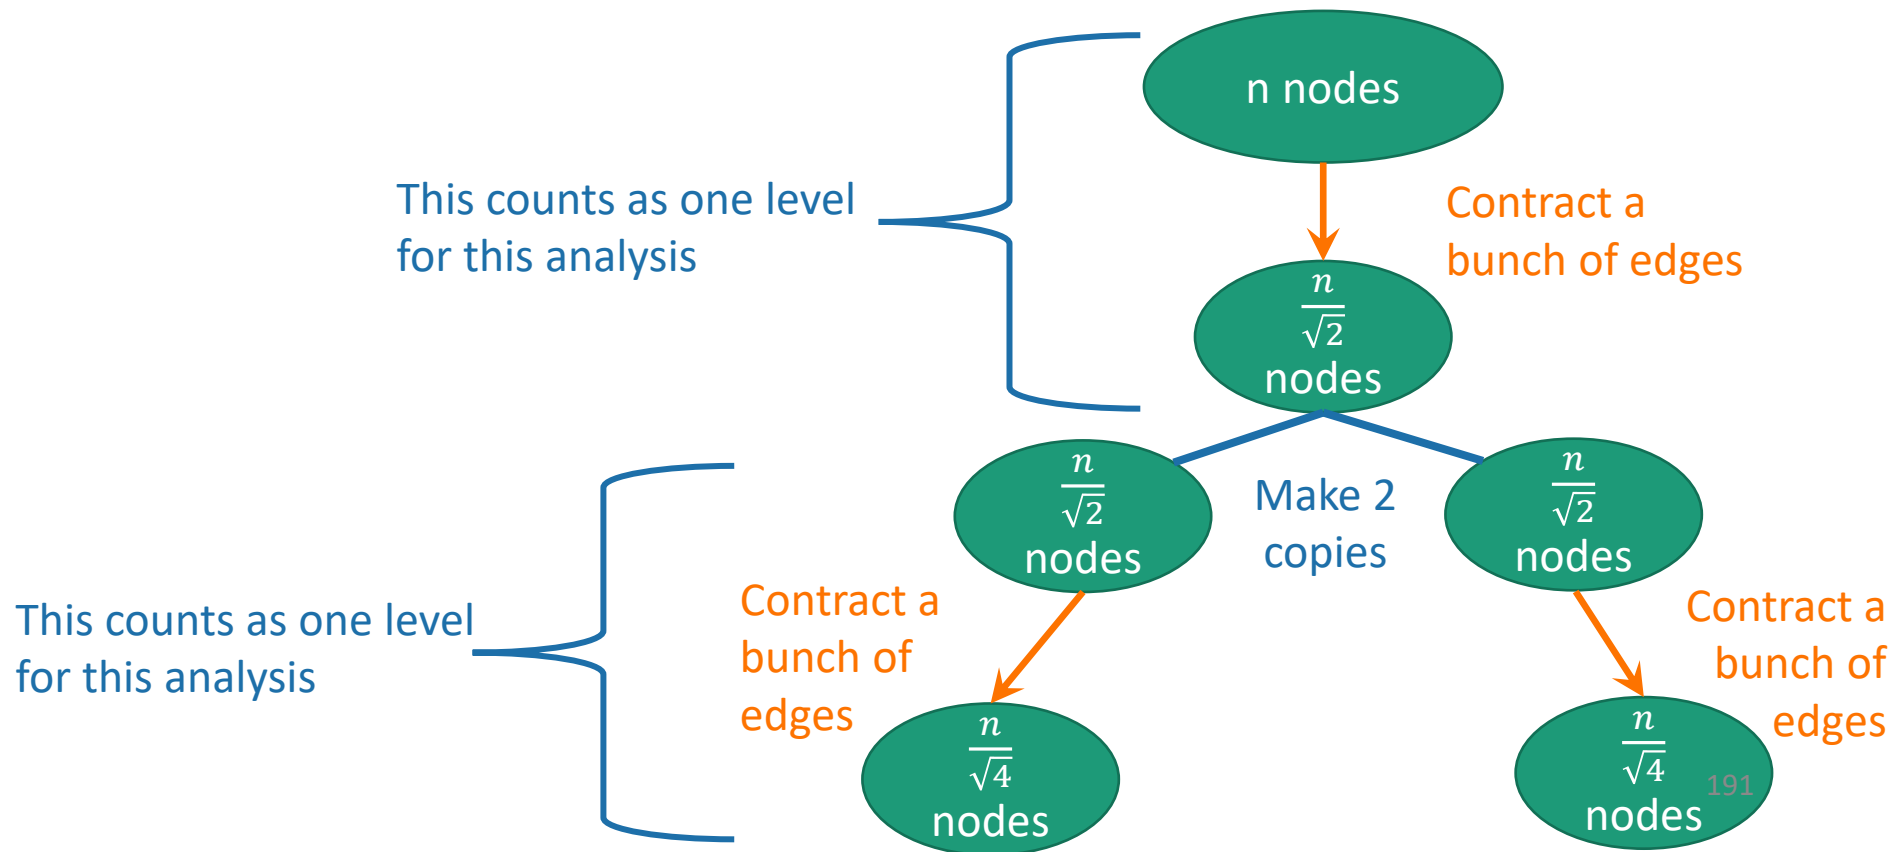


Recursion tree



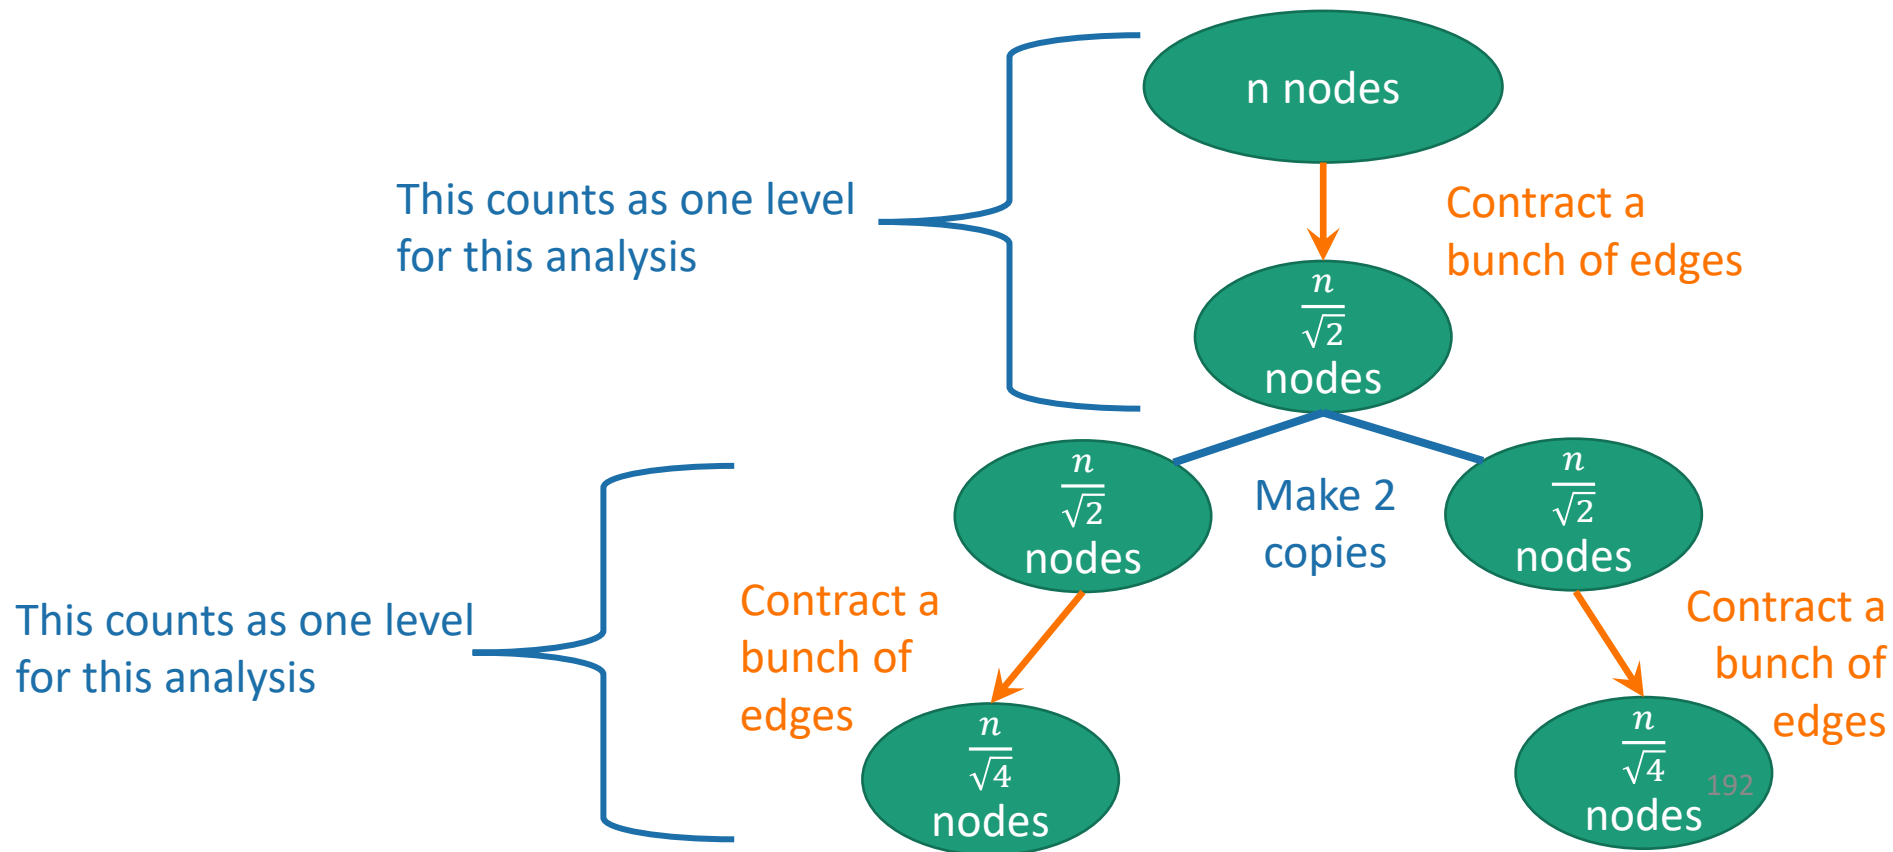
Recursion tree

- depth is $\log_{\sqrt{2}}(n) = \frac{\log(n)}{\log(\sqrt{2})} = 2\log(n)$




Recursion tree


- depth is $\log_{\sqrt{2}}(n) = \frac{\log(n)}{\log(\sqrt{2})} = 2\log(n)$
- number of leaves is $2^{2\log(n)} = n^2$




Two questions

- Does this work?
- Is it fast? 

At the j^{th} level




At the j^{th} level




- The amount of work per level is the amount of work needed to reduce the number of nodes by a factor of $\sqrt{2}$.

At the j^{th} level



- The amount of work per level is the amount of work needed to reduce the number of nodes by a factor of $\sqrt{2}$.
- That's at most $O(n^2)$.
 - since that's the time it takes to run Karger's algorithm once, cutting down the number of supernodes to **two**.


At the j^{th} level



- The amount of work per level is the amount of work needed to reduce the number of nodes by a factor of $\sqrt{2}$.
- That's at most $O(n^2)$.
 - since that's the time it takes to run Karger's algorithm once, cutting down the number of supernodes to **two**.
- Our recurrence relation is...

$$T(n) = 2T(n/\sqrt{2}) + O(n^2)$$

At the j^{th} level



- The amount of work per level is the amount of work needed to reduce the number of nodes by a factor of $\sqrt{2}$.
- That's at most $O(n^2)$.
 - since that's the time it takes to run Karger's algorithm once, cutting down the number of supernodes to **two**.
- Our recurrence relation is...

$$T(n) = 2T(n/\sqrt{2}) + O(n^2)$$

- The Master Theorem says...

$$T(n) = O(n^2 \log(n))$$



Two questions

- Does this work?



- Is it fast?
 - Yes, $O(n^2 \log(n))$.

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

e_1, e_2, \dots, e_{n-t}

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

$$e_1, e_2, \dots, e_{n-t}$$

- $\text{PR}[\text{none of } e_1, e_2, \dots, e_{n-t} \text{ cross } S^*]$

$$= \text{PR}[e_1 \text{ doesn't cross } S^*]$$

$$\times \text{PR}[e_2 \text{ doesn't cross } S^* \mid e_1 \text{ doesn't cross } S^*]$$

...

$$\times \text{PR}[e_{n-t} \text{ doesn't cross } S^* \mid e_1, \dots, e_{n-t-1} \text{ don't cross } S^*]$$

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

$$e_1, e_2, \dots, e_{n-t}$$

- **PR**[none of e_1, e_2, \dots, e_{n-t} cross S^*]

$$= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{t+1}{t+3}\right) \left(\frac{t}{t+2}\right) \left(\frac{t-1}{t+1}\right)$$

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

$$e_1, e_2, \dots, e_{n-t}$$

- **PR**[none of e_1, e_2, \dots, e_{n-t} cross S^*]

$$\begin{aligned} &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{t+1}{t+3}\right) \left(\frac{t}{t+2}\right) \left(\frac{t-1}{t+1}\right) \\ &= \frac{t \cdot (t-1)}{n \cdot (n-1)} \end{aligned}$$

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

$$e_1, e_2, \dots, e_{n-t}$$

- **PR**[none of e_1, e_2, \dots, e_{n-t} cross S^*]

$$\begin{aligned} &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{t+1}{t+3}\right) \left(\frac{t}{t+2}\right) \left(\frac{t-1}{t+1}\right) \\ &= \frac{t \cdot (t-1)}{n \cdot (n-1)} \end{aligned}$$

Choose $t = n/\sqrt{2}$

Why $n/\sqrt{2}$?

Suppose we contract $n - t$ edges, until there are t supernodes remaining.

- Suppose the first $n-t$ edges that we choose are

$$e_1, e_2, \dots, e_{n-t}$$

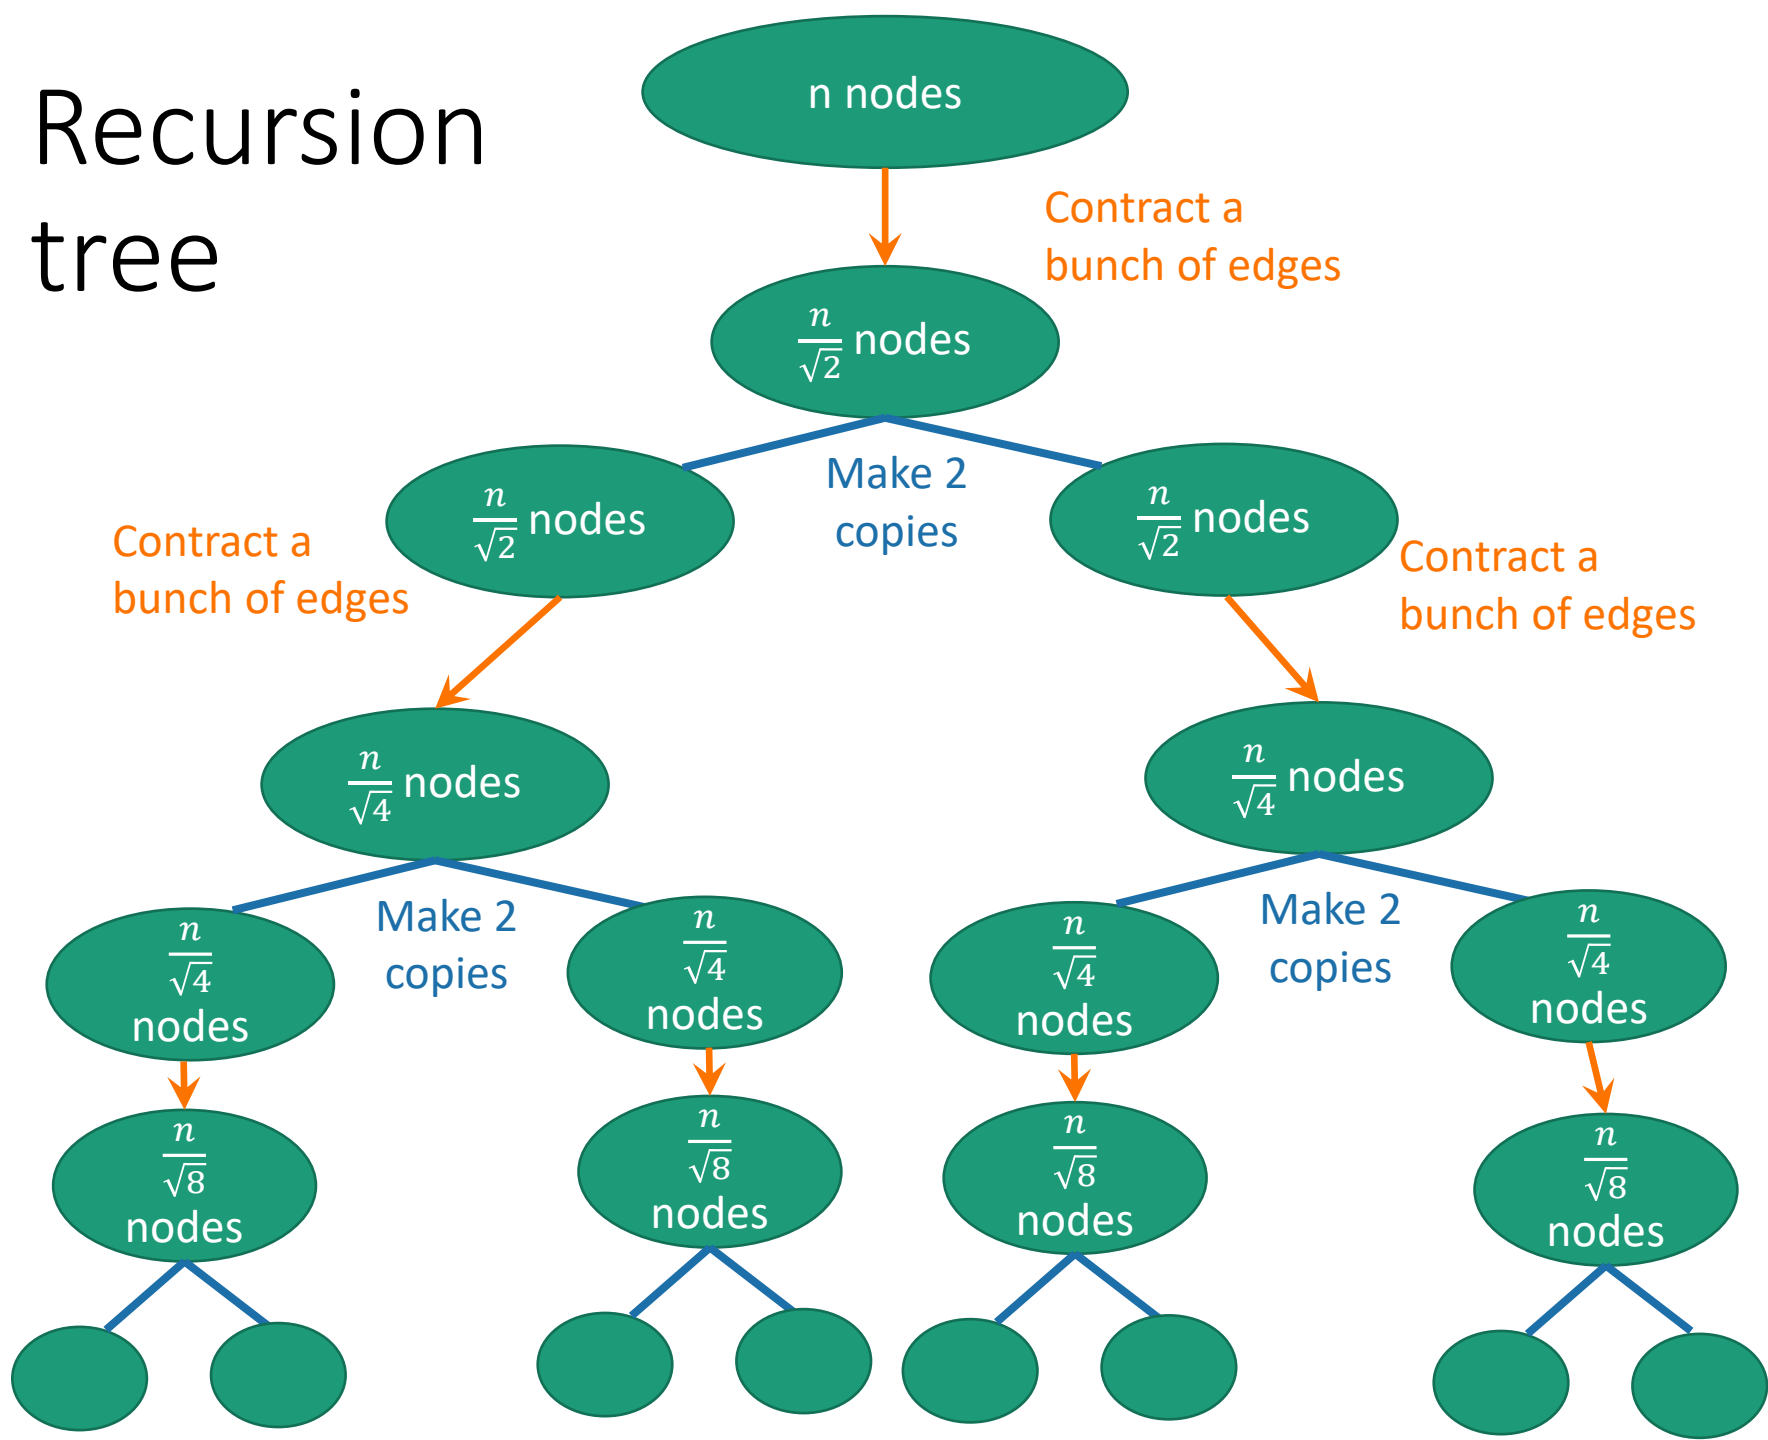
- **PR**[none of e_1, e_2, \dots, e_{n-t} cross S^*]

$$= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \left(\frac{n-6}{n-4}\right) \cdots \left(\frac{t+1}{t+3}\right) \left(\frac{t}{t+2}\right) \left(\frac{t-1}{t+1}\right)$$

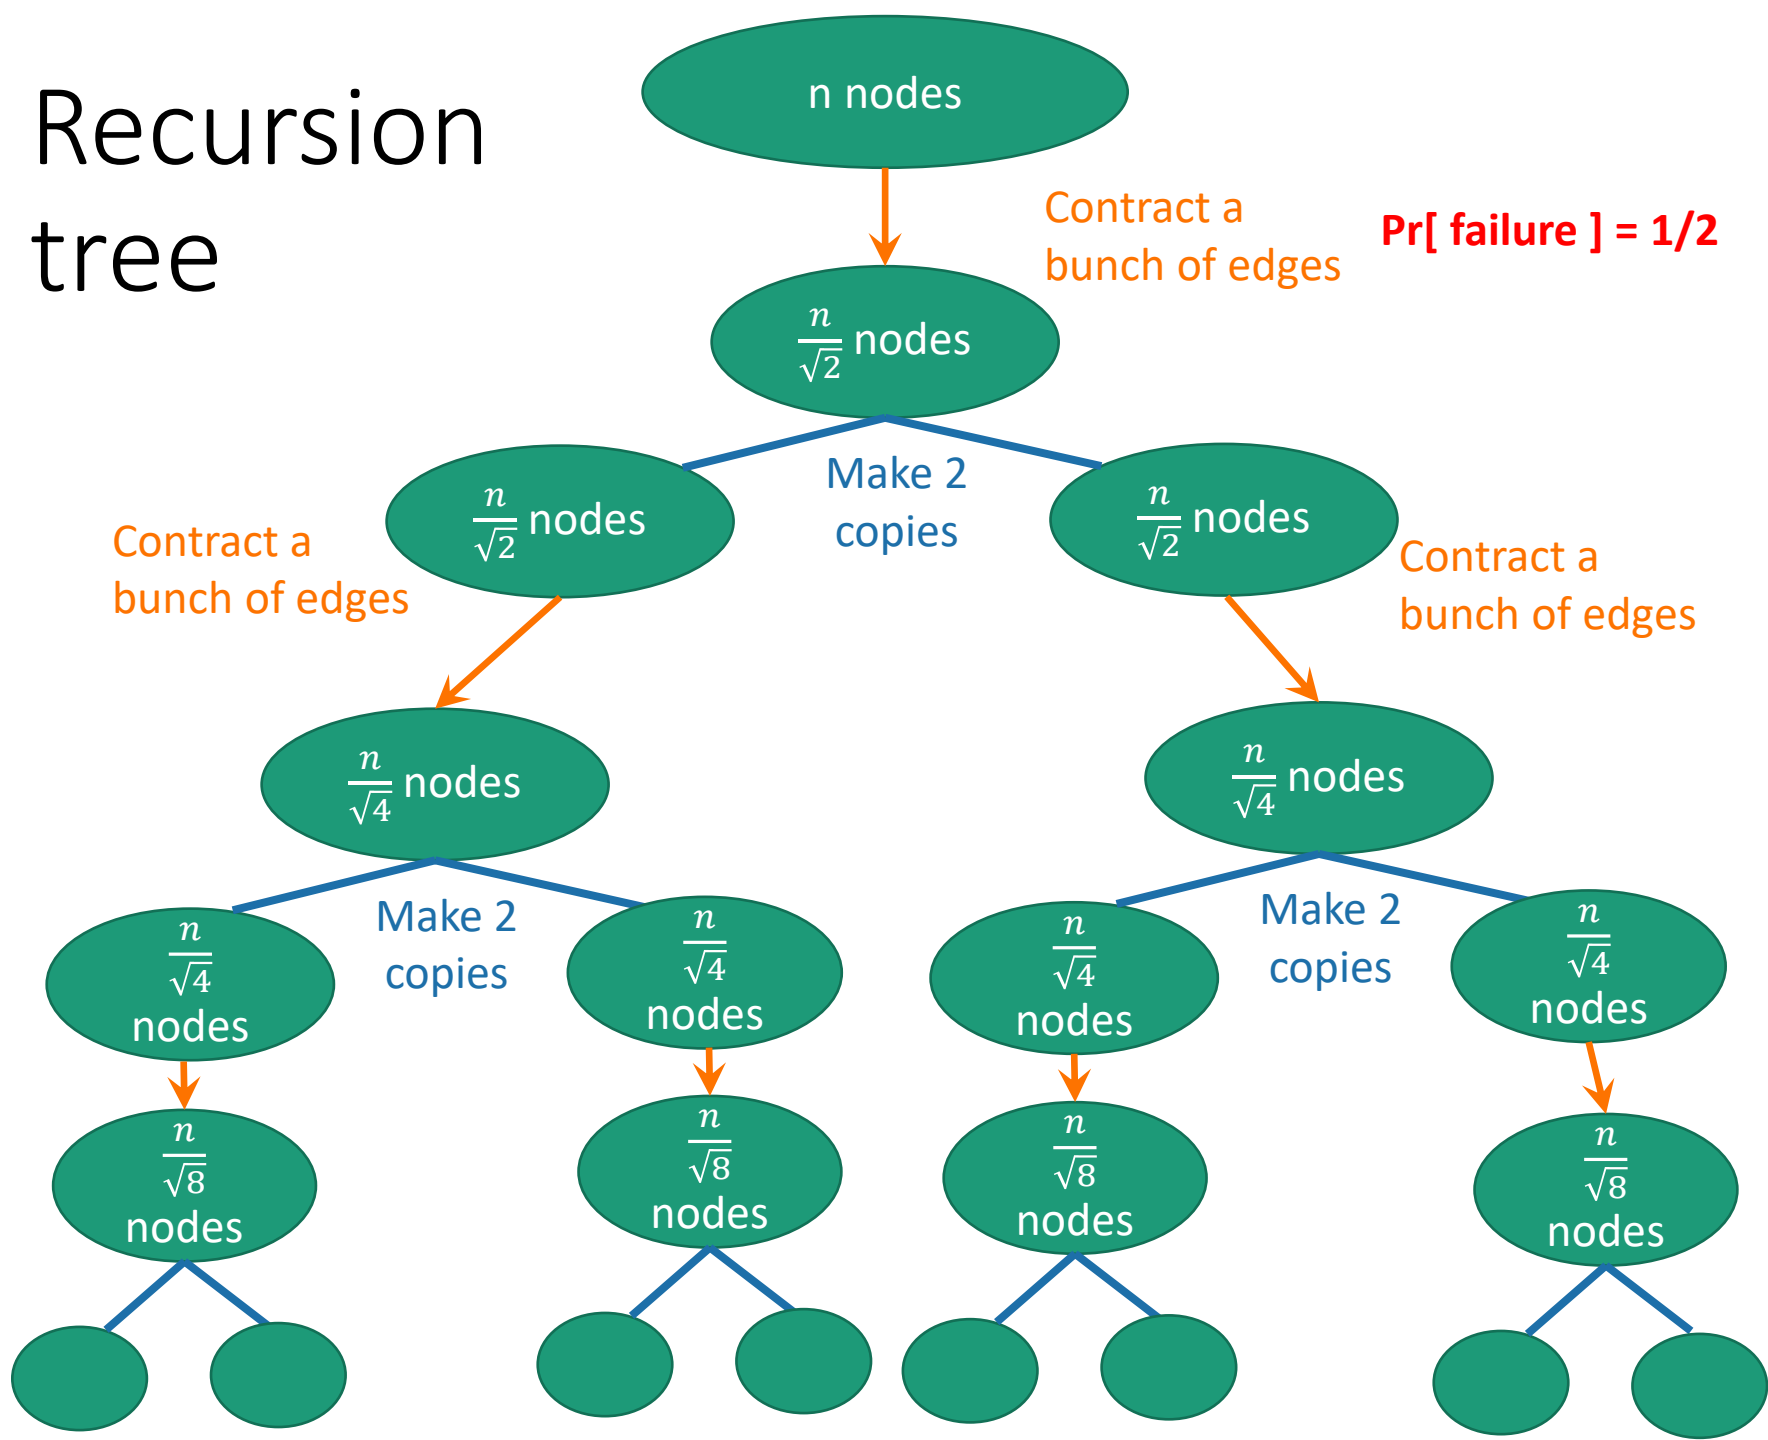
$$= \frac{t \cdot (t-1)}{n \cdot (n-1)} \quad \text{Choose } t = n/\sqrt{2}$$

$$= \frac{\frac{n}{\sqrt{2}} \cdot \left(\frac{n}{\sqrt{2}} - 1\right)}{n \cdot (n-1)} \approx \frac{1}{2} \quad \text{when } n \text{ is large}$$

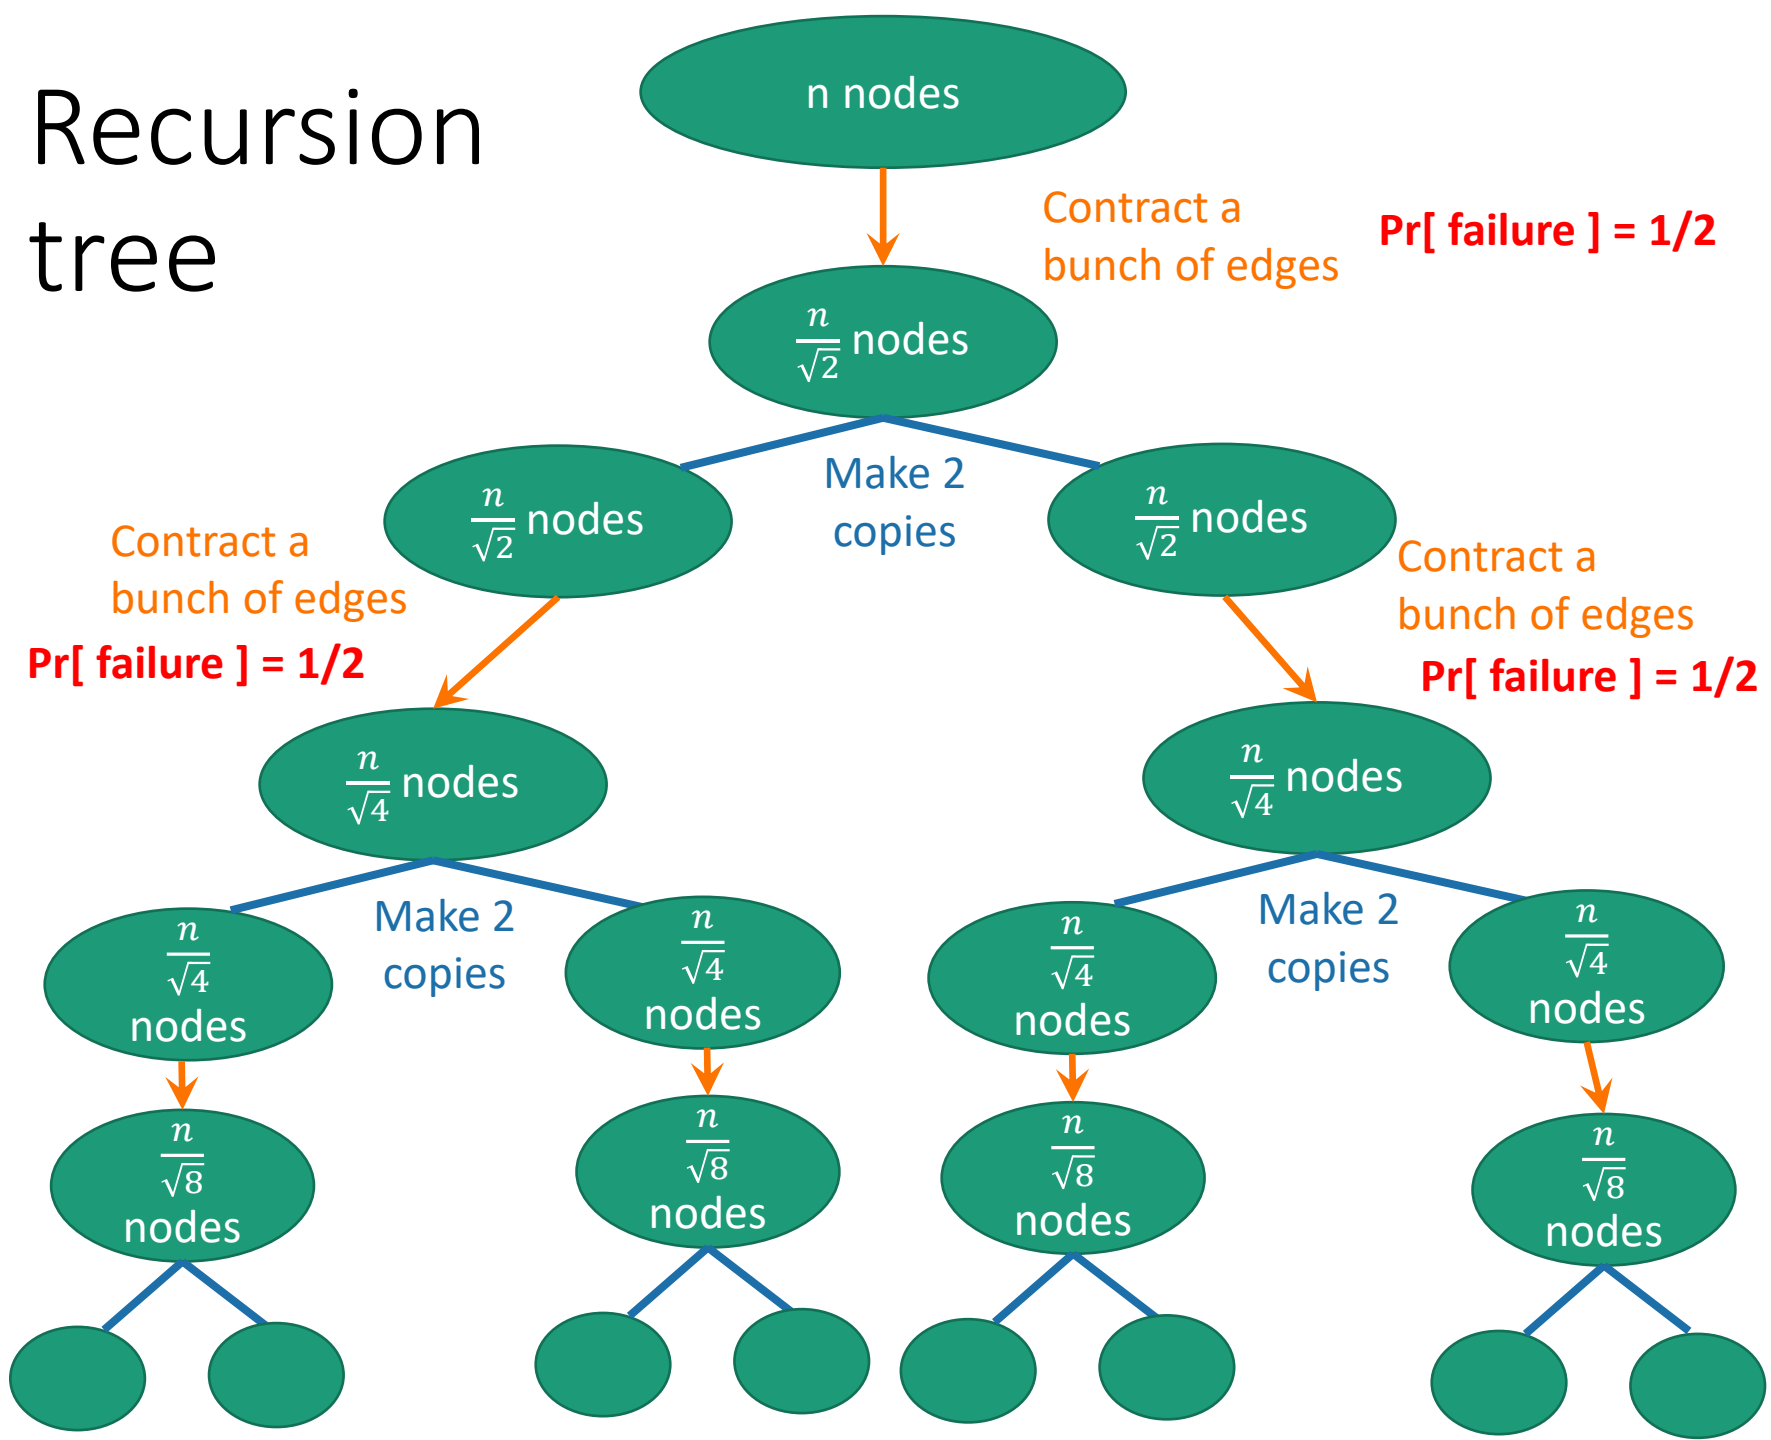
Recursion tree



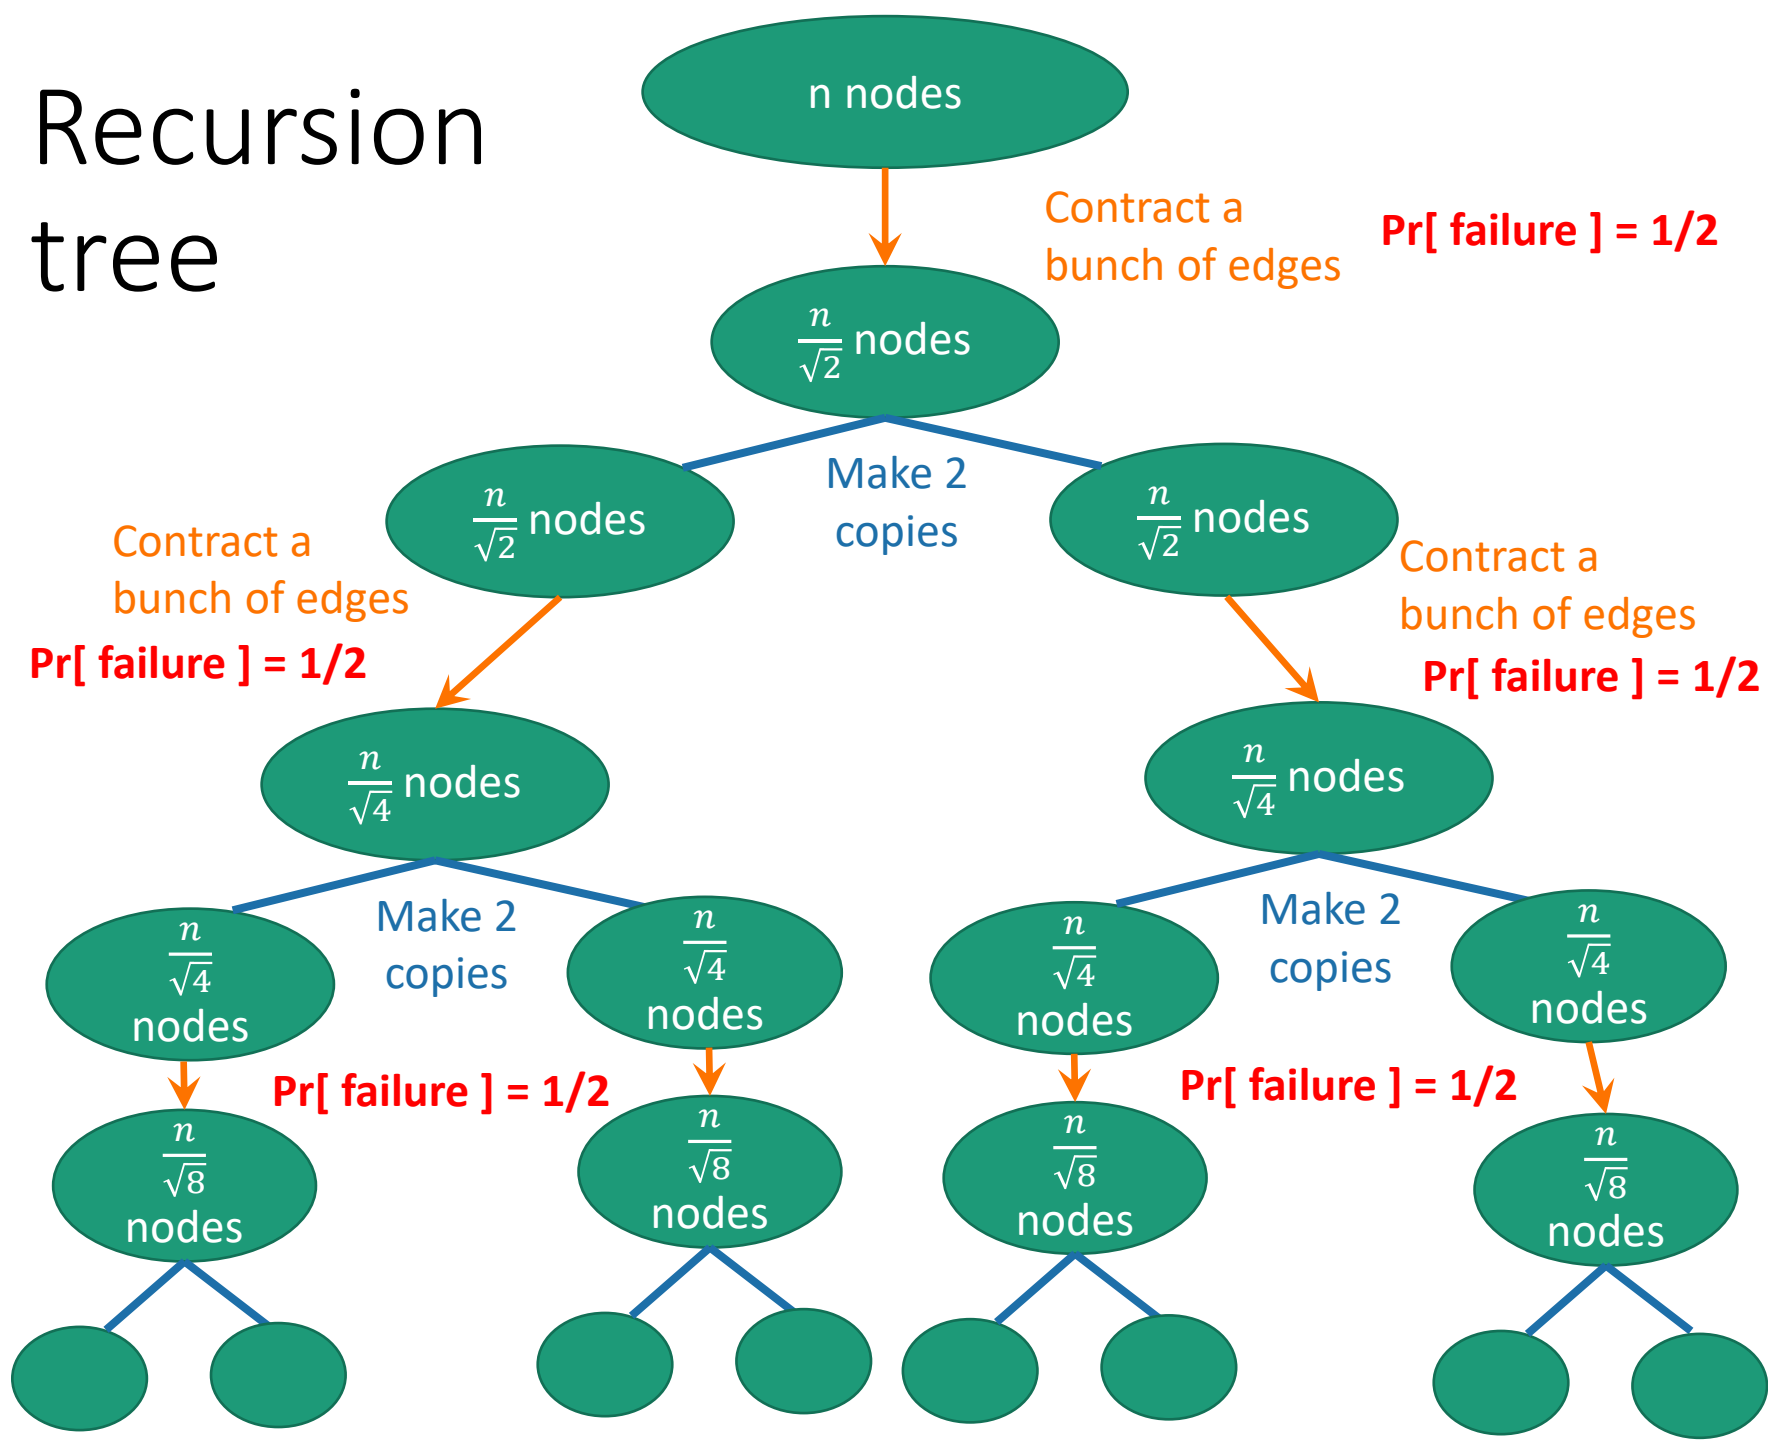
Recursion tree



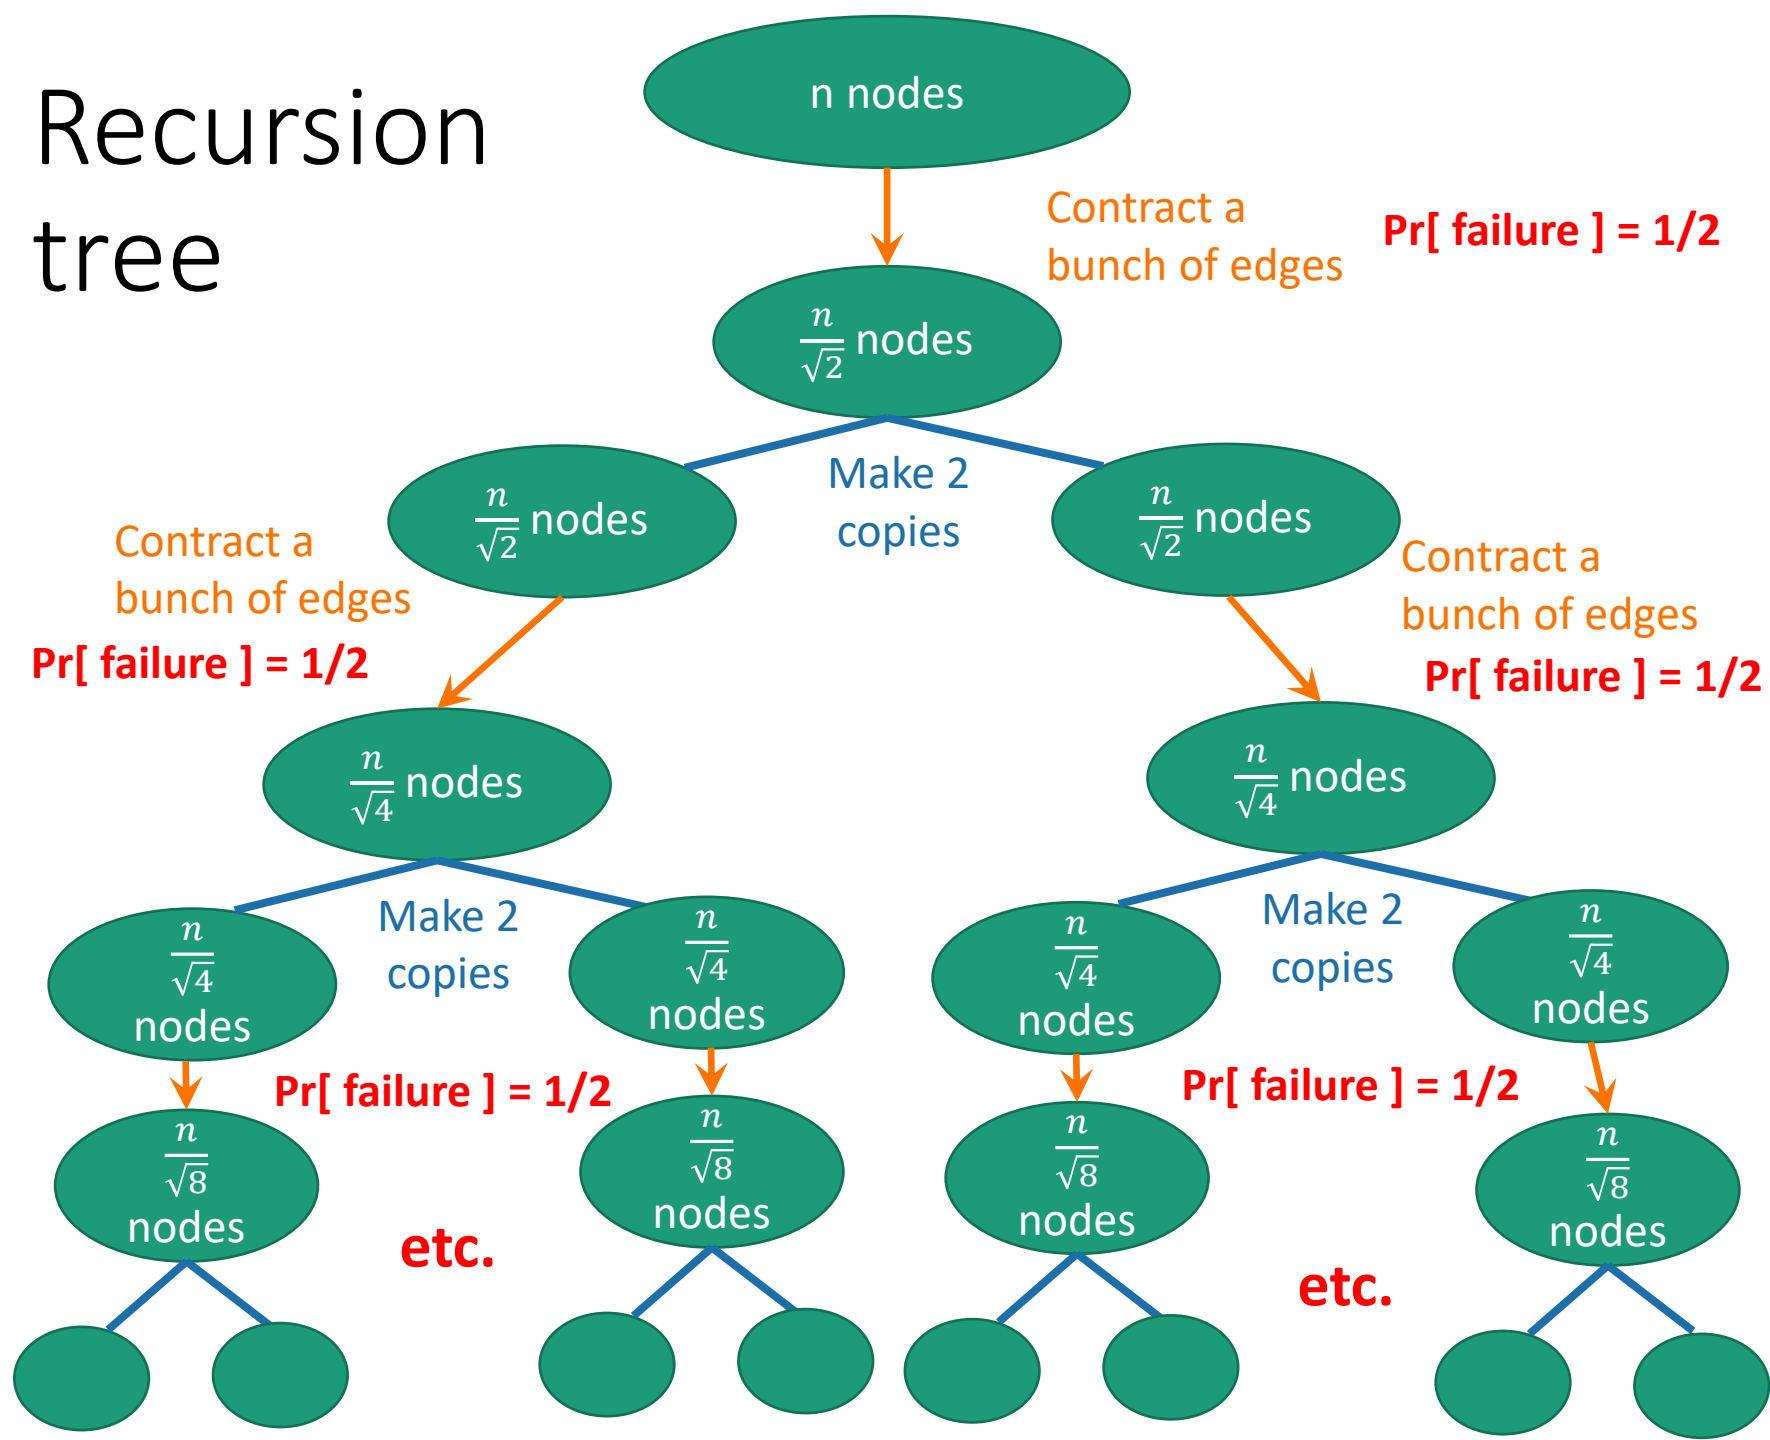
Recursion tree



Recursion tree

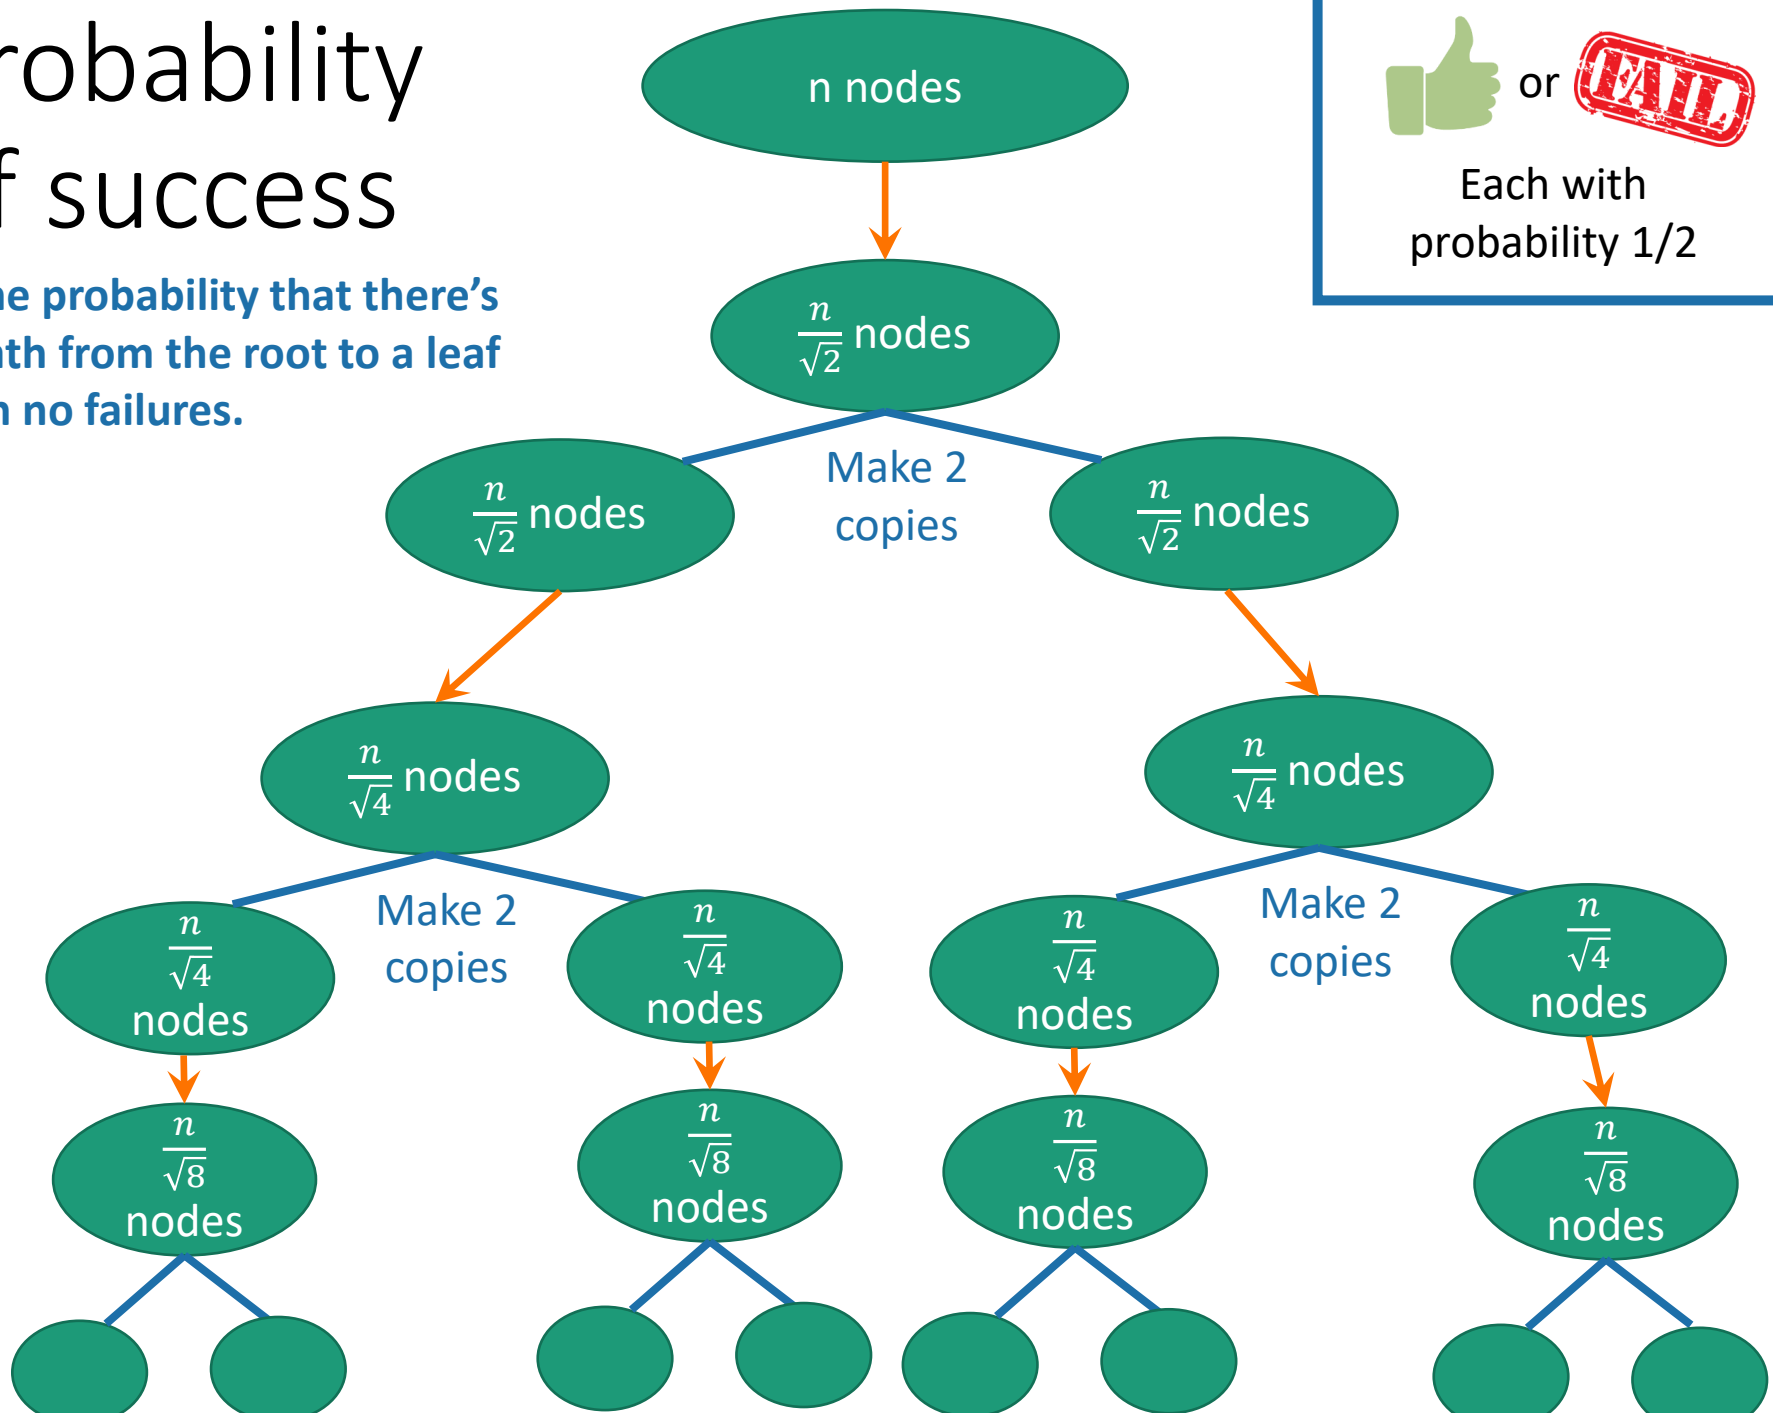


Recursion tree



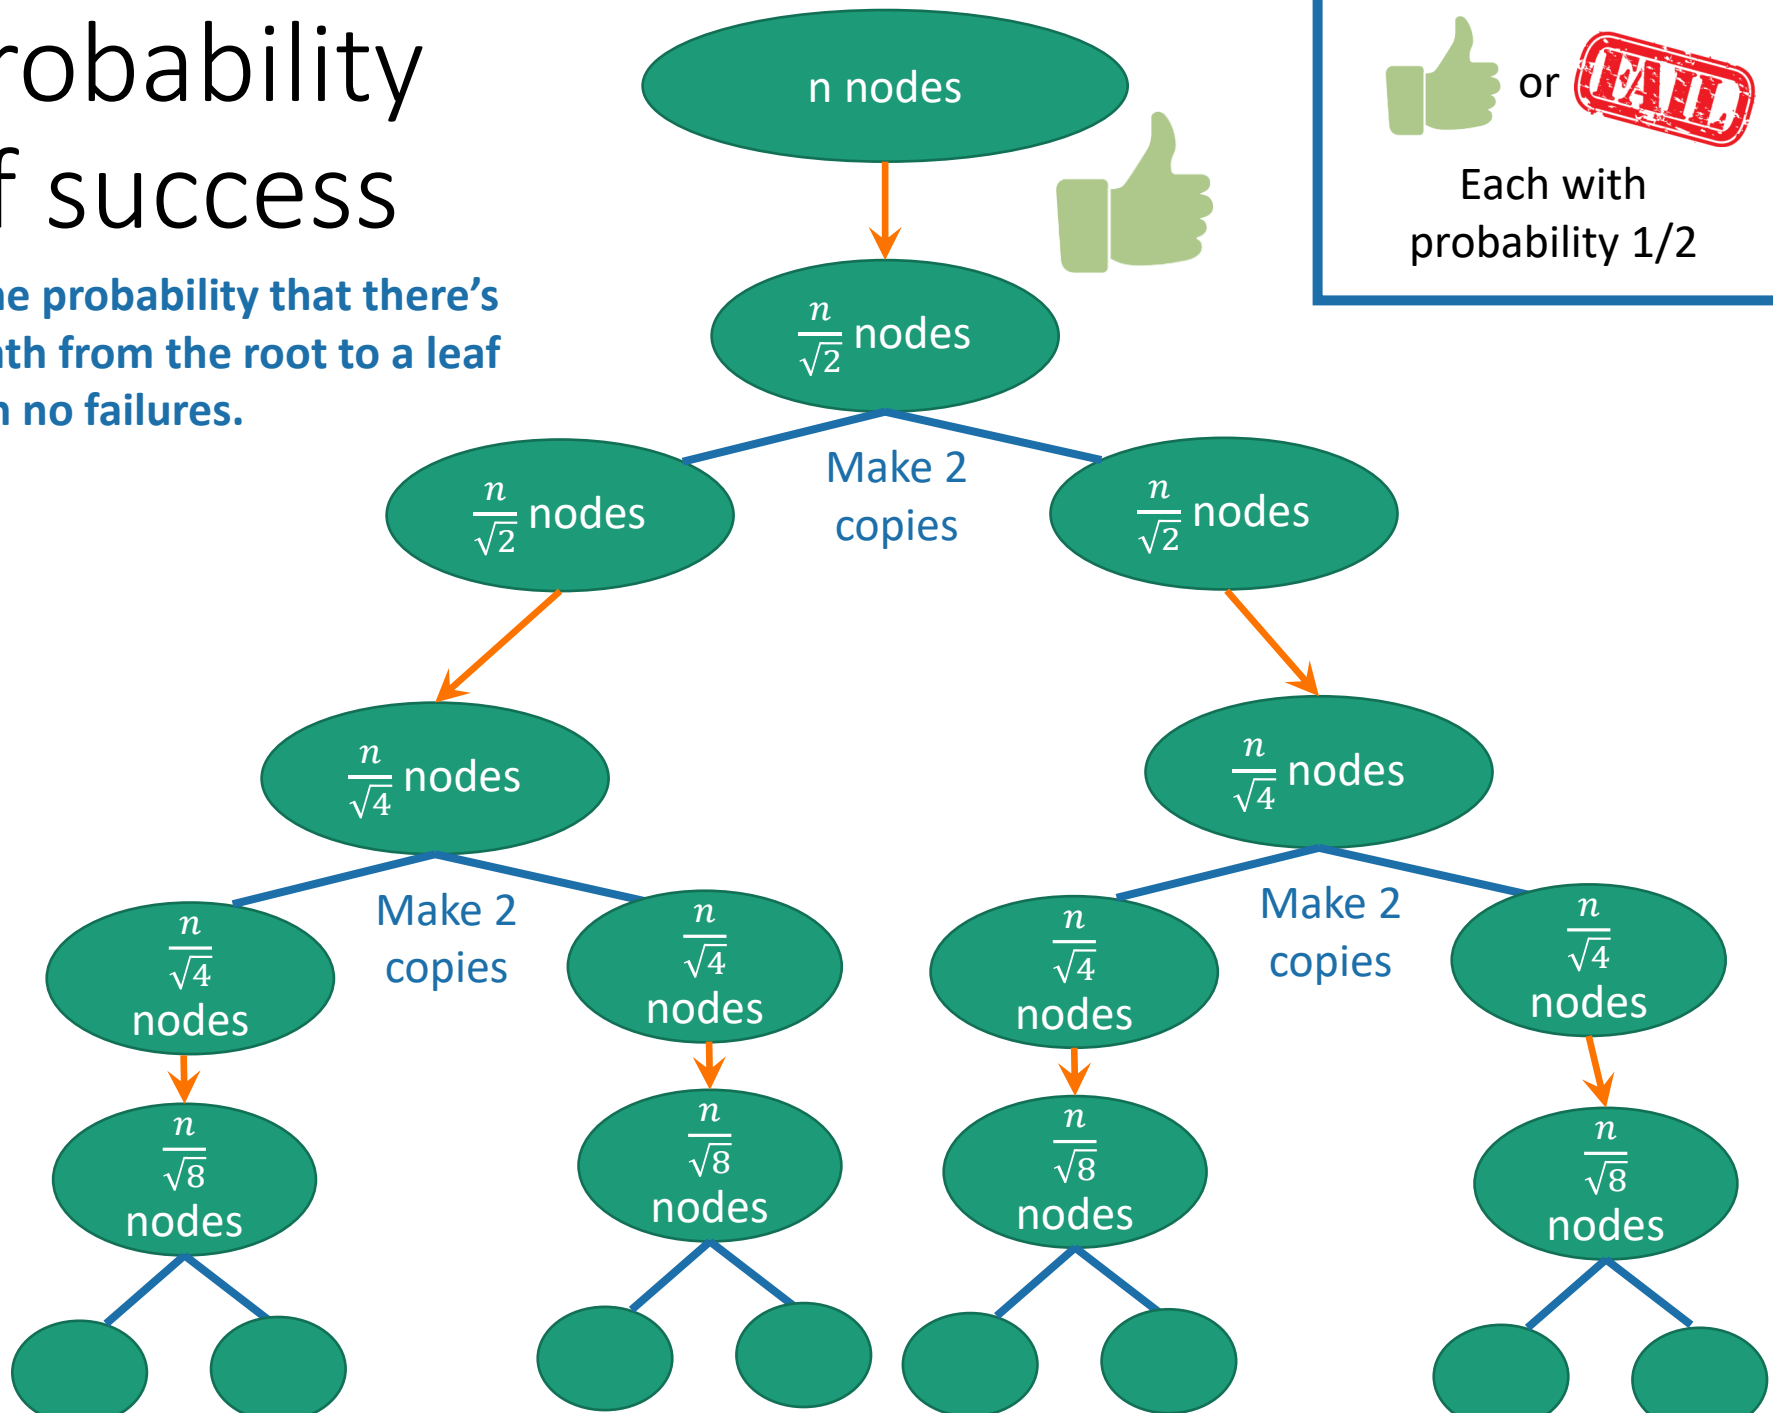
Probability of success

Is the probability that there's a path from the root to a leaf with no failures.



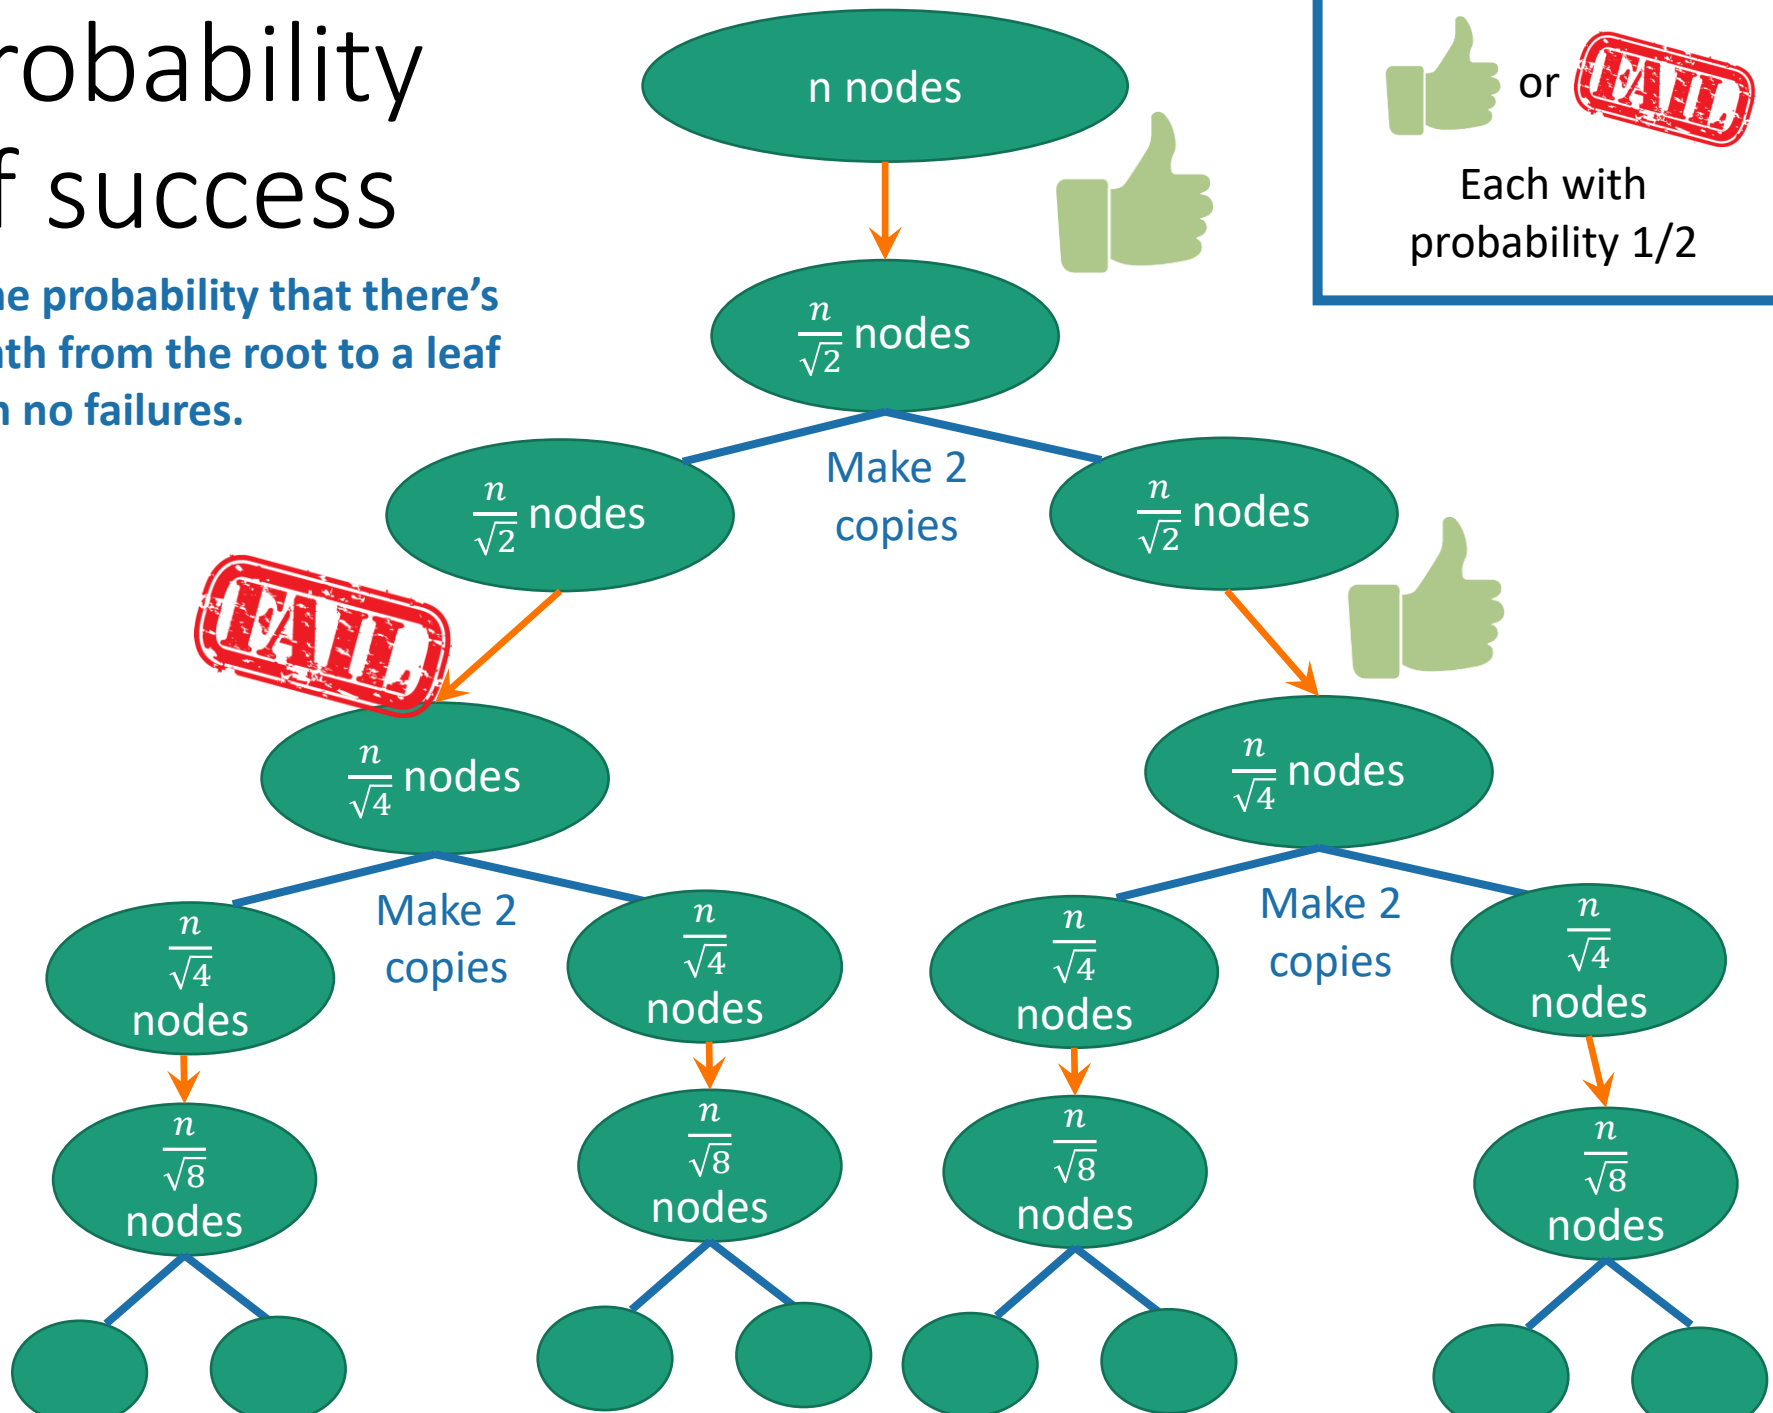
Probability of success

Is the probability that there's a path from the root to a leaf with no failures.



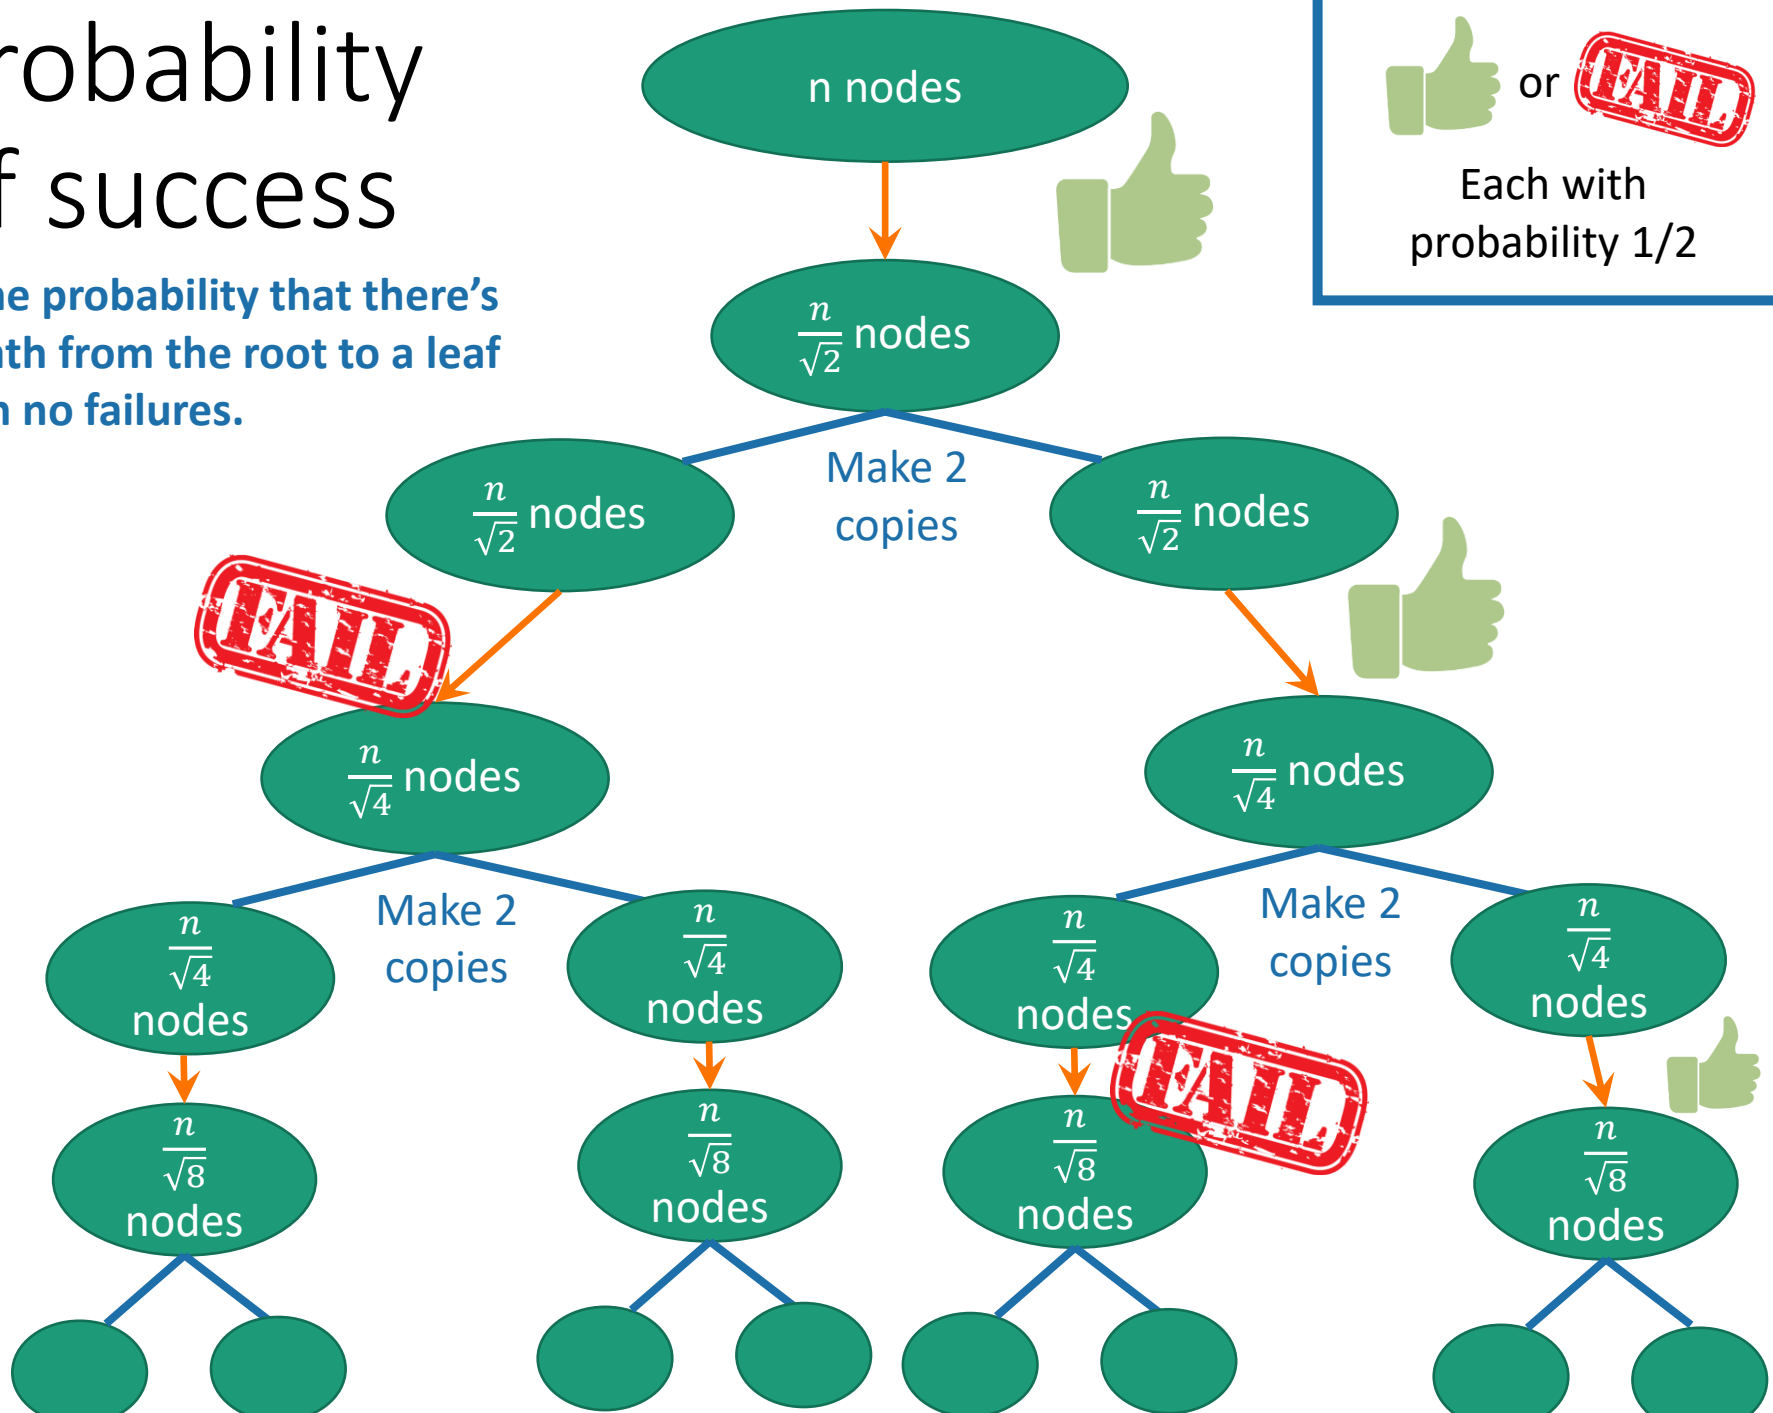
Probability of success

Is the probability that there's a path from the root to a leaf with no failures.



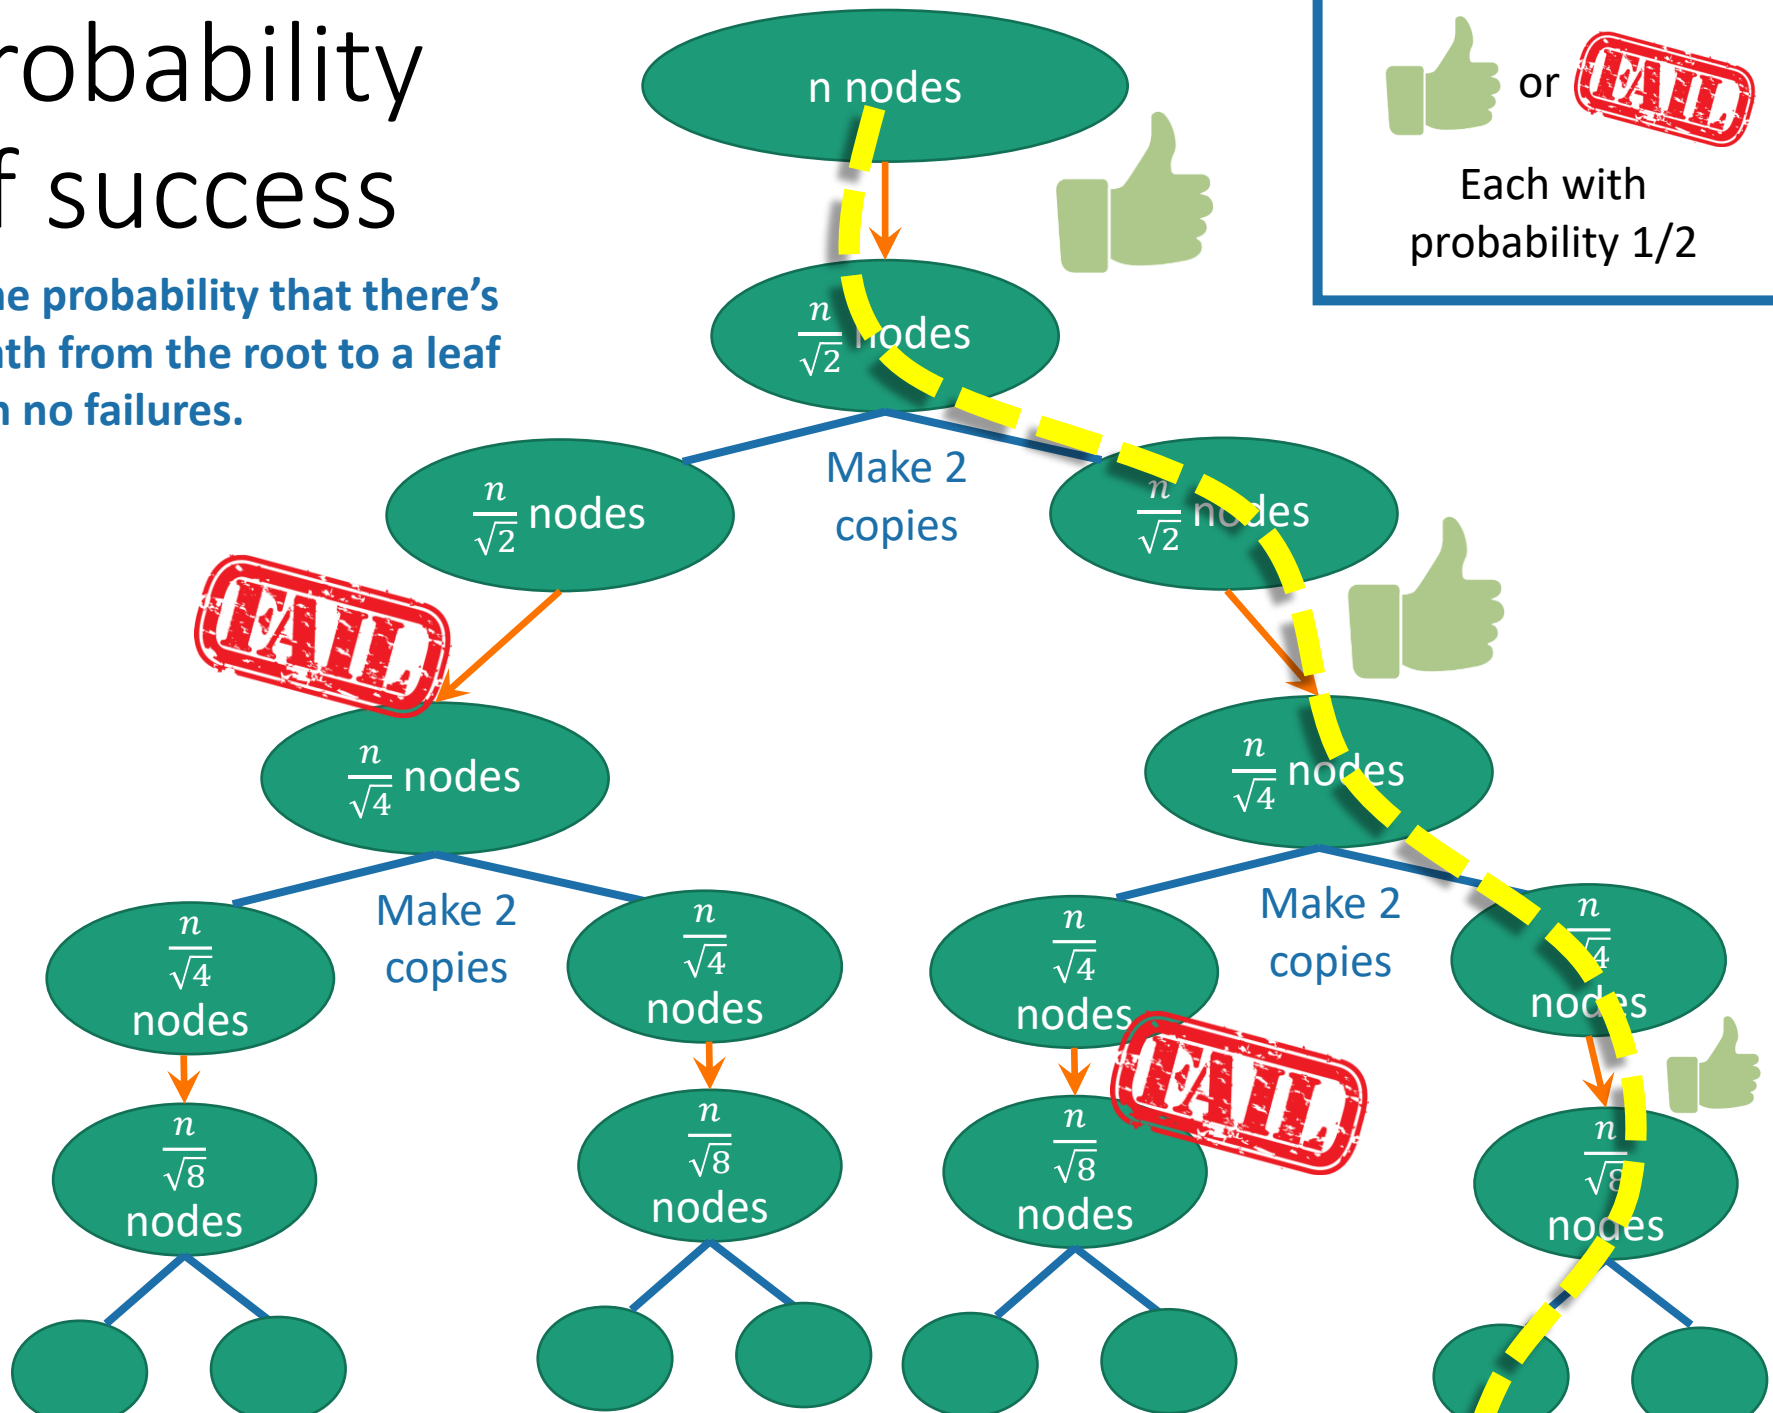
Probability of success

Is the probability that there's a path from the root to a leaf with no failures.



Probability of success


Is the probability that there's a path from the root to a leaf with no failures.




The problem we need to analyze

- Let T be binary tree of depth $2\log(n)$
- Each node of T succeeds or fails independently with probability $1/2$
- What is the probability that there's a path from the root to any leaf that's entirely successful?

Analysis




Analysis



Analysis


- Say the tree has height d .
- Let p_d be the probability that there's a path from the root to a leaf that **doesn't fail**.



Analysis

- Say the tree has height d .
- Let p_d be the probability that there's a path from the root to a leaf that **doesn't fail**.

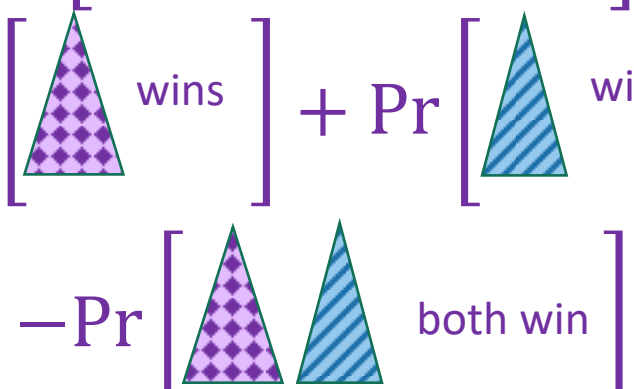
$$p_d = \frac{1}{2} \cdot \Pr \left[\begin{array}{l} \text{at least one subtree} \\ \text{has a successful path} \end{array} \right]$$




Analysis

- Say the tree has height d .
- Let p_d be the probability that there's a path from the root to a leaf that **doesn't fail**.

$$p_d = \frac{1}{2} \cdot \Pr \left[\begin{array}{l} \text{at least one subtree} \\ \text{has a successful path} \end{array} \right]$$


$$= \frac{1}{2} \cdot \left(\Pr \left[\begin{array}{l} \text{wins} \end{array} \right] + \Pr \left[\begin{array}{l} \text{wins} \end{array} \right] - \Pr \left[\begin{array}{l} \text{both win} \end{array} \right] \right)$$




Analysis

- Say the tree has height d .
- Let p_d be the probability that there's a path from the root to a leaf that **doesn't fail**.


$$\begin{aligned}
 p_d &= \frac{1}{2} \cdot \Pr \left[\begin{array}{c} \text{at least one subtree} \\ \text{has a successful path} \end{array} \right] \\
 &= \frac{1}{2} \cdot \left(\Pr \left[\begin{array}{c} \text{wins} \end{array} \right] + \Pr \left[\begin{array}{c} \text{wins} \end{array} \right] - \Pr \left[\begin{array}{c} \text{both win} \end{array} \right] \right) \\
 &= \frac{1}{2} \cdot (p_{d-1} + p_{d-1} - p_{d-1}^2)
 \end{aligned}$$



Analysis

- Say the tree has height d .
- Let p_d be the probability that there's a path from the root to a leaf that **doesn't fail**.

$$\begin{aligned}
 p_d &= \frac{1}{2} \cdot \Pr \left[\begin{array}{c} \text{at least one subtree} \\ \text{has a successful path} \end{array} \right] \\
 &= \frac{1}{2} \cdot \left(\Pr \left[\begin{array}{c} \text{wins} \end{array} \right] + \Pr \left[\begin{array}{c} \text{wins} \end{array} \right] - \Pr \left[\begin{array}{c} \text{both win} \end{array} \right] \right) \\
 &= \frac{1}{2} \cdot (p_{d-1} + p_{d-1} - p_{d-1}^2) \\
 &= p_{d-1} - \frac{1}{2} \cdot p_{d-1}^2
 \end{aligned}$$



It's a recurrence relation!

- $p_d = p_{d-1} - \frac{1}{2} \cdot p_{d-1}^2$
- $p_0 = 1$
- We are really good at those.

It's a recurrence relation!

- $p_d = p_{d-1} - \frac{1}{2} \cdot p_{d-1}^2$
- $p_0 = 1$
- We are really good at those.
- In this case, the answer is:
 - **Claim:** for all d , $p_d \geq \frac{1}{d+1}$

Recurrence relation

- $p_d = p_{d-1} - \frac{1}{2} \cdot p_{d-1}^2$
- $p_0 = 1$

- **Claim:** for all d , $p_d \geq \frac{1}{d+1}$
- **Proof:** induction on d .
 - **Base case:** $1 \geq 1$. **YEP.**
 - **Inductive step:** say $d > 0$.
 - Suppose that $p_{d-1} \geq \frac{1}{d}$.
 - $p_d = p_{d-1} - \frac{1}{2} \cdot p_{d-1}^2$
 - $\geq \frac{1}{d} - \frac{1}{2} \cdot \frac{1}{d^2}$
 - $\geq \frac{1}{d} - \frac{1}{d(d+1)}$
 - $= \frac{1}{d+1}$

What does that mean for Karger-Stein?

Claim: for all d , $p_d \geq \frac{1}{d+1}$

What does that mean for Karger-Stein?

Claim: for all d , $p_d \geq \frac{1}{d+1}$

- For $d = 2\log(n)$
 - that is, d = the height of the tree:

$$p_{2\log(n)} \geq \frac{1}{2\log(n) + 1}$$

What does that mean for Karger-Stein?

Claim: for all d , $p_d \geq \frac{1}{d+1}$

- For $d = 2\log(n)$
 - that is, d = the height of the tree:

$$p_{2\log(n)} \geq \frac{1}{2\log(n) + 1}$$

- aka,

$$\Pr[\text{Karger-Stein is successful}] = \Omega\left(\frac{1}{\log(n)}\right)$$

Altogether now



- We can do the same trick as before to amplify the success probability.
 - Run Karger-Stein $O\left(\log(n) \cdot \log\left(\frac{1}{\delta}\right)\right)$ times to achieve success probability $1 - \delta$.

Altogether now



- We can do the same trick as before to amplify the success probability.
 - Run Karger-Stein $O\left(\log(n) \cdot \log\left(\frac{1}{\delta}\right)\right)$ times to achieve success probability $1 - \delta$.
- Each iteration takes time $O(n^2 \log(n))$
 - That's what we proved before.

Altogether now



- We can do the same trick as before to amplify the success probability.
 - Run Karger-Stein $O\left(\log(n) \cdot \log\left(\frac{1}{\delta}\right)\right)$ times to achieve success probability $1 - \delta$.
- Each iteration takes time $O(n^2 \log(n))$
 - That's what we proved before.
- Choosing $\delta = 0.01$ as before, the total runtime is
$$O(n^2 \log(n) \cdot \log(n)) = O(n^2 \log^2(n))$$

Much better than $O(n^4)$!

What have we learned?

- Just repeating Karger's algorithm isn't the best use of repetition.
 - We're probably going to be correct near the beginning.
- Instead, Karger-Stein repeats when it counts.
 - If we wait until there are $\frac{n}{\sqrt{2}}$ nodes left, the probability that we fail is close to $\frac{1}{2}$.
- This lets us (probably) find a global minimum cut in an undirected graph in time **$O(n^2 \log^2(n))$** .
 - Notice that we can't do better than n^2 in a dense graph (we need to look at all the edges), so this is pretty good.

Recap

- Some algorithms:
 - Karger's algorithm for global min-cut
 - Improvement: Karger-Stein
- Some concepts:
 - Monte Carlo algorithms:
 - Might be wrong, are always fast.
 - We can boost their success probability with repetition.
 - Sometimes we can do this repetition very cleverly.

Acknowledgement

- Stanford University