

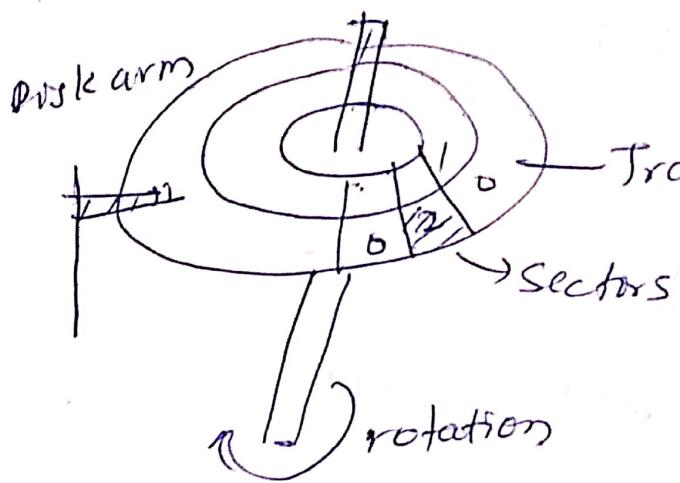
II(11B20)

4NF

## DBMS &

### Disk structure +

Each disk platter has a flat, circular shape



Sector = 512 bytes  
tracks = 50,000+  
Platter = 1 to 5

Note + Block  
contains sectors

### Data storage & Indexing + R/W in terms of block

(1) Disk structure

(2) Indexing

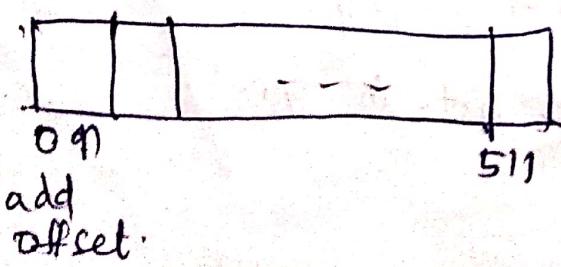
(3) B<sup>+</sup> tree

(4) Hashing.

=

Block Add = (Track No, Sector No)

Block Size = 512 bytes



# Student

SID	Sname	Course	Dept	add
1	S1			
2	S2			
3	S3			
4	S4			
5	S5			
6	:			
7	:			
8	:			
100	S100			

# Student

SID = 10 bytes

Sname = 5 bytes

Course = 8 bytes.

Dept = 10 bytes

add = 50 bytes

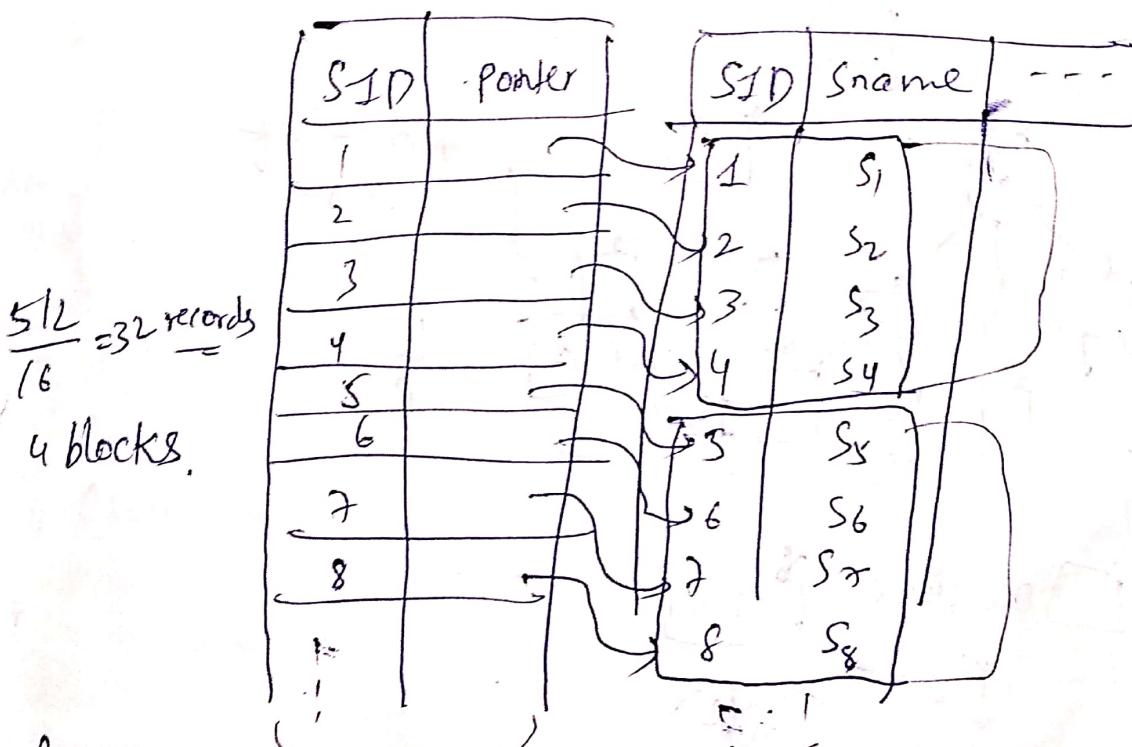
Total = 128 bytes

Assume block size - 512 bytes

(Q) How many student records we can store in one block?

$$\frac{512}{128} = 4 \text{ records} \quad \text{1 block} = 4 \text{ records}$$

→ We need 25 blocks for 100 records

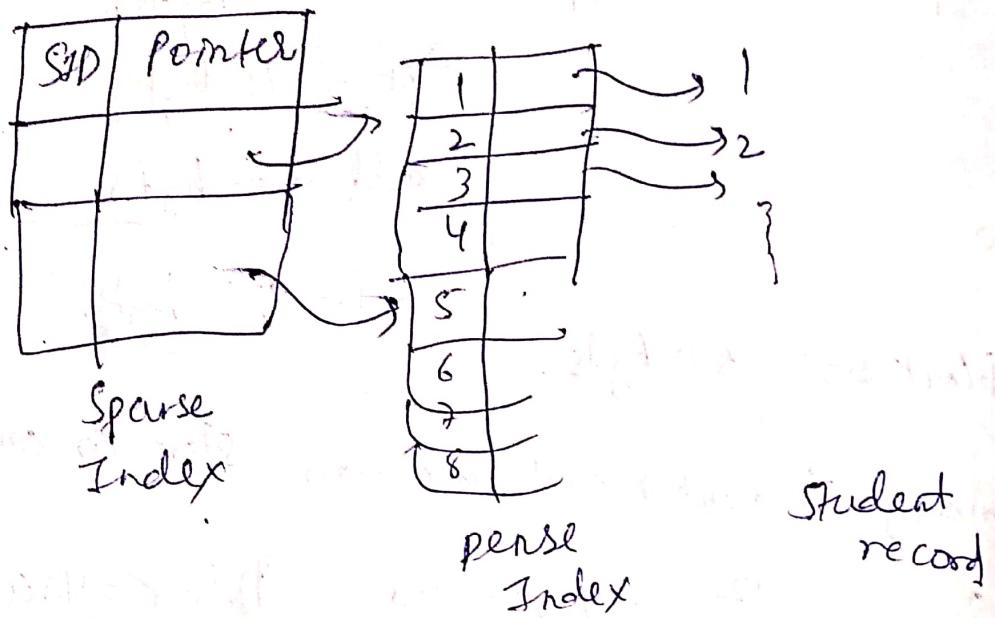
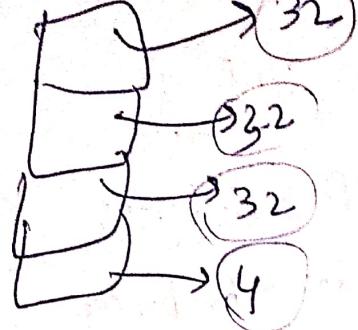


SID = 10 bytes

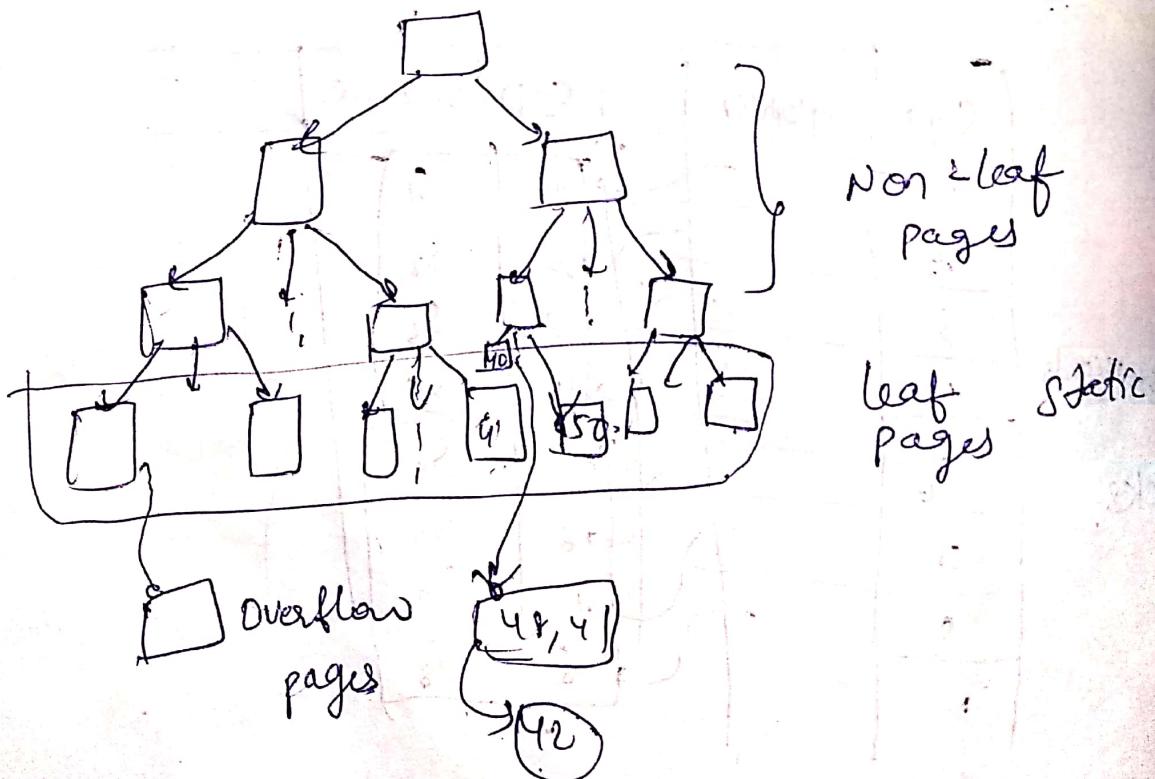
Pointer = 16 bytes  
16

Dense  
Index

100 records



### Index sequential Access model (ISAM)



## B+tree

→ It is a balanced tree

\* Internal nodes direct the search

\* Leaf nodes contain the data entry.

\* Leaf nodes organized into doubly linked list

## Characteristics

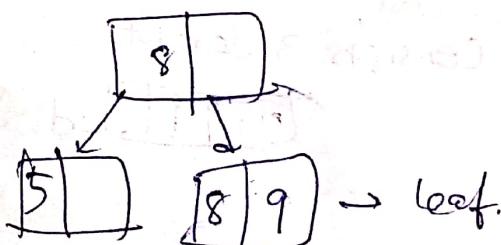
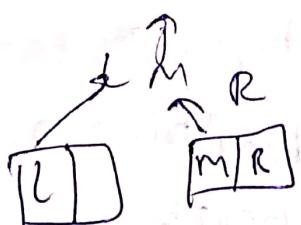
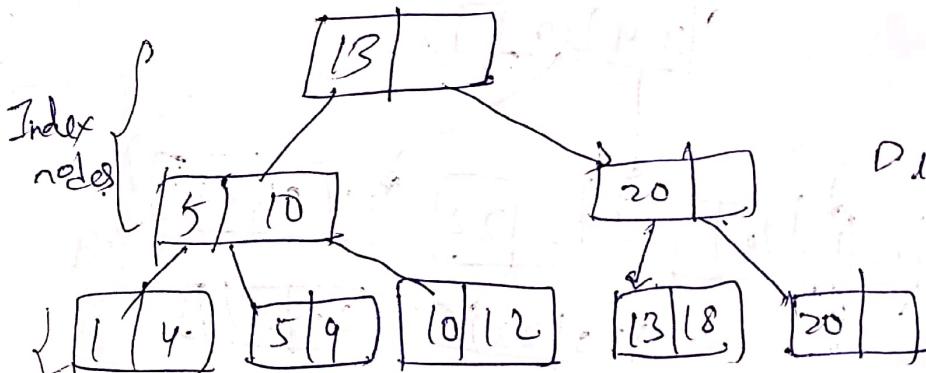
(1) Minimum occupancy is 50%  
each node, except root

(2) each node contains  $m$  entries  
where  $d \leq m \leq 2d$  entries.

[ $d$  - order of tree]

Ext order  $d = 1$

$k_1 \leq k \leq k_r$  - Index



Leaf node

L M R



Index node

L M R



650

799

499

Lab r

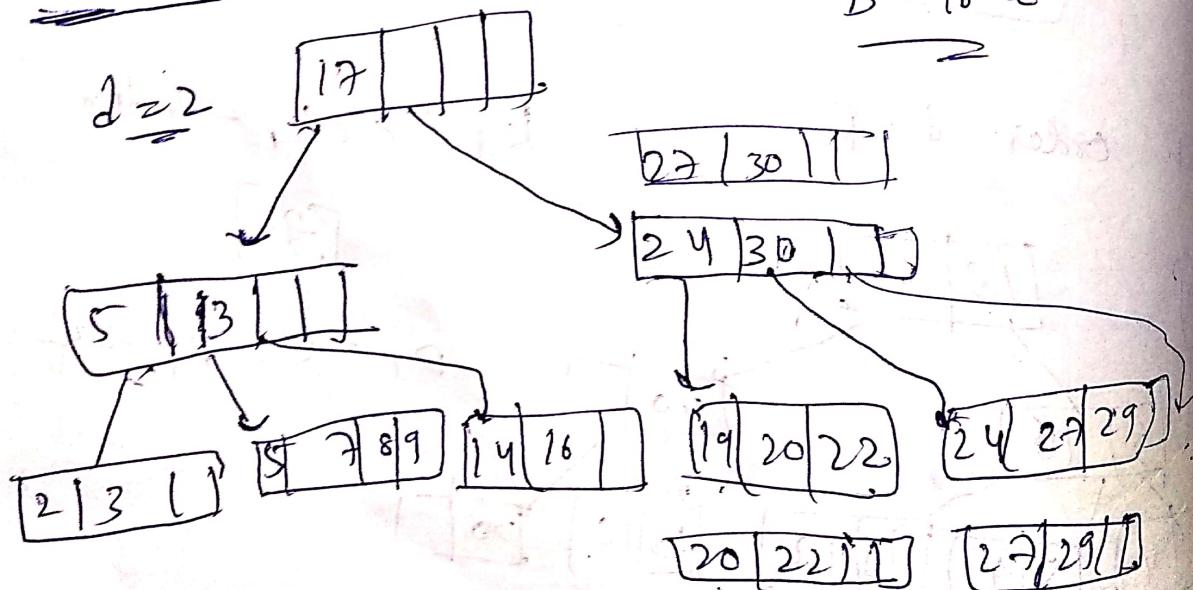
product

Error  
02000 → NO Row Found

Create trigger ensure direction before insertion

table-name FOR EACH ROW

BT tree



Deleted 19, 20

→ Redistribution

→ merge

Delete 27.

22 24 11 cd

Sibling  
Consists 3 elements

29 11 13 cd

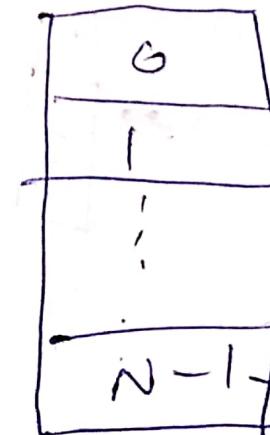
→ Static

→ Dynamic

### Static Hashing +

\* 0, 1, ..., N-1 buckets

\* Buckets store data entries



primary  
bucket  
pages

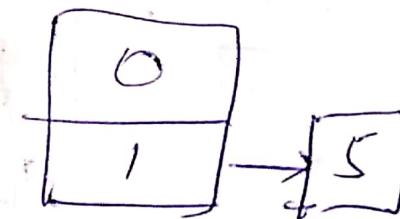
overflow  
pages

Search Apply hash,  $\lambda$

$$h(r) = r \pmod N$$

$$h(3) = 3 \pmod 2 = 1$$

$$h(5) = 5 \pmod 2 = 1$$



• Use of directory of pointers to buckets

→ Double the size of no. of buckets

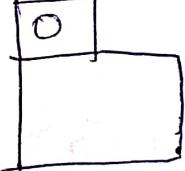
just by doubling the directory and split  
the bucket overflowed

$$\text{Expt } h(r) = r \bmod 2 = \{0, 1\}$$

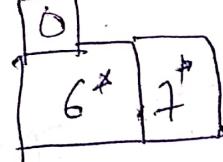
Q1 Add 6, 7, 8, 10.

bucket size is 2

Global  
So it



Local



d - depth

{ Global  
local }

6 - 110

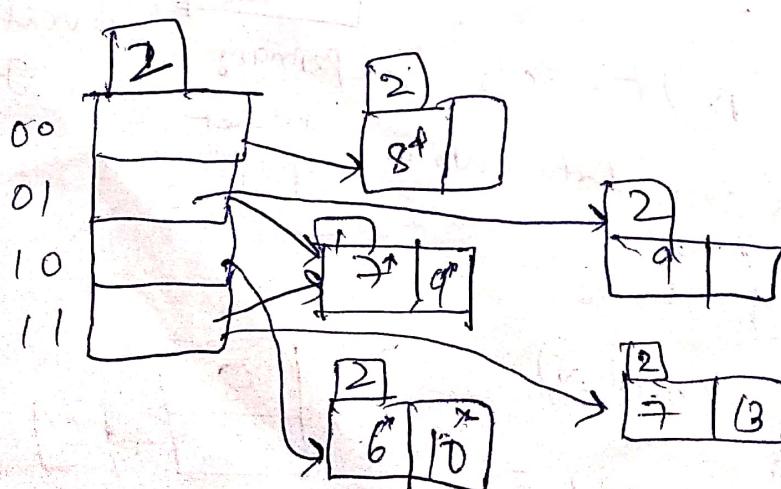
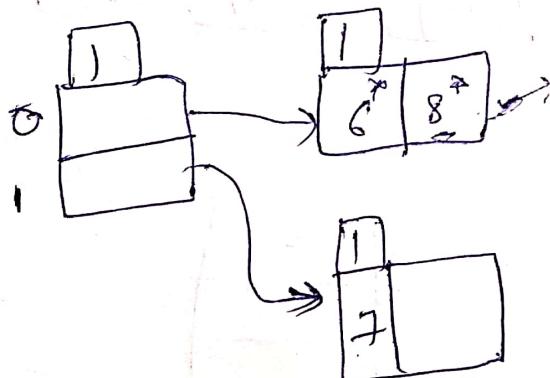
7 - 111

8 - 1000

10 - 1010

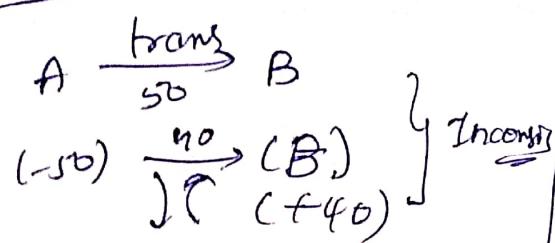
9 - 1001

13 - 1011



# 11/12/2023 Transaction Management

## Consistency & Isolation



User program's logic

$T_1$  &  $T_2$  concurrently.

$T_1 ; T_2$  } should be  
 $T_2 ; T_1$  } equivalent

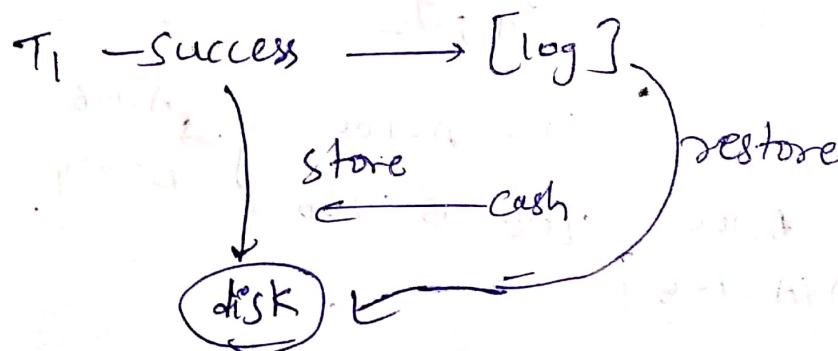
Isolation

## Atomicity & Durability

- Transaction incomplete
- Aborted
- System crash
- Read unexpected data
- Unable to access disk



[log] - Used, remember & restore



## Transactions & Schedules

Transaction - list of actions

Actions : Read, write of database object

O - object

$R_T(O)$  - action of T is read object O.

$w_T(O)$  - action of T is write object O.

Transaction must specify : commit, abort.

# Schedule (read, writes, abort, commit)

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	Result
R(A)		200
W(A)		200
	R(B)	200
	W(B)	200
R(C)	Commit	200
W(C)		200
Commit		200

T<sub>1</sub> & T<sub>2</sub> conflict each other

→ WR

T<sub>1</sub>: transfer Rs. 100. from A to B

→ RW

T<sub>2</sub>: increment both A & B by 6

→ WW

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
A: 200 R(A)	
A: 100 W(A)	
R(A)	A: 100
W(A)	A: 106
R(B)	B: 300
W(B)	B: 316
commit	
B: 318 R(B)	
(B: 418) W(B)	
Commit	

<u>T<sub>1</sub>; T<sub>2</sub></u>
100 A - 100, $\Rightarrow$ A: 100
400 B + 100, $\Rightarrow$ B: 400
100 A - 100, $\Rightarrow$ A: 100
400 B + 100, $\Rightarrow$ B: 400

Unrepeatable reads  $\rightarrow$  (rw conflict).

A - No. of copies available for a book.

Integrity constraint  $A \geq 0$

T - places an order, first read A.

then check  $A > 0$

then decrement 1

$T_1 : R(A)$  and sees  $A = 1$

$T_2 : R(A)$  and sees  $A = 1$

$w(A)$ , decrement by 1,  $A = 0$

commit

$T_1 : w(A)$  and decrement by 1,  $A = -1$ .

Overwriting Uncommitted Data (WW conflict)

A & B are two employees, their salaries must be kept equal

$T_1 : \text{sets } A \text{ and } B \text{ Salaries RS. 1000}$

$T_2 : \text{sets } A \text{ and } B \text{ Salaries RS. 2000}$

$T_2 : w(A)$

$T_1 : w(B)$

$T_2 : w(B)$  - commit

$T_1 : w(A)$  - commit

A	B
2000	1000
	2000 ←
	1000 ←

## Schedule Involving Aborted Transactions

	T <sub>1</sub>	T <sub>2</sub>
R(A)		
W(A)		
		T <sub>1</sub> : deduct Rs 1/100 from A
		T <sub>2</sub> : read A & B
R(A)		
W(A)		
R(CB)		
		add 6 to B's balance
W(B)		Commit
		T <sub>1</sub> : Aborted
		Commit
		Left 6.99 balance (A)
		Left 6.99 balance (B)
Aborted		

## Unrecoverable Schedule

Strict Two-phase locking (Strict 2PL) → By using this we can avoid unrecov— schedule

### Strict 2PL

lock → shared lock S<sub>Tf0</sub> → read  
lock → exclusive lock X<sub>f(s)</sub> → write, read

Expt	T <sub>1</sub>	T <sub>2</sub>
	R(A)	
	W(A)	
	R(A)	
	W(A)	
	R(B)	
	W(B)	
	commit	
	R(B)	
	W(B)	
	commit	

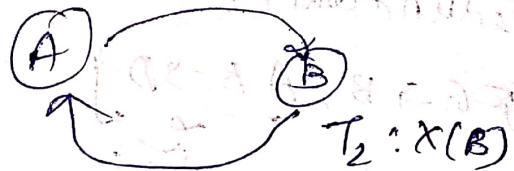
	T <sub>1</sub>	T <sub>2</sub>
X(A)		
R(A)		
W(A)		
		(a) W : T <sub>1</sub>
X(B)		
R(B)		
W(B)		
commit		
		Request X(A), not allowed to support
		R(A)
		W(A)
		X(B)
		R(B)

Once committed, the lock can be transferred to another transaction

Ext	T <sub>1</sub>	T <sub>2</sub>
	S(A)	
	R(C)	(S, A) scan
Serial Interleaving	S(A)	
	R(A)	
	X(B)	
	R(B)	
	W(B)	commit
	X(C)	
	R(C)	
	W(C)	
	commit	

### Deadlock

T<sub>1</sub> : X(A)



T<sub>1</sub> waits until T<sub>2</sub> completion  
and T<sub>2</sub> waits until T<sub>1</sub> completion

This is a deadlock!

Ext	T <sub>1</sub>	T <sub>2</sub>
	X(A)	
	R(A)	
	W(B)	
		X(B)
		R(B)
		W(B)
		X(A) ← request held in queue
	request → X(B)	
	hold queue	

→ If deadlock happens,  
we have to terminate it

→ We have to see that  
deadlock doesn't happen

→ We take preventive  
mechanisms.

20/11/2020

T<sub>1</sub>

X(A)

R(A)

W(B)

request → X(B)

hold queue

T<sub>2</sub>

lock graph

deadlock → 280

Q3 at 4 AM (4 AM = 76)

Q3 at 4 AM (4 AM = 76)

Q3 at 4 AM (4 AM = 76)

X(B)

R(B)

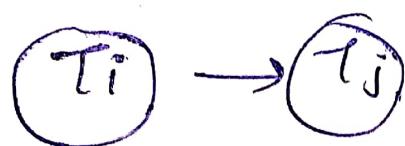
W(B)

X(A) ← request

hold queue

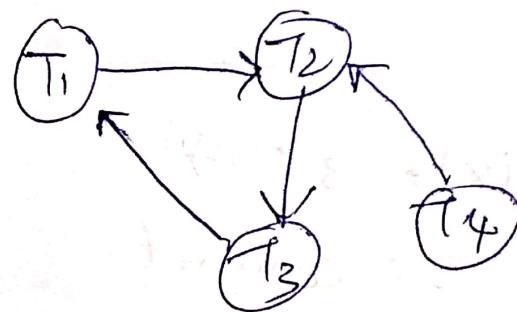
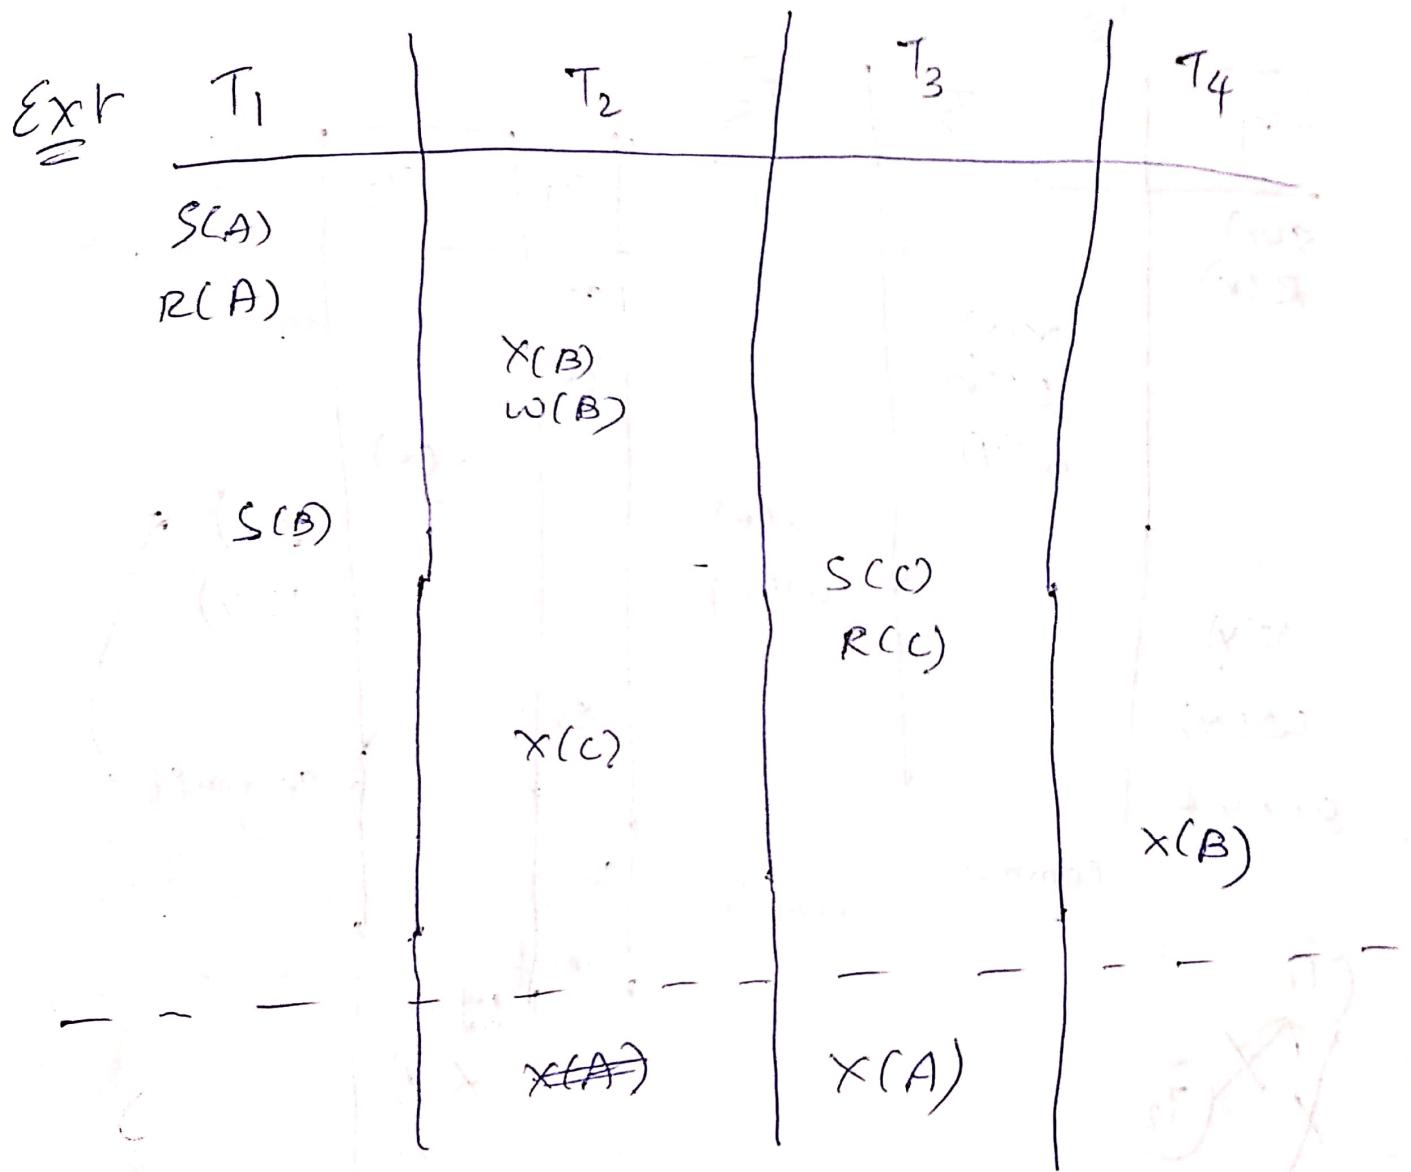
wait-for graph

the lock manager maintains a structure called a wait-for graph to detect deadlock cycles.



$T_i$  is waiting for  $T_j$  to release a lock

- (i) Lock manager adds edge to the graph when it queues lock request  
 ii) and removes the edge when it grants lock request

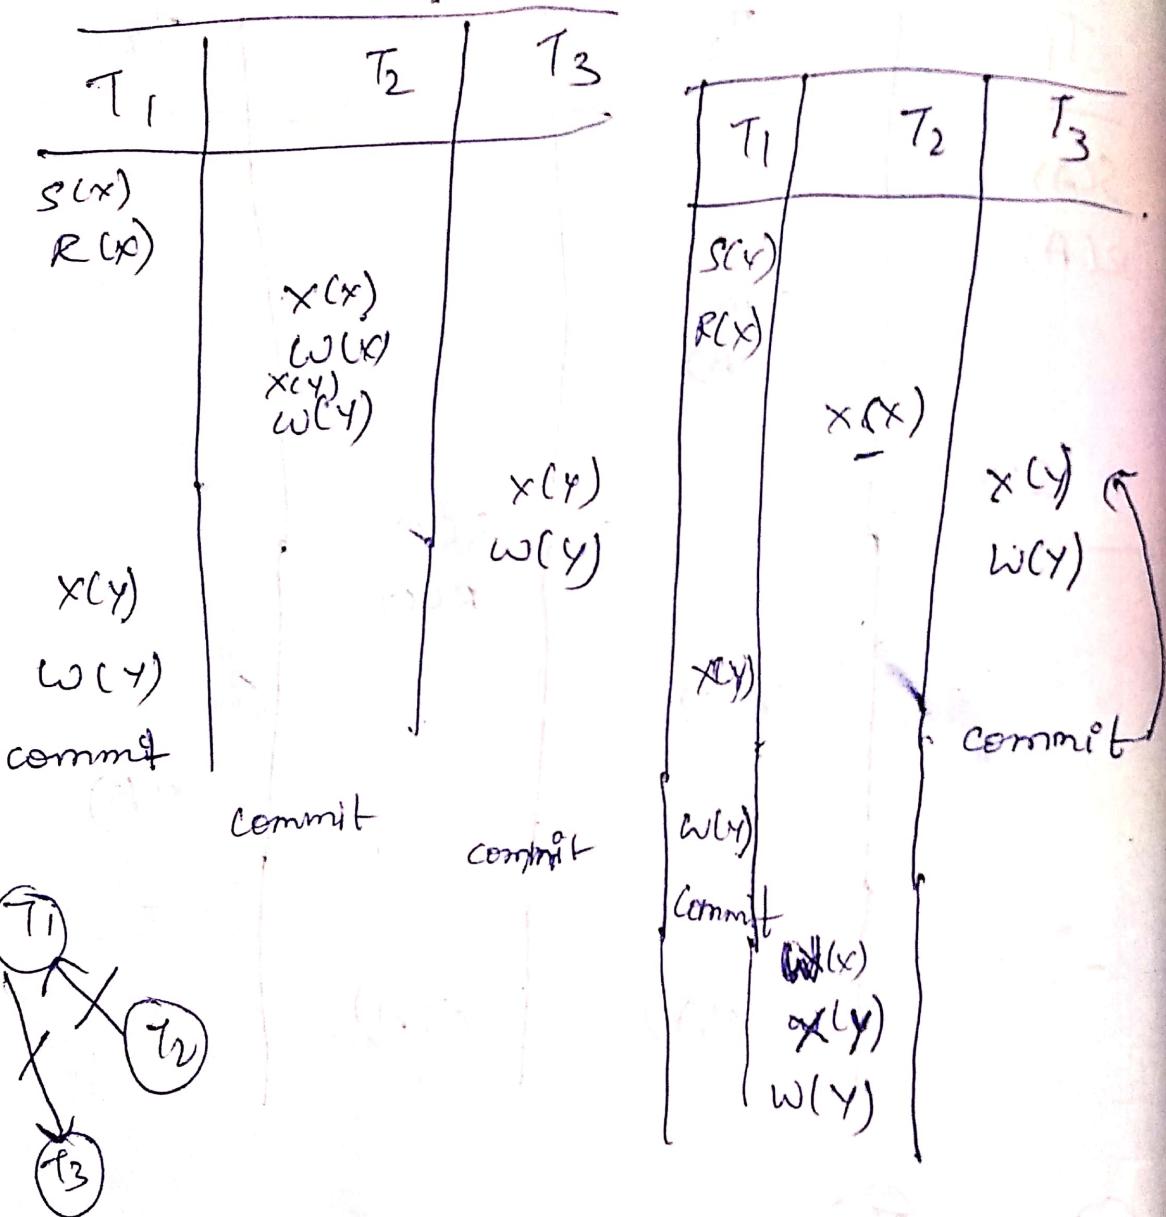


Q)  $T_1 : R(x), T_2 : w(y), T_3 : w(y)$

$T_1 : w(y)$

$T_1 : \text{commit}, T_2 : \text{Commit}, T_3 : \text{Commit}$

Strict 2PL and wait-for graph.



Q)  $T_1 : R(x), T_2 : w(y), T_3 : w(y),$

$T_1 : w(y), T_1 : \text{Commit}, T_2 : \text{Commit}$ .

$T_3 : \text{Commit}$

